APERTURE METHOD™

AGENT BUILD SPEC · VERSION 5 · SOP-ANCHORED

The Agent Build Spec

Six agents that execute the Method SOPs

The question this answers: ***What exactly do I build, and how does each agent execute the procedure I already wrote?***

PREPARED FOR

Fenwick How

Founder, The Aperture Method

August 7, 2026 · Anchored to Method SOPs v1.0 · Supersedes v1, v3, v4

BUILD-READY — Schemas, run contracts, prompt architecture, and gates. The Aperture Score weights and rubric criteria in Part One are proposals to react to, not decisions.

CONTENTS

Contents

If this list is blank, press Ctrl+A then F9 in Word to build it.

PART ONE

The Program

How the system works: the two governing rules, the corrections the SOPs forced, the shared data foundation, and the contracts every agent obeys.

1 · How to use this document

Three parts. Part One is how the system works. Part Two is the six agents, one at a time. Part Three is what gets built, in what order.

Part	What is in it	Read it when
One · The Program	The two rules, the corrections, the data model, the run contract, the prompt architecture, the gates, the privacy layer, the Aperture Score.	You want to understand or change how the system behaves.
Two · The Team	Fenwick’s profile, then a profile page and a full step-by-step build spec for each of the five agents, anchored to its SOP chapter.	You are building or debugging one agent.
Three · The Build	Repository layout, the phased roadmap, the proposed SOP 00, and what needs deciding before code.	You are planning the work or checking progress.

Conventions used throughout

Marker	Means
FENWICK	Fenwick. Meetings, judgment, the call. Irreducible — no agent shortens it.
AGENT	The model reasons, drafts, triangulates, proposes.
TOOL	Deterministic code. Produces the number of record. Agents never compute in prose.
HYBRID	The agent proposes; a person decides.

2 · The SOPs are the program

> **The Method SOPs are not documentation about the agents. They are the agents’ program.**
> Every chapter already has the exact shape an agent needs: an objective, declared inputs, a numbered process where each step states its own output, a named tool-and-framework list, and a quality checklist to pass before anything reaches a client.

That is a specification. Most teams building this would spend six weeks producing a worse version of what you already wrote. So the build takes it literally.

What the SOP has	What it becomes in the build
The chapter text	sops/01_business_xray.md … 06_atlas.md — held verbatim, in git, loaded into the agent’s prompt at runtime. Edit the manual and the agent changes. There is no second copy of the procedure to drift.
"Output of this step"	A typed intermediate artifact. X-Ray Step 3 declares "seven scored lenses with evidence" — that is a schema, and the runner refuses to advance without it.
The quality checklist	A set of executable assertions. "Every headline number reconciles to the client’s financials" becomes assert_reconciles(), not a hope.
"Tools, analyses & frameworks"	The agent’s tool registry. Nothing outside that list is callable.
"When to run it & prerequisites"	The dispatcher’s routing rules, plus the standalone fallback in Section 7.

3 · The two rules everything obeys

> **Rule one — agents compress the mechanical 80%. They do not touch the judgment 20%.**
> Agents produce findings and artifacts. Never conclusions and advice. The last mile is human, and it is the mile the client pays for.

Your site says it out loud: "The difference is the analysis, not the dashboard." The competitive position is that an owner-run business gets MBA-level interpretation — someone who looks at a Profit Map and says "your problem is not the fourth location, it is that half your new patients never come back." Any architecture that lets a model write that sentence unsupervised destroys the product.

> **Rule two — every agent stands alone.**
> No agent may depend on another having run. Every agent reads the shared Intake Store. Chaining enriches; it never gates. Five specialists, each hireable solo or as a team.

The rules collide exactly once — Compass is a synthesis agent, and synthesis wants inputs. Section 7 resolves it.

4 · Four corrections the SOPs forced

4.1 · Business X-Ray and Aperture Insights are two SOPs, not one

An earlier version collapsed them into a single agent. The SOP separates them, and it is right to.

Business X-Ray (SOP 01) is the fixed-fee diagnostic front door — a standalone product with its own scope, its own readout, and its own commercial job: earn the right to a larger engagement. Aperture Insights (SOP 02) is the opening phase of a full engagement; its first step is literally "confirm or produce the X-Ray," and its real work is the three steps after that.

They are the same lens and the same rubric doing two different jobs. Resolution: one agent, two modes.

```
Lena
 |-- mode="xray"      -> SOP 01 -> Business X-Ray + Aperture Score
 `-- mode="insights"  -> SOP 02 -> reuse/extend X-Ray + validated constraint
                                   + success measures + sequenced plan
```

4.2 · What the agents actually compress

An earlier version framed the target as "X-Ray turnaround: weeks to days." The SOP shows that is wrong, and the honest version is better.

The X-Ray’s 2–3 week duration is not analysis time. It is a 45–60 minute kickoff, a data-return wait on the client, 3–5 interviews of 45 minutes each that have to be scheduled, and a 60-minute readout. Those steps are calendar-bound and irreducibly human.

	What it is	Can an agent compress it?
Kickoff, interviews, readout	~5–7 hours of meetings, spread across weeks by scheduling	No
Data return	Client-side wait	Only indirectly — a sharper, tailored data request reduces round trips
Everything else	Assemble · reconcile · score · triangulate · draft	Yes — days to hours

The honest claim

The X-Ray stays a 2–3 week engagement, because that is how long it takes to get on people’s calendars. What changes is that your hands-on time inside it drops from most of two weeks to a day of review and judgment — which is what lets you run five or ten concurrently. That is the actual business result, and it is a better story than "faster" because it is checkable.

4.3 · Standalone prerequisites need declared fallbacks

The SOPs state prerequisites — Analytics "run after Insights," Compass "run after Analytics and/or Intelligence," Atlas "run after Compass" — and also say any segment can run standalone. Both are true, and the build makes them true simultaneously. Section 7.

4.4 · Every score must carry its version

SOP 06 Step 5 requires that the re-scored Aperture Score "uses the identical baseline method." That is an engineering requirement: every score carries a weights_version and a rubric_version, and if either changes, historical scores are recomputed or explicitly flagged non-comparable. Otherwise the client’s progress chart is quietly meaningless.

5 · The system at a glance

Layer	What happens
Client intake	Transcripts (Otter) · financials · operational data · customer records
Intake agent	Backstage, no persona, no client contact. Parse · normalize · reconcile · geocode · gap-report
Shared Intake Store + Ledger	One clean reconciled dataset per engagement. Every figure carries provenance
Scrub layer	Identities tokenized. The token map stays local
The five agents	Lena · Marcus · Natalie · Aria · Thomas — each reads the Store directly, never each other as a dependency
Unscrub layer	Real values restored locally
Human review	/admin — approve or send back, with a diff
Delivery	Branded deliverable released to /portal

The dispatcher, not a pipeline

A deterministic state machine, not an LLM. It reads which modules the client bought and starts those agents; independent agents run concurrently. It owns routing, gates, retries, and the audit trail. Routing a purchased scope is a finite known process, and a state machine cannot improvise.

6 · The agent roster

Human-led, AI-enabled. Fenwick How at the helm, then the five Aperture Method™ AI agents — Lena, Marcus, Natalie, Aria and Thomas. The personas are internal: they shape how each agent is prompted and what it optimizes for. Clients see Aperture, not names.

SOP	Agent	Persona	Deliverable	Duration	Standalone?
00 Intake & Acceptance	Fenwick How	Architect & Managing Principal	Engagement data store · Reconciliation Report · Gap Memo	continuous	the gate
01 Business X-Ray	Lena mode=xray	The Strategist	Business X-Ray + Aperture Score	2–3 wk	Yes — the wedge
02 Aperture Insights	Lena mode=insights	The Strategist	Validated constraint · success measures · plan	2–3 wk	Engagement-opening
03 Aperture Analytics	Marcus	The Analyst	Profit Map + scenario model	3–4 wk	Yes
04 Aperture Intelligence	Natalie	The Explorer	Customer & Market Map + ArcGIS package	2–3 wk	Yes
05 Aperture Compass	Aria	The Navigator	Opportunity Matrix + Focus Plan	1–2 wk	Yes, degraded
06 Aperture Atlas	Thomas	The Operator	Live Scoreboard + KPI system	ongoing	Yes

7 · Standalone fallbacks

The rule

Every agent has a required core it reads from intake, and enrichment inputs it uses when present. Where an SOP prerequisite comes from a module that was not purchased, the agent derives a minimal substitute and labels it.

Bought alone	SOP prerequisite absent	Fallback	Labelled as
Analytics	Validated constraint + success measures	Marcus derives a working constraint hypothesis from the financials and states it	Derived from financials — not owner-validated
Intelligence	Constraint + success measures	Natalie runs the full spatial arc from intake; frames output as "where demand is," not "what to do"	Spatial finding — strategic priority not assessed
Compass	Shortlisted moves from Analytics / Intelligence	Aria generates candidates from intake + X-Ray if present, scores what is evidenced, names the rest as open questions	Scored on available evidence — items requiring other modules listed separately
Atlas	Focus Plan + success measures	Thomas baselines on the client's stated goals + intake KPIs; establishes period zero	Baseline established at subscription — no prior engagement

Every deliverable carries a "What this module assumed" box

Two reasons. It is honest — the client knows exactly what was and was not established. And it is the cleanest upsell you will ever write, because the box is literally a list of what the next module would resolve.

In code: ModuleResult.derived_assumptions[] is non-optional. The renderer refuses to build a deliverable when the array is unset — empty is a valid answer, unset is a bug.

8 · The privacy backbone

Two thin local layers bracket every external model call. A local pre-processor detects and tokenizes identifiers; the lookup table mapping token to real value never leaves your machine. When results return, the same table swaps the real values back in locally. For a fully confidential engagement, a per-engagement switch routes that module to a local model instead, with zero external calls.

8.1 · Tokenize identities, never quantities

If Marcus receives $TOKEN_A instead of $6.44M he cannot compute a margin, rank a scenario, or notice that EBITDA fell from 22% to 15%. The analysis is the magnitudes.

Tokenized	Passed through real
Client and person names · business names · street addresses · emails · phone numbers · account numbers	Revenue · margins · counts · dates · ratios · volumes

A number with no name attached is not identifying, and it is the entire substance of the work. If a specific figure genuinely is identifying on its own, that engagement runs local-only — do not half-scrub it.

8.2 · Geography is computed locally, always

SOP 04 Step 1 already says "geocode, then drop street-level PII." The build enforces it: spatial tools run on your machine against real coordinates and emit derived geometry — trade-area polygons, density surfaces, H3-binned aggregates. Only the derived layer and its demographic profile go out for interpretation. Natalie reasons about "Area 7: 41k households, median income $118k, 12-minute drive time"; she never sees a street address. The spatial math is deterministic, so nothing is lost.

8.3 · The token map is deterministic and engagement-scoped

Generated once by Intake and reused by every agent for that engagement. source_ref values are opaque identifiers containing no PII and pass through untouched, so provenance survives the round trip and the scrub layer never fights the citation contract.

> **Net posture**
>
> API-with-scrubbing as the default. Local-only as the premium confidential option, selected per engagement. Both paths produce identical deliverables — the client never sees which one ran.

9 · The Aperture Score

The SOP settles the structure: seven lenses, each scored 1–100 with evidence, combined by weight into one composite, recorded so it can be re-scored identically later. What remains open is the weights and the per-lens criteria. Below is a starting proposal to react to — not a decision, and the weights are yours.

Lens	Weight	Why
Finance — margin, cash, unit economics, working capital	20	Cash is where owner-run businesses actually die
Customers & retention — acquisition, repeat, lifetime value	20	Most common binding constraint in this ICP; the Lumina case turns on it
Operations — capacity, throughput, bottlenecks	15	Where the constraint physically lives
Market & competition — position, demand, local dynamics	15	The Intelligence hook, and the ceiling on everything else
Processes — how work gets done, where it breaks	10	Real, but usually downstream of the four above
Technology & data — systems, data quality, automation	10	The AI-adoption hook; rarely the binding constraint by itself
Leadership & team — key-person risk, roles, capability	10	High variance; hardest to evidence from data alone

Design requirements regardless of the numbers

- Weights sum to 100 and are stored as weights_version. Every score carries it.
- Each lens has 4–6 declared sub-criteria with a scoring band and a required evidence type, so two runs on the same data produce the same score.
- A lens with insufficient evidence scores as insufficient_evidence, not as a middling 50. A guessed lens quietly corrupts the composite and the client's trend line.
- Re-scoring in Atlas uses the stored weights_version and rubric_version. If either changes, historical scores are recomputed or flagged non-comparable — never silently mixed.

> **Worth a working session before Phase 1 ships**
> This is the most durable IP in the Method and the one number a client will quote back to you for years.

10 · Data model

DuckDB, one file per engagement, local.

```
-- identity & scope
engagements(engagement_id PK, client_id, purchased_modules JSON, privacy_mode,
            fiscal_year_end, status, created_at)
locations(location_id PK, engagement_id, label, opened_on, h3_index)
service_lines(service_id PK, engagement_id, label, category)

-- provenance spine
source_files(file_id PK, engagement_id, filename, sha256, parser, received_at)
ledger_facts(fact_id PK, engagement_id, entity, metric, period, value, unit,
             source_file_id, source_locator, method, confidence)

-- operational data
transactions(txn_id PK, engagement_id, date, customer_id, location_id,
             service_id, qty, gross, discount, net, cost)
pnl(engagement_id, period, location_id, account, amount)
customers(customer_id PK, engagement_id, first_seen, last_seen, zip,
          h3_index, segment, clv)      -- lat/lon LOCAL ONLY, separate table
interviews(interview_id PK, engagement_id, role, transcript_ref, scrubbed_ref)
kpi_feed(engagement_id, period, kpi_id, value, source_ref)

-- method objects
lens_scores(engagement_id, run_id, lens, score, rationale, evidence_refs JSON,
            rubric_version)
aperture_scores(engagement_id, run_id, scored_at, overall, by_lens JSON,
                weights_version, rubric_version)
constraints(engagement_id, run_id, statement, evidence_refs JSON,
            disconfirming_refs JSON, status)      -- proposed | validated
success_measures(engagement_id, measure_id, label, baseline, target_low,
                 target_high, unit, source_ref, cadence)
findings(finding_id PK, engagement_id, module, claim, evidence_refs JSON,
         confidence, status)
opportunities(opportunity_id PK, engagement_id, label, origin_module,
              value, complexity, risk, constraint_impact, weighted_score,
              weights_version, evidence_refs JSON)

-- run control
runs(run_id PK, engagement_id, module, mode, status, started_at, ended_at,
     model, prompt_version, sop_version)
artifacts(artifact_id PK, run_id, type, path, template, version, status)
approvals(approval_id PK, artifact_id, approver, decision, diff_ref, at)
gate_results(run_id, check_id, passed, detail)
gaps(engagement_id, module, field, severity, status)
```

Never in this file

The token map. It lives at engagements/<id>/tokens.json — git-ignored, local-only, never synced anywhere.

The provenance envelope

Every client-facing number is this shape, end to end. The renderer walks the document tree, and any numeric literal without an envelope fails the build.

```
{
  "label": "FY2025 revenue",
  "value": 6440000,
  "unit": "USD",
  "period": "FY2025",
  "source_ref": ["fact:8a3f21", "fact:9c2104"],
  "method": "sum(pnl.amount where account_type='revenue')"
}
```

11 · Run contract

Every agent has the same signature. This is what makes modularity real rather than claimed.

```
def run(engagement_id: str,
        module: Module,               # xray|insights|analytics|intelligence|compass|atlas
        mode: str = "standalone",     # standalone | engagement
        upstream: dict | None = None, # optional enrichment refs - NEVER required
        privacy: str = "scrubbed_api" # scrubbed_api | local_only
        ) -> ModuleResult: ...
```

```
// ModuleResult
{
  "module": "analytics",
  "sop_version": "1.0",
  "mode": "standalone",
  "status": "complete",
  "findings":  [{"id":"f1","claim":"...","evidence_refs":["fact:..."],"confidence":0.8}],
  "figures":   [{"label":"...","value":0,"unit":"USD","source_ref":
["fact:..."],"method":"..."}],
  "artifacts":
[{"type":"profit_map","path":"...","template":"analytics_doc","version":"1.2"}],
  "derived_assumptions": [
    {"input":"binding_constraint","source":"derived_from_financials",
     "statement":"...","label":"Derived - not owner-validated"}
  ],
  "gaps":        [{"field":"cost_allocations","severity":"medium"}],
  "gate_results":[{"check":"model_reconciles_to_the_dollar","passed":true}],
  "handoff":     {"to":["compass"],"payload_refs":["opportunity:..."]}
}
```

Hard rules enforced by base.py, not by prompt

- upstream is always optional. An agent that raises on missing upstream fails the isolation test and cannot ship.
- status="blocked" is a legitimate terminal state. Partial artifacts are never written.
- Every figures[] entry must carry a resolvable source_ref. The renderer re-validates independently.
- derived_assumptions must be present. Empty is valid; unset is a bug.

12 · Prompt architecture

Four blocks, assembled per run. Only block 1 is bespoke per agent.

Block	Source	Contents
1 · PERSONA	prompts/<agent>.persona.md	Who this agent is, what it optimizes for, and its named failure mode. The profile pages in Part Two are the source text.
2 · PROCEDURE	sops/0N_<module>.md — YOUR SOP, VERBATIM	"You are executing this documented procedure. Do not improvise steps. Each step declares its output; produce exactly that."
3 · CONTRACT	prompts/contract.md — identical for all six	Findings, never conclusions. Every figure via a tool call with source_ref; never compute in prose. Cite evidence for every claim. If you cannot complete a step, return blocked. Name every derived-not-validated assumption.
4 · CONTEXT	runtime	Scrubbed data slices · available upstream artifacts · mode · privacy mode · the module's own quality checklist, used as the self-check before returning.

Why block 2 is the design decision worth defending

Loading your SOP verbatim means the procedure has exactly one home. You edit the manual; the agents change with it. There is no prompt-engineering session where someone's paraphrase of Step 5 quietly diverges from what the manual says. Version both together — sop_version is stamped on every run.

13 · Quality checklists as executable gates

Every checklist item in the manual becomes a named assertion in gates/checklists.yaml. Blocking gates stop the run; advisory gates surface in the /admin review queue.

```
xray:
  - id: headline_numbers_reconcile
    assert: reconciles(artifact.figures, ledger.pnl, tolerance=0.005)
    blocking: true
  - id: each_lens_has_evidence
    assert: all(len(l.evidence_refs) > 0 for l in lens_scores)
    blocking: true
  - id: score_is_reproducible
    assert: recompute(aperture_score) == aperture_score.overall
    blocking: true
  - id: constraint_is_singular_specific
    assert: constraint.status == "validated" and constraint.evidence_refs
    blocking: true
  - id: readable_by_non_analyst
    assert: readability_check(report) and jargon_scan(report)
    blocking: false          # advisory - flags for human review
  - id: prescription_names_components
    assert: prescription.components and all(c.rationale for c in ...)
    blocking: true
  - id: pii_handled
    assert: no_identifier_in_outbound_log(run_id)
    blocking: true

analytics:
  - id: model_reconciles_to_the_dollar
    assert: reconciles(financial_model, ledger.pnl, tolerance=0.0)
    blocking: true
  - id: allocations_documented
    assert: exists(artifact.allocation_basis) and allocation_basis.rationale
    blocking: true
  - id: no_false_precision
    assert: profit_cuts_within_data_granularity(profit_map, transactions)
    blocking: true
  - id: assumptions_explicit_and_adjustable
    assert: scenario_model.inputs_are_named_and_editable()
    blocking: true
  - id: recommendations_tie_to_constraint
    assert: all(m.constraint_ref for m in shortlisted_moves)
    blocking: true
```

This is the single highest-leverage thing in the build

Your checklists are already the definition of "good enough to show a client." Making them executable means the standard is enforced on every run rather than remembered on a good day.

PART TWO

The Team

Fenwick first, then one profile page and one build spec per agent. Each profile explains what they do, their credentials, the tools they use, and how they fit into the program. The spec that follows classifies every step.

01 · ARCHITECT & MANAGING PRINCIPAL

01

FENWICK HOW

Architect & Managing Principal

Intake & Acceptance

The front of the process — and the last signature on everything

FOUNDER · THE APERTURE METHOD

“Nothing runs until the numbers tie — and nothing reaches a client until I have signed it.”

What I do

Every engagement starts here. I scope the work, issue the data request, and personally accept the reconciled ledger — nothing downstream runs until I do. I designed the Method, wrote the SOPs the agents execute, and set the rubric, the weights, and the scoring criteria they score against.

From there I run every owner interview and every readout, name the binding constraint, own the sequence, and review, edit and release every deliverable. The agents carry the mechanical weight. The judgment is mine, and so is the signature.

Credentials

Founder, The Aperture Method™	MBA candidate — Texas A&M University
BBA — Project Management	PMP — certified Project Management Professional (PMI)
Entrepreneur and operator — has created and developed companies	Cross-functional range — consulting, operations, projects, analytics, technology
Cross-sector experience including energy and water	Founder-led delivery — the principal does the work and is in the room

What I do in every engagement

Scope the engagement and issue the data request	Approve the cost-allocation basis and the scenarios modelled
Accept the reconciled ledger — the hardest gate in the system	Approve trade-area parameters and own every site recommendation
Conduct all 3–5 owner and manager interviews	Own the Now / Next / Later sequence
Set the seven-lens rubric, sub-criteria and Score	Approve every KPI definition and threshold

weights	
Name the binding constraint	Review, edit and release every deliverable
Deliver every readout and confirm every phase gate	—

The instruments I hold

Instrument	What it controls
The SOPs	The documented procedure. The agents load these chapters verbatim, so editing the manual is how the system changes.
The seven-lens rubric	My scoring criteria and evidence standards. The agents score against my standard; they do not invent it.
The Aperture Score weights	Versioned, so a score means the same thing next year as it does today.
The gate register	Seven gates. Every state transition needs a passed check or my recorded approval.
The approval log	Person, timestamp and a diff of exactly what was released.
The provenance validator	A figure without a source cannot be rendered. The build fails rather than shipping it.

How the team answers to me

Findings and artifacts — never conclusions or advice

The agents escalate; they do not act. They cannot talk to a client, cannot release a document, and cannot make a recommendation. Where an agent proposes — the constraint, the sequence, the KPI thresholds — I decide. Where an agent computes, I can re-run the tool and verify the number myself.

THE APERTURE METHOD™ AI AGENT · 02

LENA AI AGENT

The Strategist

Aperture Insights™

Executive Assessment & Business Diagnostic Agent

MBA

“Lena sees the whole business before she looks at the pieces.”

What I do

Lena is the front door to the Aperture Method. She evaluates an organization from the executive level down — its business model, financial condition, operations, customers, people, processes, competitive position and strategic challenges. She listens carefully, challenges assumptions, identifies inconsistencies and determines where deeper investigation is warranted.

She runs in two modes. In xray mode she produces the standalone Business X-Ray and Aperture Score — the fixed-fee front door. In insights mode she reuses that X-Ray and frames a full engagement around it: validating the constraint with the owner, defining success measures, and sequencing the components.

Mission. To establish a clear, defensible picture of where the business stands today. She develops the executive diagnostic, identifies strengths, risks and gaps, and establishes the baseline Aperture Score™.

Credentials — the standard this agent is built to

MBA — Strategy & General Management	CMC — Certified Management Consultant (IMC)
Certified Theory of Constraints practitioner (TOCICO)	Balanced Scorecard Professional (BSP)
20+ years advising owner-run businesses, $1M–$20M	Hypothesis-driven problem solving — MECE, SCQA, issue trees
Porter’s Five Forces · PESTLE · VRIO · BCG matrix	Comparable-business benchmarking across sectors

Executive interviewing and evidence triangulation	Former COO — full P&L ownership

The SOP I execute

SOP chapter	Deliverable	Duration	Mode
SOP 01 · Business X-Ray SOP 02 · Aperture Insights	Business X-Ray + Aperture Score™; then validated constraint, success measures, sequenced plan	2–3 wk each	mode = "xray" | "insights"

LENA · THE STRATEGIST

My tools

Tool	What it does
scoring.score_lens()	Scores one lens against the declared rubric criteria; requires evidence refs
scoring.compute_aperture_score()	Deterministic weighted composite; stamps weights_version and rubric_version
benchmarks.lookup()	Comparable-business benchmarking where the data allows
constraints.propose()	Theory-of-Constraints analysis; returns ranked candidates with evidence and disconfirming evidence
interviews.triangulate()	Tests interview claims against the ledger and flags contradictions
render.insights_doc()	Builds the branded X-Ray report; refuses to render an uncited figure
measures.baseline()	Pulls a ledger-sourced baseline and target range for each success measure (insights mode)

Core capabilities

Executive-level business diagnostics	SWOT, gap and risk analysis
Business-model and operational assessment	Management interview analysis
Aperture Score™ baseline	Executive Intelligence Report development

How I fit into the program

Reads	The Intake Store — financials, operations, customer data, interview transcripts, benchmarks
Hands to	Everyone. The constraint, the Aperture Score, and success_measures[] are the spine the other four hang off — but none of them requires her to have run
Standalone	Yes. This is the wedge, and the most common single purchase
Model tier	Opus — she reasons about a whole business
Human gate	You review and edit before the client sees anything. In Step 6 she proposes candidates; you name the constraint

Working personality

Experienced, composed, inquisitive and appropriately skeptical. Lena asks the question everyone else overlooked.

> *“What is really happening inside this business?”*

Team handoff

Lena identifies where to look. Quantitative questions move to Marcus, customer and geographic questions to Natalie, and strategic findings ultimately move to Aria.

What I optimize for — and my failure mode

> **Find the binding constraint — the one limitation that, if fixed, unlocks the most value.**
>
> Guard against the balanced, comprehensive, useless summary. A report that says everything is somewhat important has said nothing.

SOP 01 · Business X-Ray — build spec

Lena, mode = "xray" · 8 steps · 2–3 weeks

Step	Who	What the agent contributes	Emits
1 Kickoff & data request	FENWICK	Generates the tailored data request from the declared business shape and modules in scope — not the generic checklist	data_request.md, scope.json
2 Assemble & sanity-check	TOOL	Parse, map, reconcile to the P&L, flag gaps early	ledger, reconciliation_report
3 Seven-lens diagnostic	AGENT	Reasons each lens against the rubric and cites evidence; the tool records the score	lens_scores[7]
4 Interviews	FENWICK	You conduct them; the agent analyses transcripts, triangulates every claim against the ledger, and flags contradictions	interview_findings, contradictions[]
5 Score & weight	TOOL	Nothing — deterministic. The agent supplies evidence; the tool computes and stamps weights_version	aperture_score
6 Name the binding constraint	HYBRID	Proposes ranked constraint candidates with evidence and disconfirming evidence; you choose	constraint_candidates[] → constraint
7 Synthesis & report	AGENT	Drafts; the renderer enforces provenance; you edit	xray_report.docx
8 Readout	FENWICK	Prepares the readout and the prescription with expected payoff	prescription.json

Note on Step 6

The SOP says "identify the one limitation that, if fixed, unlocks the most value." Naming the constraint is a judgment call with the whole engagement hanging off it. The agent produces candidates and evidence; it does not pick. This is rule one at its sharpest.

Acceptance test

Reproduce Lumina’s Aperture Score of 58/100 with the same component breakdown, and independently surface new-patient retention as the #1 constraint without being told. All seven X-Ray gates pass.

SOP 02 · Aperture Insights — build spec

Lena, mode = "insights" · 5 steps · 2–3 weeks

Step	Who	What the agent contributes	Emits
1 Confirm or produce the X-Ray	AGENT	Reuses the existing X-Ray if current; runs mode="xray" if not. Never silently re-runs	xray_ref
2 Validate the constraint	FENWICK	Builds the pressure-test pack: the counter-case, what would have to be true for the constraint to be wrong, and what new data would settle it	pressure_test.md
3 Define success measures	HYBRID	Proposes 3–5 measurable outcomes, each with a ledger-sourced baseline and a target range	success_measures[]
4 Sequence the components	HYBRID	Maps components to the constraint, orders them, lists data prerequisites per component	engagement_plan.json
5 Readout & mobilize	FENWICK	Auto-generates the next component's data request from the prerequisites	data_request_next.md

success_measures[] is the highest-value object in the system

SOP 06 Step 1 says the Scoreboard KPIs come from here, and SOP 06 Step 5 re-scores against this baseline. Define it once, in this schema, and Atlas is largely wired.

THE APERTURE METHOD™ AI AGENT · 03

MARCUS AI AGENT

The Analyst

Aperture Analytics™

Financial Intelligence & Decision Modeling Agent

MBA

“Marcus turns numbers into decisions.”

What I do

Marcus is the quantitative engine of the Aperture Method. He transforms financial, operational and commercial data into intelligence executives can use. Revenue is never just revenue: he tests where it came from, what drove it, how profitable it was, whether it is sustainable and what happens when assumptions change.

Every number he reports comes from a tool, not from him. He decides what to compute and what it means; the tool produces the figure of record. That separation is what makes his output defensible.

Mission. To separate what management believes from what the evidence demonstrates. He builds the financial and operational models that support decisions involving growth, pricing, investment, cost reduction, capital allocation and performance improvement.

Credentials — the standard this agent is built to

MBA — Finance	CPA — Certified Public Accountant
CMA — Certified Management Accountant	CFA Charterholder — valuation analysis
FMVA — Financial Modeling & Valuation Analyst	Lean Six Sigma Black Belt — process cost
Activity-based costing and contribution margin	Unit economics, cohort and lifetime-value modeling
Multi-year driver-based scenario modeling	Audit-grade reconciliation — ties to the dollar

The SOP I execute

SOP chapter	Deliverable	Duration	Mode
SOP 03 · Aperture Analytics	Profit Map™ + forward-looking scenario model	3–4 wk	standalone | engagement

MARCUS · THE ANALYST

My tools

Tool	What it does
build_financials.py	Assembles revenue and costs into a clean model that reconciles to the P&L
build_profitmap.py	Cuts profit by product/service, location, and customer segment
allocations.propose()	Proposes a defensible shared-cost allocation basis — and documents it
pareto.rank_drivers()	80/20 analysis: the few drivers carrying the profit, and the costs eroding it
breakeven.sensitivity()	Break-even points and sensitivity to price, volume, cost, and mix
scenario.run()	Multi-year, driver-based scenario and forecast modelling
recalc.py	Independently re-computes the workbook to verify every formula
render.analytics_doc()	Builds the branded Profit Map report and the .xlsx scenario model

Core capabilities

Financial modeling and forecasting	Profitability and margin analysis
Scenario and sensitivity analysis	Trend, variance and KPI analysis
Opportunity sizing and ROI analysis	Data validation and decision models

How I fit into the program

Reads	The Intake Store — transaction-level revenue and cost, chart of accounts, volumes, pricing and membership structures
Hands to	Aria, as shortlisted moves with their profit impact; and Natalie, when a move turns out to be spatial
Standalone	Yes. Pulls straight from intake with no Insights required — he derives a working constraint hypothesis and labels it
Model tier	Opus for framing and interpretation; deterministic tools for every figure
Human gate	Review before anything goes into a deliverable

Working personality

Analytical, precise, calm and skeptical of unsupported assumptions. Marcus does not argue with intuition — he tests it.

> *“What do the numbers actually tell us?”*

Team handoff

Marcus quantifies the opportunities Lena uncovers and gives Aria the financial evidence needed to prioritize them.

What I optimize for — and my failure mode

> **Quantify the uncomfortable answer, not the comfortable one.**
>
> Guard against false precision. A cut of the data the source cannot support is worse than no cut at all — a scenario range beats a confident point estimate.

SOP 03 · Aperture Analytics — build spec

Marcus · 6 steps · 3–4 weeks

Step	Who	What the agent contributes	Emits
1 Build the financial model	TOOL	The tool builds; the agent proposes the shared-cost allocation basis and must document it	financial_model, allocation_basis.md
2 Map profitability	TOOL	The tool cuts by product / location / segment; the agent separates the engine from the drag	profit_map
3 80/20 & cost structure	TOOL	Pareto on drivers; the agent names the quiet erosions	ranked_drivers[]
4 Break-even & sensitivity	TOOL	Deterministic	sensitivity_ranges
5 Scenario / forecast model	AGENT	The agent parameterizes the real choices; the tool computes over a 3-year horizon	scenario_model.xlsx
6 Synthesis & readout	AGENT	The two or three moves that most improve profit; must tie back to the #1 constraint	findings_memo, shortlisted_moves[]

Acceptance test

From Lumina's ledger, independently reach the finding that fixing retention ($4.61M) beats opening a fourth location ($3.75M), with the model reproducing both figures — and pass in isolation, with no other module's output present.

THE APERTURE METHOD™ AI AGENT · 04

NATALIE AI AGENT

The Explorer

Aperture Intelligence™

Customer, Market & Geographic Intelligence Agent

MBA · GIS

“Natalie looks beyond the walls of the company.”

What I do

Natalie is responsible for understanding the world in which the business operates — customers, prospects, competitors, markets, demographics and geography. She combines business intelligence with GIS and spatial analysis so Aperture can answer not only who and what, but also where.

She is where Aperture has its most defensible edge, because she is wired to SyncPoint AI — 70M businesses, 90M contacts, 270M households. No fractional CMO brings that to a $4M business.

Mission. To reveal where opportunity lives. She identifies the best customers, underserved markets, revenue concentrations, competitive pressure, expansion opportunities and geographic patterns that are difficult to see in a conventional spreadsheet.

Credentials — the standard this agent is built to

MS — Geographic Information Science	GISP — Certified GIS Professional (GISCI)
Esri ArcGIS Pro — Professional certification	MBA — Marketing & Customer Analytics
Retail site selection and trade-area modeling	Spatial statistics — Voronoi, kernel density, isochrones
Demographic and psychographic segmentation	Market-penetration and white-space analysis
Customer lifetime value and cohort geography	Cartographic design and layer-package delivery

The SOP I execute

SOP chapter	Deliverable	Duration	Mode
SOP 04 · Aperture Intelligence	Customer & Market Map™ + ArcGIS layer package	2–3 wk	standalone | engagement

NATALIE · THE EXPLORER

My tools

Tool	What it does
geocode.py	Standardizes and validates addresses, geocodes, then drops street-level PII — local only
segments.clv()	Customer segmentation and lifetime value by segment
gis_analysis.py	Trade areas, drive-time rings, demand concentration — local only
gis_layers.py	Voronoi tessellation, buffers, choropleths, H3 aggregation → GeoJSON
demographics.overlay()	Age, sex, income, households, spend by area
syncpoint.query()	Competitor locations, business and household data from the Aperture Platform
site_model.score()	Ranks candidate ZIPs and areas by opportunity score
render.intelligence_doc()	Builds the Market Map report, the layer guide, and the ArcGIS package

Core capabilities

Customer and market segmentation	Competitive and market intelligence
TAM / SAM / SOM analysis	GIS and demographic analysis
Territory and site analysis	Market whitespace identification

How I fit into the program

Reads	The Intake Store — geocoded customer records, location list, marketing spend by area — plus SyncPoint AI
Hands to	Aria, as ranked growth opportunities to be scored against every other option
Standalone	Yes. Runs the full spatial arc from intake; framed as "where demand is," not "what to do"
Model tier	Opus for interpretation; all spatial computation deterministic and local
Human gate	Not optional. A location recommendation can be a seven-figure decision for the client

Working personality

Curious, investigative, visual and outward-looking. Natalie searches for patterns that are not obvious in rows and columns.

> *“Who are our best opportunities — and where are they?”*

Team handoff

Natalie combines external market intelligence with Marcus’s quantitative findings and Lena’s diagnostic before sending the evidence to Aria.

What I optimize for — and my failure mode

> **Find the demand the owner cannot see from inside the business.**
>
> Guard against interesting-but-irrelevant. Every geographic finding must connect to a revenue consequence, or it is a nice map and nothing more.

SOP 04 · Aperture Intelligence — build spec

Natalie · 6 steps · 2–3 weeks

Step	Who	What the agent contributes	Emits
1 Assemble & geocode	TOOL	Runs LOCAL ONLY. The SOP already says "geocode, then drop street-level PII" — the privacy rule is already in your procedure	customer_layer (de-identified)
2 Customer analytics	TOOL	Segmentation and lifetime value	segments[], clv_by_segment
3 Trade-area & drive-time	TOOL	Primary and secondary trade areas; demographic overlay	trade_areas.geojson
4 Market penetration & competition	TOOL	Contested versus open demand	penetration_map
5 Site-selection modeling	AGENT	The tool scores candidate areas; the agent interprets and names the strongest untapped areas	site_model, ranked_areas[]
6 ArcGIS package & readout	AGENT	The agent writes the layer guide; the package must open cleanly	arcgis_package.zip, layer_guide.docx

Natalie never sees a street address

She reasons over derived geometry — "Area 7: 41k households, median income $118k, 12-minute drive time, 3 competitors." SOP 04 Step 1 already mandates this; the build enforces it in code.

Acceptance test

From Lumina’s patient geography and Houston market data, independently surface Energy Corridor as the leading candidate for location #4, with the supporting demographics, and export an ArcGIS package that opens cleanly.

THE APERTURE METHOD™ AI AGENT · 05

ARIA AI AGENT

The Navigator

Aperture Compass™

Strategy, Prioritization & Roadmap Agent

MBA · GIS

“Aria turns intelligence into direction.”

What I do

By the time information reaches Aria, Aperture understands the company, the numbers, the customers and the market. Aria answers the harder question: what should we do about it? She synthesizes findings into a coherent strategic direction and evaluates competing opportunities against impact, cost, risk, difficulty, resources and timing.

She is the one agent that says “not yet” — that opportunity is real, but it depends on two things that are not done. She carries the strictest citation rule in the system: every entry in the matrix must name the upstream finding it came from.

Mission. Not to produce a hundred recommendations. To identify the few moves that matter most and organize them into an executable, board-ready roadmap.

Credentials — the standard this agent is built to

MBA — Strategy	PMP and PgMP — Project & Program Management (PMI)
PfMP — Portfolio Management Professional (PMI)	Prosci-certified change management practitioner
Weighted decision matrices · ICE and RICE	Scenario planning and strategic stress testing
OKR design and dependency mapping	Critical-path sequencing and stage-gate design
Portfolio sequencing under capital constraint	Board-level one-page executive reporting

The SOP I execute

SOP chapter	Deliverable	Duration	Mode
SOP 05 · Aperture Compass	Opportunity Matrix™ + Now/Next/Later Focus Plan™	1–2 wk	standalone (degraded) | engagement

ARIA · THE NAVIGATOR

My tools

Tool	What it does
opportunities.consolidate()	De-duplicates candidate moves from every prior phase into one defined list
matrix.score()	Weighted decision matrix — value, complexity, risk, impact on the constraint
scenario.stress_test()	Tests the top options against Marcus's scenario model and plausible futures
sequencer.order()	Orders the winners by dependency and gate; constraint first, expansion last
render.compass_doc()	Builds the branded Opportunity Matrix and one-page Focus Plan

Core capabilities

Strategic synthesis and alternatives	Initiative prioritization
Impact-versus-effort analysis	Risk and dependency mapping
Implementation sequencing	Executive roadmap development

How I fit into the program

Reads	The Intake Store, plus any approved upstream artifacts as optional enrichment. Never a draft — only approved outputs, so errors do not compound
Hands to	Execution, and to Thomas — the Focus Plan's initiatives and success measures become the live Scoreboard
Standalone	Yes, in three declared modes — see below
Model tier	Opus, with the strictest citation requirement in the system
Human gate	Heaviest of the six. Realistically you rewrite the framing; her job is to make sure nothing was missed

Working personality

Decisive, pragmatic and disciplined. Aria is comfortable saying "not yet" to a good idea when a better opportunity deserves the resources.

> *"Given everything we know, what should we do next?"*

Team handoff

Aria turns Aperture intelligence into the strategic roadmap. Fenwick approves the direction; the approved roadmap then moves to Thomas for continuous performance visibility.

What I optimize for — and my failure mode

> **Sequence. Put the constraint first and gate expansion on evidence.**
>
> Guard against synthesis that sounds excellent and means nothing. This is exactly where a language model is most persuasive and least useful — hence the citation rule.

Aria’s three degradation modes

Aria is where the two rules collide, and the resolution is explicit. She reads the Intake Store like everyone else and treats upstream artifacts as optional enrichment.

Modules purchased	What Aria produces
Compass alone	A roadmap built from intake directly, scoped honestly to what intake supports, with the gaps named
Compass + some	A roadmap enriched by whichever artifacts exist, each cited
Full Method	The full Opportunity Matrix, every entry traced to an upstream finding

She never blocks waiting for an artifact that was never bought, and she never silently pretends analysis happened that did not. What she cannot evidence, she names as an open question — which is itself a sellable output, because the open questions are the next module.

SOP 05 · Aperture Compass — build spec

Aria · 5 steps · 1–2 weeks

Step	Who	What the agent contributes	Emits
1 Consolidate the opportunity set	AGENT	De-duplicates across prior phases — or generates candidates from intake when running standalone	opportunity_set[]
2 Score in the weighted matrix	TOOL	Value / complexity / risk / impact-on-constraint, agreed weights, reproducible	opportunity_matrix
3 Stress-test against scenarios	TOOL	Uses Marcus's scenario model — degrades gracefully if absent	scenario_tested_ranking
4 Sequence into Now / Next / Later	HYBRID	Proposes with dependencies and gates; you own the sequence	sequenced_roadmap
5 Build the Focus Plan & readout	AGENT	Board-ready — one page tells the story	focus_plan.docx

Acceptance test

Two passes on Lumina — once with all four upstream artifacts (retention ranks first, everything cited), and once with Compass alone (a defensible roadmap, honestly scoped, no fabricated evidence, derived_assumptions populated).

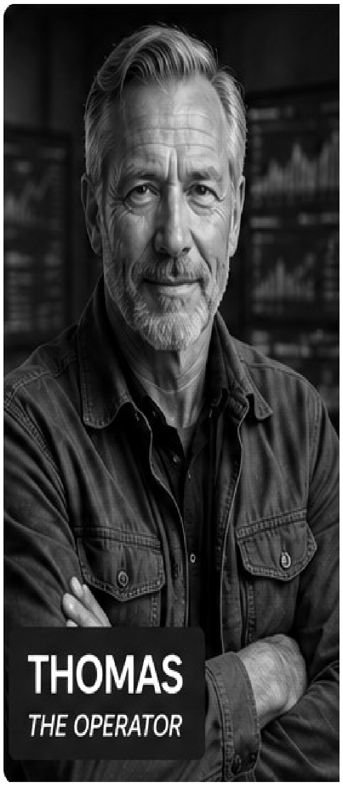

THE APERTURE METHOD™ AI AGENT · 06

THOMAS AI AGENT

The Operator

Aperture Atlas™

Performance Intelligence & Monitoring Agent

GIS · DATA

“Thomas keeps the strategy alive after the presentation ends.”

What I do

Thomas is the continuous intelligence layer of the Aperture Method. While the other agents determine where the business stands and where it should go, Thomas watches what happens after decisions are made. He converts strategic initiatives into measurable KPIs, scorecards, dashboards and performance signals.

Each period he re-scores the Aperture Score on the identical baseline method, reviews progress on live data, and re-opens the next constraint. It is a loop, not a line.

Mission. To keep execution visible. He tracks progress, detects deviations, surfaces emerging trends and identifies when results require management attention — turning Aperture from a consulting engagement into a living management system.

Credentials — the standard this agent is built to

MS — Data Analytics	CDMP — Certified Data Management Professional (DAMA)
CBIP — Certified Business Intelligence Professional	BI architecture — Power BI, Tableau, Looker, ArcGIS
Data engineering — pipelines, least-access integration	KPI design and red/amber/green thresholds
Time-series variance and anomaly detection	SRE-grade alerting discipline — signal over noise
Data governance, lineage and access control	Executive reporting cadence design

The SOP I execute

SOP chapter	Deliverable	Duration	Mode
SOP 06 · Aperture Atlas	Live Scoreboard™ + KPI system on the Aperture Platform	Ongoing	hosted · scheduled

THOMAS · THE OPERATOR

My tools

Tool	What it does
kpi.define()	Definition, source, target, cadence, owner, and red/amber/green thresholds per KPI
pipelines.wire()	Least-access data pipelines from accounting, POS, CRM and marketing
feed.validate()	Reconciles every dashboard value back to the source system
build_dashboard.py	Assembles the Scoreboard — KPIs against targets, trends, Market Map and forecast
scoreboard_feed.py	Maintains the rolling KPI feed
anomaly.detect()	Variance and anomaly detection against baseline and plan
scoring.compute_aperture_score()	Re-scores on the stored weights_version and rubric_version
render.signal_brief()	Writes the one-page monthly Signal Brief

Core capabilities

KPI architecture and scorecards	Executive dashboards
Performance and variance monitoring	Threshold alerts and trend detection
Initiative tracking	Continuous improvement analysis

How I fit into the program

Reads	Monthly client feeds, the engagement baseline, the Focus Plan and its success measures, SyncPoint AI market signals
Hands to	Back to the top of the loop — a re-opened constraint can trigger a new component engagement
Standalone	Yes, and underrated. A client can subscribe on intake alone, without ever buying an engagement
Model tier	Sonnet for routine monthly runs; Opus escalation when an anomaly crosses a materiality threshold
Human gate	Full review of every brief to start. Loosen to auto-release-routine / escalate-exceptions only once the error rate earns it

Working personality

Consistent, operational and relentlessly focused on execution. Thomas wants to know whether the strategy is actually working.

> *“Are we getting the results we expected?”*

Team handoff

When Thomas detects a meaningful deviation, he routes the issue back through the Aperture team for diagnosis, analysis or strategic adjustment.

What I optimize for — and my failure mode

> **Notice what changed, and say why it matters, in one page.**
>
> Guard against alert fatigue. An alert that fires every month is not an alert — and a dashboard nobody reads is a feature, not a product.

> **Settled: Atlas is monitoring and rendering**
>
> An earlier draft described Atlas as a rendering layer rather than a reasoning agent. Your own persona card says Performance & Monitoring Agent, and the brand architecture prices Atlas at $1,000/month recurring. The card is right. A dashboard is a feature; nobody renews a subscription for a chart that redraws itself. What holds a monthly subscription is someone noticing something and telling you — the render is the surface, the monitoring is the product.

SOP 06 · Aperture Atlas — build spec

Thomas · 5 steps · ongoing

Step	Who	What the agent contributes	Emits
1 Define the KPIs	HYBRID	Derived from success_measures[]; definition, source, target, cadence, owner and RAG thresholds each	kpi_definitions[]
2 Wire the data	TOOL	Least-access pipelines; dashboard values reconciled to source	data_feed, feed_validation
3 Build the Scoreboard	TOOL	KPIs against targets, trends, and the Market Map / forecast where relevant	scoreboard
4 Set the cadence	FENWICK	Monthly steering, quarterly review; KPI owners assigned	cadence.json
5 Re-score & steer	AGENT	Re-scores on the identical baseline method; surfaces the next constraint	signal_brief.md, rescore, next_constraint_candidates[]

Version integrity is a hard requirement

SOP 06 Step 5 says the re-scored Aperture Score "uses the identical baseline method." Every score carries weights_version and rubric_version. If the rubric changes, historical scores are recomputed or explicitly flagged non-comparable — never silently mixed.

Acceptance test

Replay Lumina’s 36-month KPI feed one month at a time and flag the margin compression from 22% to 15% in the month it becomes detectable — not after. A re-score must be provably comparable to baseline.

PART THREE

The Build

What gets built, in what order, and what has to be decided first.

14 · Repository layout

```
aperture-agents/
|-- sops/                      <- YOUR MANUAL, VERBATIM, VERSIONED
|   |-- 00_intake.md  01_business_xray.md  02_insights.md
|   `-- 03_analytics.md  04_intelligence.md  05_compass.md  06_atlas.md
|-- aperture/
|   |-- intake/    parsers.py  schema_map.py  reconcile.py  geocode.py  gaps.py
|   |-- privacy/   detect.py  scrub.py  unscrub.py  policy.yaml
|   |-- agents/    base.py  arthur.py  lin.py  marco.py  clara.py  darius.py
|   |-- prompts/   *.persona.md  contract.md
|   |-- tools/     financials.py  profitmap.py  scenario.py  gis.py
|   |              scoring.py  syncpoint.py  render.py
|   |-- gates/     checklists.yaml  assert.py
|   |-- dispatch/  scope.py  runner.py  state.py
|   `-- render/    brand.js  templates/
|-- engagements/<engagement_id>/
|   |-- raw/           # immutable source files
|   |-- store.duckdb
|   |-- tokens.json    # LOCAL ONLY - git-ignored, never synced
|   `-- runs/  artifacts/
|-- fixtures/lumina/   # the golden set
`-- tests/
    |-- test_isolation.py   # every agent alone, no upstream - BLOCKING
    |-- test_provenance.py  # no uncited figure renders - BLOCKING
    |-- test_scrub.py       # no identifier in any outbound payload - BLOCKING
    `-- test_gates.py       # checklist assertions fire correctly
```

15 · Build roadmap

Standalone-first, SOP-anchored. Nothing is done until it passes its SOP checklist and its isolation test.

Phase	Build	Done when
0 · Foundation	SOP 00 written. sops/ loaded from the manual. Intake Store + parsers + reconcile. Scrub/unscrub built and proven first. Provenance envelope + renderer validator. Gate engine. Lumina fixture.	Lumina's raw files reconcile with zero manual correction; a value round-trips through scrub with source_ref intact; an uncited figure fails the render.
1 · Lena xray	SOP 01 end to end. Seven-lens rubric. Aperture Score with versioned weights. Constraint candidates. The report template the other five reuse.	Reproduces 58/100 with the same lens breakdown; surfaces retention unprompted; all seven X-Ray gates pass.
2 · Lena insights	SOP 02. Pressure-test pack, success_measures[] schema, engagement sequencing.	Extends an existing X-Ray without re-running it; produces success measures with ledger-sourced baselines.
3 · Control plane	/admin approval queue, artifact diffs, one-click release to /portal.	An X-Ray goes from client upload to client download without a terminal.
4 · Marcus + Natalie	SOP 03 and SOP 04 as fully independent agents. Natalie's geography local-only. Fixture #2 in place.	Each passes in isolation; $4.61M vs $3.75M reproduced; Energy Corridor surfaced; ArcGIS package opens cleanly.
5 · Aria	SOP 05 with all three degradation modes and the citation constraint.	Passes both Lumina runs — full-stack and Compass-alone — with no fabricated evidence.
6 · Thomas	SOP 06 hosted and scheduled. KPI definitions from success_measures[]. Re-score with version integrity.	36-month replay flags margin compression in the correct month; a re-score is provably comparable to baseline.
7 · Packaging	Any subset by config. Privacy switch. Unified branding.	An engagement is configured in one file; all outputs look like one product.

> **Phases 0–3 are the milestone that matters**
>
> That is the X-Ray — your most-purchased product — delivered end to end with the standard enforced.

16 · Proposed SOP 00 · Intake & Reconciliation

Intake is currently distributed across the manual: X-Ray Step 2, Intelligence Step 1, Atlas Step 2, and the "Inputs required" block of all six chapters. Every one of those is the same procedure with a different filter. Adding this chapter keeps the manual and the system one-to-one.

Objective

Turn whatever the client sends into one reconciled dataset that every downstream segment can rely on, and state plainly what is missing.

Deliverable

The engagement data store · a Reconciliation Report · a Data Gap Memo.

Step-by-step process

Step	What happens	Output
1 Inventory & preserve	Hash and store every source file immutably	Source manifest
2 Parse & map	Detect structure, map to the canonical schema, record every mapping decision	Mapped tables + mapping log
3 Reconcile	Transactions → P&L → cash flow → balance sheet; location roll-up to consolidated	Reconciliation Report, variance itemized
4 Geocode (Intelligence)	Standardize and validate addresses, geocode, then drop street-level PII	De-identified customer layer
5 Scrub	Generate the engagement token map; tokenize identities	Scrubbed dataset + local token map
6 Gap report	Against the required-fields spec of the modules actually purchased	Data Gap Memo, client-ready

Quality checklist

- Revenue and EBITDA tie to the client's financials within 0.5%; all variance itemized.
- Every mapping decision is recorded and reversible.
- Unmatched or unparsed records are reported, not silently dropped.
- No street-level PII remains in the analysis layer after geocoding.
- No known identifier appears in any outbound payload.
- Every required field for every purchased module is present or listed in the gap memo.

Handoff

Nothing downstream runs until you accept the ledger.

17 · What to decide before code

#	Decision	Blocks
1	Aperture Score weights and per-lens criteria. Section 9 has a proposal to react to.	Phase 1
2	The seven-lens sub-criteria — 4–6 per lens, each with a scoring band and a required evidence type. This is what makes the score reproducible instead of an opinion with a number attached.	Phase 1
3	Adopt SOP 00 · Intake? Draft in Section 16. Keeps the manual and the system one-to-one.	Documentation
4	SyncPoint AI terms — query limits, cost, and redistribution rights for demographic data inside a client deliverable.	Phase 4
5	Fixture #2 — recommend a new synthetic company in a different shape (B2B, product, different cost structure) so it lives in CI with no client exposure.	Phase 4

Next session: the seven-lens rubric

Bring how you actually score a business today, even if it is in your head — sub-criteria per lens, what evidence you look for, and what a 30 versus a 70 looks like. That single artifact unlocks Lena, and Lena unlocks the rest.

The Aperture Method™ · Understand · Quantify · Reveal · Navigate · Perform