

The Seven-Lens Rubric & the Aperture Score™

Version 1.2.0 · Ratified

The seven lenses are taken verbatim from SOP 01 — Business X-Ray. The sub-criteria, anchor descriptions and weights below were reasoned by construction from the target segment — owner-run, \$1M–\$20M, place-based or data-rich — and were **ratified by Fenwick How on 2026-08-10**. They are the only weights in use: the website mirrors them, and its build fails if the two disagree. Any change is a rubric version change, and every score stamps the version it used, so history stays comparable.

How scoring works

Each sub-criterion is scored **0–100** against its anchor text, or marked **I** for insufficient evidence. A lens score is the mean of its scored sub-criteria. The Aperture Score is the weighted mean of the seven lens scores.

Band	Meaning
0-20	Critical — actively threatens the business
21-40	Weak — a material drag; likely constraint territory
41-60	Adequate — functioning, not an advantage
61-80	Strong — a genuine asset
81-100	Distinctive — a durable competitive advantage

Insufficient is not zero. A sub-criterion marked **I** is excluded from its lens mean. Scoring it zero would assert a critical weakness the evidence does not support.

Weight is never redistributed. An insufficient lens lowers coverage rather than quietly inflating the remaining lenses. The gap stays visible, because the gap is a finding.

Below 70% coverage, no Score is published. The deliverable in that case is a data request, not a number.

Claims are not evidence. Intake answers and interview statements enter the ledger as claims. They support a score only once the ledger corroborates them.

The weights

The column on the right is for your own numbers. These are the single most consequential judgment in the rubric — they decide what the Score is actually measuring.

Lens	Proposed	Reasoning	Yours
Finance	20%	Highest weight. Cash and unit economics are what actually kill businesses in this segment, and it is the most reliably evidenced lens — the P&L exists whether or not anything else does.	
Operations	18%	In place-based businesses the constraint most often binds physically: capacity, throughput, a bottleneck someone can point at.	
Customers & Retention	18%	Retention is the highest-leverage and most commonly mismeasured area in owner-run businesses. Owners routinely believe their repeat rate is far higher than it is.	
Market & Competition	14%	Sets the ceiling on everything else. A well-run business in a shrinking market has a strategy problem no operational fix reaches.	
Leadership & Team	12%	Key-person risk is existential here and is the single most common reason an owner-run business cannot be sold. Weighted above Processes for that reason.	
Processes	10%	Lower weight deliberately. Weak process is usually a symptom of Operations or Leadership rather than a root constraint.	
Technology & Data	8%	Lowest weight, deliberately. Technology is almost never the binding constraint in this segment — it is an enabler. Weighting it higher produces 'buy software' recommendations, which is the failure mode of this entire industry.	
Total	100%	The engine refuses to run unless these sum to 100%.	

The seven lenses

Intake shows which questionnaire answers feed each sub-criterion — SHR shared, INS Insights, ANL Analytics, INT Intelligence, CMP Compass, ATL Atlas. Every sub-criterion has at least one source, so nothing is scored on impression. Note that several lenses draw on later-phase answers: on a standalone Business X-Ray those will not exist, so score what you have and let coverage tell the honest story.

Finance 20%

fin-01 — Margin visibility

Whether the business knows where it makes and loses money.

Intake: ins_fin1, ins_fin6, an_num2, an_num3 · Evidence: measured, derived

20	Profit is known only at whole-business level. The owner is guessing about product or location profitability.
50	Gross margin known by major line; overhead unallocated. The owner has a view but low confidence in it.
80	Contribution and fully-absorbed profit known by product/service and location, on a documented allocation basis.

fin-02 — Cash position and cycle

Whether cash is predictable and adequate through the year.

Intake: ins_fin2, ins_fin4 · Evidence: measured

20	Cash is unpredictably tight. The owner discovers shortfalls rather than forecasting them.
50	Seasonal tight points are known from experience. No rolling forecast exists.
80	A rolling cash forecast exists; the cash conversion cycle is measured and actively managed.

fin-03 — Unit economics

Whether the economics of one unit, job or customer are understood and positive.

Intake: ins_fin5, an_price1, an_price3, an_price4, an_dec3 · Evidence: measured, derived

20	Unit economics unknown. Pricing set by feel or by reference to a competitor.
50	Rough unit margin known for the main lines, untested against fully loaded cost.
80	Unit economics known and tested, including cost to serve and acquisition cost.

fin-04 — Cost trend control

Whether cost increases are understood and acted on.

Intake: ins_fin3 · Evidence: measured, derived

20	Costs have grown materially and the owner cannot explain why.
50	The main areas of cost growth are identified; no systematic review.
80	Cost trends reviewed on a cadence with named owners; drivers understood and challenged.

fin-05 — Books reliability

Whether the accounts can be relied on for decisions.

Intake: an_num4, data_missing, data_basis · Evidence: measured, claimed

20	Books closed late or not at all. The owner does not trust them — and is right not to.
50	Closed monthly with some lag. Broadly trusted, with known issues that get worked around.
80	Closed promptly, reconciled, and actually used by the owner to make decisions.

Operations 18%

ops-01 — Capacity headroom

Whether the business can absorb demand growth without a step change in cost.

Intake: ins_ops2, ins_ops3, ins_ops4 · Evidence: measured, derived

- | | |
|----|--|
| 20 | At or over capacity now. Demand is being turned away or service quality is degrading. No headroom in the binding resource. |
| 50 | Roughly 10-25% headroom in the binding resource. Peaks already strain it. |
| 80 | 40%+ headroom, and the next step-change in capacity is identified and costed. |

ops-02 — Bottleneck identification

Whether the constraining step is known, measured and actively managed.

Intake: ins_ops1, ins_ops2 · Evidence: measured, claimed

- | | |
|----|---|
| 20 | Nobody can name where work backs up — or different people name different places. Treat disagreement here as a finding in its own right. |
| 50 | The bottleneck is named consistently but not measured. Management is reactive. |
| 80 | Named, measured, and explicitly protected and scheduled around. |

ops-03 — Throughput reliability

Whether work completes predictably, on time and to standard.

Intake: ins_ops1, ins_ops4, an_vol1 · Evidence: measured, claimed

- | | |
|----|--|
| 20 | Frequent rework, missed dates or expediting. On-time completion is not measured at all. |
| 50 | Generally reliable, but exceptions are absorbed by individual effort rather than by system. Measured informally. |
| 80 | On-time and quality both measured, variance understood, exceptions rare and root-caused. |

ops-04 — Demand and capacity matching

Whether capacity is planned against demand, including seasonality.

Intake: ins_ops3, ins_ops4, an_vol3 · Evidence: measured, derived

- | | |
|----|--|
| 20 | Capacity is fixed and demand swings are absorbed by overtime or lost sales. |
| 50 | Seasonality is known from experience and informally planned for. |
| 80 | Capacity is flexed against a demand forecast; the plan is written and reviewed on a cadence. |

ops-05 — Qualified capacity

Whether the supply of licensed, certified or specifically-trained people is the binding constraint on billable volume, and how long that constraint takes to relieve.

Intake: ins_ops5, ins_ops4 · Evidence: claimed, measured

20	Billable volume is capped by qualified headcount, the cap is not measured, and the lead time to qualify another person is long or unknown.
50	The constraint is understood and roughly measured, but hiring or qualification is reactive — it starts when someone leaves or demand already exceeds supply.
80	Qualified capacity is measured against forecast demand, the qualification pipeline is deliberate, and lead time to add a qualified person is known and planned for.

Customers & Retention 18%

cus-01 — Acquisition understanding

Whether it is known where customers come from and what each channel costs.

Intake: ins_cust1, ins_cust5, int_grow2, int_grow3 · Evidence: measured, derived

20	Sources unknown or purely anecdotal. No spend attribution.
50	Sources ranked by impression. Marketing spend known in total but not by channel.
80	Sources measured, spend attributed by channel, acquisition cost known per channel.

cus-02 — Retention measurement

Whether repeat behaviour is measured rather than assumed.

Intake: ins_cust2 · Evidence: measured

20	Repeat rate unknown. The owner believes customers return but cannot show it. Expect the belief to be optimistic — test it early.
50	Repeat rate estimable from the system but not tracked over time or by cohort.
80	Cohort retention tracked over time; changes are noticed and investigated.

cus-03 — Customer value concentration

Whether value concentration is known and deliberately managed.

Intake: ins_cust3, b_buyers, data_have · Evidence: measured, derived

20	No view of which customers carry the revenue. Concentration risk unassessed.
50	Top customers known by name. No systematic segmentation or lifetime value.
80	Segmented with LTV. Concentration quantified and managed as a deliberate risk position.

cus-04 — Churn causes

Whether the reasons customers leave are known and addressed.

Intake: ins_cust4 · Evidence: measured, claimed

20	Departures go unnoticed. No exit information of any kind.
50	Reasons known anecdotally. Nothing systematic.
80	Churn reasons captured systematically and fed back into the offer or operations.

Market & Competition 14%

mkt-01 — Competitive position

Whether there is a clear, defensible reason customers choose this business.

Intake: ins_mkt1, int_comp3 · Evidence: claimed, measured

20	No articulable difference. Competing on price by default.
50	A real difference exists but is not consistently communicated or defended.
80	A clear, evidenced advantage that customers can name and competitors struggle to copy.

mkt-02 — Competitive awareness

Whether the business knows who it loses to and why.

Intake: ins_mkt1, ins_mkt4, int_comp1, int_comp2 · Evidence: measured, claimed

20	Competitors unknown or unexamined. Losses unexplained.
50	Competitors named. Loss reasons anecdotal.
80	Losses tracked with reasons. Competitor moves monitored and responded to deliberately.

mkt-03 — Demand outlook

Whether the direction of the market is known and reflected in the plan.

Intake: ins_mkt2 · Evidence: measured, derived

20	No view of market direction. Continuity assumed.
50	Directional view held from experience, not evidenced externally.
80	Market trend evidenced from external data, and the plan visibly reflects it.

mkt-04 — Geographic and segment position

Whether the served market, its penetration and its whitespace are understood.

Intake: ins_mkt3, int_geo2, int_geo3, int_grow1 · Evidence: measured, derived

- | | |
|----|---|
| 20 | No view of catchment, penetration or whitespace. |
| 50 | Catchment known by impression. Strong and weak areas sensed but never measured. |
| 80 | Catchment and penetration measured. Whitespace quantified and sized. |

Leadership & Team 12%

ldr-01 — Key-person risk

What happens when the owner or a key person is absent.

Intake: ins_lead1, ins_lead4 · Evidence: claimed, measured

- | | |
|----|---|
| 20 | The business stops or degrades badly within days. The owner cannot take a month off. This is the most common reason an owner-run business cannot be sold. |
| 50 | Runs for weeks, but decisions queue and some functions have no backup. Single-person dependencies are known but not quantified. |
| 80 | Runs a month without the owner. Decision rights are delegated and documented, and the revenue exposed to any single-person dependency is quantified and being deliberately reduced. |

ldr-02 — Role clarity and capability

Whether the right roles are defined and filled with capable people.

Intake: ins_lead2 · Evidence: claimed, measured

- | | |
|----|---|
| 20 | Roles undefined or badly mismatched to the people in them. Critical gaps unfilled and unacknowledged. |
| 50 | Roles broadly defined. One or two critical gaps acknowledged but not being actively closed. |
| 80 | Roles defined against the plan, filled with capable people, gaps identified and actively recruited. |

ldr-03 — Decision-making and execution

Whether decisions get made and executed, or stall.

Intake: d_approve, ins_lead3, comp_done1, comp_done2, comp_done3 · Evidence: claimed, measured

- | | |
|----|---|
| 20 | Initiatives routinely stall. Nobody can say why, or who decides. |
| 50 | Decisions get made but slowly. Execution depends on the owner personally pushing. |
| 80 | Clear decision rights, named delivery owners, and evidence of initiatives actually completed. |

ldr-04 — Succession and capability transfer

Whether the ability to do critical work exists in more than one person, and whether it is written down well enough to transfer. Assesses the system that holds capability, never an individual's competence.

Intake: ins_lead5, ins_proc1, ins_proc2 · Evidence: claimed, measured

20	Critical work lives in one head. Nothing is documented, nobody is cross-trained, and a resignation would remove capability the business cannot rebuild quickly.
50	Some cross-cover exists, unevenly. Documentation is partial or out of date, so transfer depends on the departing person's goodwill and time.
80	Every critical role has a named cover and current documentation good enough to work from. Capability transfer has actually been tested, not merely assumed.

ldr-05 — Structure and span of control

Whether the shape of the organisation fits its size — how much reports directly to the owner, and whether anyone else carries real management load.

Intake: ins_lead6, b_employees · Evidence: claimed, measured

20	Everyone reports to the owner. There is no management layer, so the owner is the constraint on every decision and the business cannot grow past their personal bandwidth.
50	A partial layer exists, but spans are uneven and the owner still carries direct reports who should sit under someone else.
80	Spans are deliberate and appropriate to size. Management load is genuinely distributed, and the owner's direct reports are the ones that should be.

Processes 10%

prc-01 — Process documentation

Whether the work survives the person currently doing it.

Intake: ins_proc2, ins_lead1 · Evidence: measured, claimed

20	Nothing documented. Process lives entirely in individuals' heads.
50	Key processes documented but out of date, or documented and not followed.
80	Documented, current, and actually followed. New staff can be trained from them.

prc-02 — Manual and fragile steps

How much of the work depends on memory, a spreadsheet or paper.

Intake: ins_proc2 · Evidence: measured, claimed

20	Core commercial processes run on memory, a spreadsheet or paper. Single points of failure in the critical path.
50	Manual steps remain in non-critical paths. Known and tolerated.
80	Critical paths systematised. Remaining manual steps are deliberate choices and documented.

prc-03 — Improvement capability

Whether broken processes actually get fixed.

Intake: ins_proc1, ins_lead3, comp_done1 · Evidence: measured, claimed

- | | |
|----|---|
| 20 | Known problems have persisted for years. No mechanism exists to fix them. |
| 50 | Problems get fixed when they become acute, by whoever has time. |
| 80 | A named owner, a cadence, and evidence of processes measurably improved. |

Technology & Data 8%

tec-01 — System coverage

Whether the systems capture what the business needs to run.

Intake: s_accounting, s_pos, s_crm, s_marketing, s_other, atl_wire2 · Evidence: measured

- | | |
|----|--|
| 20 | Core activity is not captured in any system. Reporting is rebuilt by hand each time. |
| 50 | Systems in place per function but not integrated. Manual joins required to answer basic questions. |
| 80 | Systems cover the core and data flows between them without manual re-entry. |

tec-02 — Data trustworthiness

Whether the numbers the business relies on are trusted, and rightly so.

Intake: ins_tech2, an_num4, data_missing, data_basis · Evidence: measured, claimed

- | | |
|----|--|
| 20 | The owner does not trust the numbers, and is right not to. |
| 50 | Numbers broadly trusted, with known problem areas that get worked around rather than fixed. |
| 80 | Data quality actively managed. Reconciliations routine, discrepancies rare and traced to source. |

tec-03 — Decision visibility

Whether decision-makers can see what they need, when they need it.

Intake: ins_tech1, ins_tech3, atl_watch1, atl_watch3 · Evidence: measured, claimed

- | | |
|----|--|
| 20 | Key questions cannot be answered without a manual exercise taking days. |
| 50 | Standard reports exist. Anything ad-hoc requires the owner or an analyst. |
| 80 | Decision-makers self-serve the numbers that matter, on a cadence that fits the decision. |

Changing this rubric

1. Edit the JSON, never a rendered copy. templates/rubric.json is the instrument. This PDF, rubric.md and the website's lens page are all generated from it.

2. Bump the version. A weight or lens change is a major version; a sub-criterion or anchor change is a minor one. Scores stamp the version they used, so past engagements stay readable against the rubric that produced them.

3. Re-run the generators and the checks. `tools/render_rubric_pdf.py`, `tools/render_rubric.py` and `tools/validate_rubric.py`, then the website's `scripts/check-lenses.mjs` — which fails the deploy if the published weights and this file disagree.

Ratified by · Fenwick How, Architect & Managing Principal

2026-08-10