

Canonical Architecture Reference

The authoritative definition of the Aperture Method's architecture. Every future form, agent, assessment, model, deliverable and analytical method is placed against this document. Where any other artifact disagrees with it, this document governs.

VERSION 1.2	EFFECTIVE 10 August 2026	OWNER Fenwick How
STATUS Adopted · living	REVIEW On material change	CLASSIFICATION Confidential

This is a living document. It is expected to change as the Method is built out. Clauses marked **OPEN** record decisions that have not yet been made; they are part of the record, not omissions from it. Amendment procedure is at §13.

CONTENTS

1 Purpose, scope and authority

2 Canonical vocabulary

3 Object taxonomy

4 The Force5 agent layer

5 The engagement spine

6 Phase contracts

7 The evidence standard

8 The gate register

9 Load-bearing objects

10 Scoring authority

11 Placement rules

12 Change control

13 Open decisions

SECTION 1

Purpose, scope and authority

- 1.1 This document defines the architecture of The Aperture Method™: its objects, phases, contracts, gates and evidence standard. Its purpose is to give every future addition a defined place, so that the Method grows by extension rather than by accretion.
- 1.2 It governs the SOPs, the agent definitions, the intake instruments, the scoring rubric, the deliverable set, the commercial documents and the codebase. It does not govern commercial terms, pricing or legal drafting.
- 1.3 **Precedence.** Where artifacts conflict, authority runs in this order: **(a)** this document; **(b)** the Method SOPs; **(c)** the Agent Build Spec; **(d)** marketing and commercial material. A conflict resolved here **MUST** be corrected at source in the losing artifact; this document is not a place to hide a contradiction.
- 1.4 **Normative language.** **MUST** is a requirement; non-conformance is a defect. **SHOULD** is a strong default; departure must be recorded with a reason. **MAY** is discretionary.
- 1.5 The companion *Master Process Architecture* is descriptive — it shows the system as built at a point in time. This document is normative — it states what the system must be. Where the map shows something this document forbids, the map is reporting a defect, not granting permission.

The governing principle

Human-led, AI-enabled. Automation compresses the mechanical work — reading, reconciling, modelling, mapping, monitoring. Judgment, context and accountability remain with a named person. Agents produce findings; only a person produces conclusions. No clause in this document may be read to permit otherwise.

SECTION 2

Canonical vocabulary

Three naming conflicts were found in circulation. This section resolves them. These names are canonical; no other name for the same object may appear in any artifact.

2.1 Modules

PHASE	CLIENT-FACING NAME	VERB	INTERNAL AGENT	SOP
—	Business X-Ray™	Diagnose	Lena	01
01	Aperture Insights™	Understand	Lena	02
02	Aperture Analytics™	Quantify	Marcus	03
03	Aperture Intelligence™	Reveal	Natalie	04
04	Aperture Compass™	Navigate	Aria	05
05	Aperture Atlas™	Perform	Thomas	06

- 2.1.1

Aperture Atlas™ is the canonical name of Phase 05, used internally and with clients alike — one name, no translation layer. Ratified by Fenwick How on 10 August 2026, reversing the position taken in v1.1. **"Aperture Atlas" is retired** and **MUST NOT** appear in any artifact. The website mirrored the old name in eight places and now does not; `scripts/check-lenses.mjs` fails the site build if it returns. The contract, Exhibit A and the intake form, previously flagged non-conformant for saying "Atlas", are conformant as written.
- 2.1.2

There are five phases and six SOPs. The Business X-Ray is a product, not a phase: it is the fixed-fee wedge, and when a client engages the Method it is delivered inside Phase 01. Any artifact stating "six phases" is non-conformant.
- 2.1.3

Agent identities are internal. Clients engage modules, not names. Agent names **MUST NOT** appear in any client-facing artifact, including document metadata, sheet tabs, chart titles, alt text and filenames.

2.2 The seven lenses

Ratified. Three lens sets were once in circulation — the SOP 01 list, the Build Spec list with different weights (Customers 20 / Ops 15), and a third set scored in a delivered X-Ray. The SOP 01 list below is canonical, is what the scoring engine implements, and is now the only one: the website mirrors it in `src/lib/lenses.ts`, and

`scripts/check-lenses.mjs` fails the build if the two disagree, if the intake form's lens headings disagree, if a retired lens name reappears, or if the rubric is ever un-ratified.

#	LENS	KEY	WEIGHT	COVERS
1	Finance	finance	20%	Margin, cash, unit economics, working capital
2	Operations	operations	18%	Capacity, throughput, bottlenecks
3	Customers & Retention	customers-retention	18%	Acquisition, repeat behaviour, lifetime value
4	Market & Competition	market-competition	14%	Position, demand, local dynamics
5	Leadership & Team	leadership-team	12%	Key-person risk, roles, capability
6	Processes	processes	10%	How work gets done; where it breaks
7	Technology & Data	technology-data	8%	Systems, data quality, automation

2.2.1 "Marketing & growth engine" is **not** a lens. Marketing evidence is scored under Customers & Retention (acquisition) and Market & Competition (position). It was scored as a lens in the delivered Lumina X-Ray; that deliverable, its kickoff pack, the bound volume and the public example report have been rebuilt, and the name is now blocked at build time.

2.2.2 Any change to the lens set, its sub-criteria or its weights is a rubric version change and is governed by §10.

SECTION 3

Object taxonomy

Everything in the Method is exactly one of six object types. An artifact that cannot be typed has not been designed yet.

TYPE	DEFINITION	MUST DECLARE
Input	Something the client supplies: a document, a data export, system access, or a person's time.	Which phase needs it · blocking or helpful · the intake question that collects it
Method	An analytical technique applied to inputs: a calculation, model, rubric or framework.	Discipline · phase · inputs consumed · output produced · the evidence class it yields
Agent	An automated worker executing a defined SOP within a fixed tool allowlist.	Phase · tool allowlist · refusal rules · which steps are human
Output	An internal artifact: a score, a finding, a dataset, a model. Not sent to a client.	Producing tool · evidence refs · consumers
Deliverable	An artifact the client receives and keeps.	Phase · format · the gate it must pass · the pre-delivery check
Gate	A point where work stops until a named human decides.	Who decides · what evidence they see · what the two outcomes are

- 3.1

Every Method **MUST** declare the inputs it consumes, and every declared input **MUST** have a collection route — an intake question, a data request item, or a system connection. A method with no collection route is not yet part of the architecture.
- 3.2

Every Input **MUST** have at least one consumer. An input collected but consumed by nothing is a burden placed on the client for no return and **SHOULD** be removed or wired.
- 3.3

Every Agent tool allowlist is closed. A capability named in an agent's description but absent from its allowlist is **not callable** and **MUST NOT** be advertised as available.
- 3.4

Every Deliverable **MUST** pass the pre-delivery check before release, including the internal-identity scrub of §2.1.3.

SECTION 4

The Force5 agent layer

Force5 is the internal name for the five specialist agents that execute the SOPs. They are a labour layer, not a decision layer. This section is their constitution.

4.1 What the agent layer is

- 4.1.1** Each agent owns exactly one phase, holds one standard of evidence, and hands typed findings to the next. The Method executes as a **sequence, not a swarm**.
- 4.1.2** Agents compress the mechanical work — reading, reconciling, modelling, mapping, monitoring. They **MUST NOT** close the loop themselves. Every phase ends at a human gate (§8).
- 4.1.3** Agent identities are internal (§2.1.3). Clients engage the Aperture modules. No agent name, and no reference to "Force5", may appear in a client-facing artifact.

4.2 The roster

AGENT	MODULE	OWNS	STANDARD OF EVIDENCE
Lena The Strategist	Aperture Insights™	Seven-lens diagnostic · Aperture Score · constraint candidates	Refuses to render an uncited figure. Proposes constraints; never names one.
Marcus The Analyst	Aperture Analytics™	Financial model · Profit Map · scenarios	Never does arithmetic in prose. Reconciles to the dollar or reports the break.
Natalie The Explorer	Aperture Intelligence™	Geocoding · segments · trade areas · site selection	Never sees a street address. PII drops before analysis begins.
Aria The Navigator	Aperture Compass™	Opportunity Matrix · sequencing · Focus Plan	Every entry cites its upstream finding, or it does not render.
Thomas The Operator	Aperture Atlas™	KPIs · Scoreboard · re-score · monthly brief	Escalates; never acts. Every dashboard value reconciles to source.

4.3 The agent constitution

These rules bind every agent. They are not style guidance; they are the terms on which automation is permitted to touch client work at all.

- 4.3.1** **Closed allowlist.** Each agent may call only the tools on its list. A capability named in an agent's description but absent from its allowlist is not callable and **MUST NOT** be advertised.

- 4.3.2 No gate may be closed by an agent.** Where a step is hybrid, the agent proposes and a person disposes.
- 4.3.3 No client contact.** Agents do not send, present, or correspond. Every readout is delivered by a person.
- 4.3.4 Findings, never conclusions.** An agent surfaces what the evidence shows. What it means, and what to do, is a human output.
- 4.3.5 Refusal is correct behaviour.** An agent that cannot meet its evidence standard **MUST** stop and say so. A blocked state is a legitimate terminal state; a plausible answer produced without evidence is a defect, not a service.
- 4.3.6 No rubric authorship.** Agents read the rubric, the weights and the criteria. They **MUST NOT** invent, adjust or extend them (§10).
- 4.3.7 Provenance.** Every figure an agent emits carries its evidence references and the tool that produced it (§8).

4.4 Tool allowlists

Status: ● implemented · ○ contract defined, body not built.

AGENT	TOOLS	BUILT
Lena	scoring.score_lens ● · scoring.compute_aperture_score ● · benchmarks.lookup ○ · constraints.propose ○ · interviews.triangulate ○ · measures.baseline ○ · render.insights_doc ○	2 / 7
Marcus	build_financials · build_profitmap · allocations.propose · pareto.rank_drivers · breakeven.sensitivity · scenario.run · recalc · render.analytics_doc — all ○	0 / 8
Natalie	geocode · segments.clv · gis_analysis · gis_layers · demographics.overlay · syncpoint.query · site_model.score · render.intelligence_doc — all ○	0 / 8
Aria	opportunities consolidate · matrix.score · scenario.stress_test · sequencer.order · render.compass_doc — all ○	0 / 5
Thomas	kpi.define · pipelines.wire · feed.validate · build_dashboard · anomaly.detect · scoring.compute_aperture_score ● · render.signal_brief — rest ○	1 / 7

- 4.4.1** Until a tool is built, its work is performed by a person. An unbuilt tool is not a missing feature — it is unautomated labour, and it is where the Method depends on Fenwick rather than on the system.

4.5 Human steps, by phase

These are canonical. Where an SOP omits one, the SOP is non-conformant (see §6).

PHASE	STEPS THAT MUST BE PERFORMED BY A PERSON
X-Ray + 01	Kickoff and data request · every interview · naming the binding constraint · the readout
02 Analytics	Approving the cost-allocation basis before computation · specifying which scenarios are modelled · the readout
03 Intelligence	Approving trade-area parameters and scoring criteria before ranking · owning and presenting every site recommendation
04 Compass	Agreeing criteria and weights with the owner · owning the sequence that goes to a board · the readout
05 Aperture Atlas	Approving every KPI definition and threshold · setting the cadence · reading every brief before the client does

Why the constitution is short

Every rule in §4.3 exists to protect one thing: that a client is buying a person's judgment, supported by machinery, rather than machinery wearing a person's name. An agent layer that closes gates, contacts clients or states conclusions is a different product with a different risk profile, and it is not this one.

SECTION 5

The engagement spine

- 5.1 The canonical sequence is: **Commercial acceptance** → **Intake** → **Data collection** → **Ledger** → **Business X-Ray** → **Phase 01** → **Phases 02 and 03** → **Phase 04** → **Phase 05** → **loop**.
- 5.2 Phases 02 and 03 **MAY** run concurrently. Both consume Phase 01's typed questions and neither depends on the other.
- 5.3 Phase 04 **MUST NOT** begin until at least one of Phases 02 or 03 has been gated. It scores what those phases produce; with neither, it has nothing to score.
- 5.4 Phase 05 **MUST NOT** begin without an approved Focus Plan and agreed success measures. A Scoreboard built without them measures activity rather than strategy.
- 5.5 **Standalone purchase.** Any module **MAY** be sold alone. When a module runs without its normal upstream input, it **MUST** declare the substitution on the face of the deliverable — for example: *"Derived from financials — not owner-validated."* Silent degradation is prohibited.
- 5.6 Every phase ends at a gate (§8). A phase **MUST NOT** begin on ungated inputs.

SECTION 6

Phase contracts

Each phase is an interface. It states what it requires, what it guarantees, and what it hands on. A phase that cannot meet its guarantee stops at its gate and says so.

5.1 Business X-Ray™ + Phase 01 · Insights

REQUIRES	P&L 3yr + YTD · balance sheet 3yr · revenue by product and location · payroll by role · capacity and volume · payment terms and debt · 3–5 interviews · read-only accounting access
GUARANTEES	Seven lenses scored against the ratified rubric with evidence · a baseline Aperture Score with stamped version · ranked constraint candidates with confirming <i>and</i> disconfirming evidence · a prescription of components in order
HANDS ON	The named constraint · 3–5 success measures with baselines · typed questions to Phases 02–04
HUMAN STEPS	Kickoff and data request · all interviews · naming the constraint · the readout

5.2 Phase 02 · Analytics

REQUIRES	Transaction-level export · chart of accounts · owner compensation and add-backs · price list · volume by product, location and month · accounting basis
GUARANTEES	A model reconciling to the P&L exactly · contribution and fully-absorbed profit reported separately · an approved and documented allocation basis · driver-based scenarios with stated assumptions · independent recomputation passed
HANDS ON	Shortlisted moves with size, confidence and failure conditions · the scenario model
HUMAN STEPS	Approve the allocation basis before any figure is computed on it · specify which scenarios are modelled · the readout

5.3 Phase 03 · Intelligence

REQUIRES	Customer records with addresses or ZIPs · purchase history joinable by customer ID · location list · marketing spend by area · competitor set
GUARANTEES	Geocoding performed locally with street-level PII dropped before analysis · match rate reported with the concentration of misses · penetration computed over addressable households · observed, modelled and estimated separately symbolised · ranking sensitivity to parameters stated

HANDS ON	Ranked areas and sites with geography, size, vintage and confidence
HUMAN STEPS	Approve trade-area parameters and scoring criteria before any ranking · own and present every site recommendation

5.4 Phase 04 · Compass

REQUIRES	A named constraint · agreed success measures · shortlisted moves from 02 and/or 03 · owner limits on capital, capacity, risk and timing
GUARANTEES	Every matrix row cites its upstream finding · criteria and weights agreed before scoring and stamped · constraint-first sequencing with expansion last · displacement stated for every "Now" item · a one-page Focus Plan
HANDS ON	Approved initiatives with intended outcomes and the measures that would show them working
HUMAN STEPS	Agree criteria and weights with the owner · own the sequence that goes to a board · the readout

5.5 Phase 05 · Aperture Atlas

REQUIRES	Approved Focus Plan · success measures · live system connections with authorisation · Aperture Score baseline
GUARANTEES	Every KPI traces to an approved initiative and has a named human owner · every dashboard value reconciles to its source on build and refresh · re-scoring uses the stored rubric version · the client can operate the Scoreboard themselves
HANDS ON	Deviations routed back to the phase that can diagnose them · re-opened constraint candidates
HUMAN STEPS	Approve every KPI definition and threshold · set the cadence · read every brief before the client does

Non-conformance on record

Phases 02, 03 and 04 currently contain no human step in their written SOPs, while the Build Spec assigns Fenwick approval duties in all three. The human steps stated in 5.2–5.4 above are canonical; the SOPs must be amended to carry them.

The evidence standard

- 7.1** Every figure in every deliverable **MUST** carry an evidence reference resolving to a ledger entry.
- 7.2** Every figure **MUST** be produced by a tool that can be re-run. A number produced by reasoning is not evidence and **MUST NOT** be published.
- 7.3** Every value carries exactly one class:

CLASS	MEANING	MAY SUPPORT A SCORE
Measured	From a client system or document	Yes
Derived	Computed by a named tool from measured inputs	Yes
Assumed	A judgment input; cites who set it and on what basis	Only where the rubric permits
Claimed	Asserted in interview, not yet corroborated	No, until promoted

- 7.4** **Reconciliation.** Two thresholds exist and **MUST NOT** be confused. **Ledger acceptance** tolerates aggregate variance of $\leq 0.5\%$ provided every variance is itemised — it is a completeness threshold. **Figure publication** tolerates nothing: any figure entering a deliverable must reconcile exactly to its source. Accepting a ledger does not license an unreconciled figure.
- 7.5** Insufficient evidence is a finding. A lens, criterion or figure that cannot be evidenced **MUST** be marked insufficient and **MUST NOT** be estimated. Insufficient is never scored as zero, and its weight is never redistributed.
- 7.6** Ledger references are append-only. A superseded entry is retained and marked; a restatement by the client is itself a finding.
- 7.7** Confidence is recorded separately from score. A well-evidenced 40 and a thinly-evidenced 40 are different findings.

SECTION 8

The gate register

"Seven gates" was referenced across the source material but never enumerated. This is the enumeration. A gate is blocking unless marked advisory.

GATE	NAME	DECIDED BY	PASSES WHEN
G1	Scope accepted	Fenwick + client	Scope, exclusions, fee and timeline are signed. Exclusions are recorded explicitly.
G2	Ledger accepted	Fenwick	Sources ingested, reconciled within tolerance, every variance itemised, gaps recorded. Nothing downstream runs until this passes.
G3	Constraint named	Fenwick, then owner	One constraint named from ranked candidates, agreed with the owner in writing.
G4	Basis approved	Fenwick	Phase-specific judgment inputs approved <i>before</i> computation: allocation basis (02), trade-area parameters and scoring criteria (03), matrix criteria and weights (04), KPI definitions and thresholds (05).
G5	Phase output accepted	Fenwick	Guarantees in the phase contract met; recomputation and reconciliation passed; insufficiencies marked.
G6	Deliverable released	Fenwick	Pre-delivery check clean, including the internal-identity scrub. Reviewed, edited and signed by a person.
G7	Continue or close	Client	Client decides the next component, or moves to Aperture Atlas, or closes. Decision and reasoning recorded.

- 8.1 Each gate **MUST** produce a record: who decided, when, on what evidence, and the reasoning. On a retainer, the reasoning is what makes the Scoreboard interpretable a year later.
- 8.2 A blocked state is a legitimate terminal state. Partial artifacts **MUST NOT** be released to work around a gate.
- 8.3 Agents **MUST NOT** close a gate. Where a step is marked hybrid, the agent produces a proposal and a person disposes of it.

SECTION 9

Load-bearing objects

Four objects are consumed by three or more downstream components. If one is wrong, everything after it is wrong and nothing will say so. Each requires a defined home, not a mention in prose.

OBJECT	CONSUMERS	REQUIREMENT
The constraint	Insights validation · Analytics tie-back · Compass scoring · Compass sequencing · Live KPI selection	MUST be a recorded artifact with its evidence, its disconfirming evidence, and the date the owner agreed it. Not prose.
Success measures	Live KPIs · re-score baseline · Compass inputs · Focus Plan · Scoreboard	3–5 measures, each with unit, baseline value, target range, cadence and owner. MUST be structured, not narrative.
Rubric version	Every lens score · every composite · every future re-score	MUST be stamped on every score. Re-scoring uses the stored version. See §10.
The ledger	Every phase, every agent	MUST be a single append-only store per engagement, holding every figure's source, class and provenance.

9.1 A load-bearing object **MUST NOT** exist only inside a document. If it cannot be read by a tool, it cannot be enforced, and the guarantees in §6 and §7 become aspirations.

Scoring authority

- 10.1** The rubric — lenses, sub-criteria, anchors and weights — is the practice's intellectual property. Agents read it; they **MUST NOT** invent, adjust or extend it.
- 10.2** An unratified rubric **MUST NOT** produce a client-facing Aperture Score. This is enforced in code, not by discipline.
- 10.3** Sub-criteria are scored 0–100 against anchor descriptions at 20, 50 and 80. A lens score is the mean of its scored sub-criteria. The Aperture Score is the weighted mean of lens scores.
- 10.4** Below 70% coverage no composite is published. The deliverable in that case is a data request.
- 10.5** **Comparability.** Re-scoring **MUST** use the stored rubric version and weights, never the current ones. Where the rubric changes, history is either restated explicitly or labelled non-comparable. Silent mixing of versions is prohibited — it destroys the only thing a baseline is for.
- 10.6** Rubric changes are versioned semantically: a weight or lens change is a major version; a sub-criterion or anchor rewording is a minor version.

SECTION 11

Placement rules

This is the operative section. Anything new — a form, an agent, an assessment, a model, a deliverable, an MBA-level method — is placed by answering these questions in order. If any answer is unavailable, the thing is not ready to be built.

10.1 The placement procedure

- 1 **What type is it?** One of the six in §3. If it is two things, it is two objects and must be split.
- 2 **Which phase owns it?** Exactly one phase owns it. Others may consume it. If no phase owns it, it belongs to a cross-cutting layer and must be declared as such.
- 3 **What does it consume?** Every input named, and for each: the collection route that already exists, or the new intake question required. A method whose inputs are not collected is not yet buildable.
- 4 **What does it produce?** The output object, its evidence class, and every downstream consumer. An output with no consumer is not an output.
- 5 **Which gate does it pass?** Name the gate from §7 and who decides. If it requires a judgment input, that input passes G4 *before* computation.
- 6 **What does it displace?** Does it duplicate or supersede something? Retire the predecessor in the same change; do not leave both.
- 7 **How is it evidenced?** The tool that produces it, and how a hostile reader would trace its figures to source.
- 8 **Register it.** Amend this document and the Master Process Architecture in the same change.

10.2 Placing a new analytical method

The Method's differentiation rests on applying graduate-level quantitative work to businesses that cannot otherwise buy it. That claim is only defensible where the work is actually done.

- 11.2.1 A method **MUST** declare: its discipline; the phase that owns it; the inputs it requires; the output it produces; the decision it informs; and the evidence class of that output.
- 11.2.2 A method **MUST NOT** be advertised as applied until it has appeared in a delivered client artifact with its working shown. Until then it is **catalogued**, not applied, and **MUST** be labelled as such in internal material.
- 11.2.3 A method **MUST NOT** be included whose inputs the intake does not collect. Adding the method and the question is one change, not two.

- 11.2.4** Sophistication is not a justification. A method earns its place by changing a decision the owner would otherwise make differently — not by being advanced.

10.3 Placing a new client input

- 11.3.1** Every new question **MUST** name its consumer — a rubric sub-criterion, an SOP step, or a named method — before it is added to any instrument.
- 11.3.2** Every input is classified **blocking** or **helpful**, and carries a plain-language reason the client can read.
- 11.3.3** References between instruments and the rubric **MUST** use stable identifiers, never positional numbering. Positional references break silently when a question is inserted.

10.4 Placing a new deliverable

- 11.4.1** A deliverable **MUST** name the phase, the executive question it answers, the gate it passes, and the format the client keeps.
- 11.4.2** If it cannot be traced to an existing phase contract's guarantee, either the contract is amended or the deliverable is not in scope.

10.5 Placing a new agent or tool

- 11.5.1** A tool **MUST** be added to exactly one agent's allowlist, with a documented contract and its refusal conditions.
- 11.5.2** An agent **MUST NOT** gain a capability that closes a gate, contacts a client, or states a conclusion.

Change control

- 12.1** This document is amended by Fenwick How. No other artifact may contradict it; a contradiction is a defect in that artifact.
- 12.2** Versioning: **major** for a change to phases, object types, gates or the evidence standard; **minor** for a phase contract, placement rule or vocabulary change; **patch** for clarification without behavioural effect.
- 12.3** Every amendment **MUST** record what changed, why, and which artifacts must be updated to conform.
- 12.4** An amendment affecting the rubric, the evidence standard or a gate **MUST** state its effect on engagements already in flight, including whether prior scores remain comparable.
- 12.5** The Master Process Architecture **SHOULD** be regenerated with each amendment so the descriptive and normative views do not drift.

Amendment record

VERSION	DATE	CHANGE
1.2	10 Aug 2026	Ratifies the seven lenses across every artifact and adds the build-time conformance guard. Reverses §2.1.1: Aperture Atlas™ becomes the single name for Phase 05, internal and client-facing; "Aperture Live" is retired. Closes D1 (the SOW fee schedule governs; Exhibit A rewritten) and D2. Records the Lumina rebuild.
1.1	10 Aug 2026	Adds §4, the Force5 agent layer and its constitution. Applies the brand system.
1.0	10 Aug 2026	Initial adoption. Resolves module naming, phase count, the lens set and the reconciliation tolerances. Enumerates the seven gates for the first time. Establishes the placement procedure.

SECTION 13


Open decisions

Recorded, owned and dated. An open decision is not a gap in this document; it is a decision that has not yet been made, held where it can be seen.

#	DECISION	BLOCKS	STATUS
D1	Reconcile the two published fee schedules. Closed 10 Aug 2026: the SOW schedule governs — X-Ray \$9,500 fixed (credited toward the Full Method within 60 days), Full Method \$41,000, Atlas \$2,950/month on a 6-month minimum. Exhibit A was rewritten to match and its "from" floors removed. Individual phase fees remain quoted per engagement; the SOW never listed them, and none were invented.	—	Closed · commercial
D2	Correct "Aperture Atlas" to "Aperture Atlas" in the contract, Exhibit A and the public intake form. Closed 10 Aug 2026, the other way: Atlas is the single name, internal and client-facing. The website was corrected to match; the contract and intake form needed no change.	—	Closed · §2.1.1
D3	Amend SOPs 03, 04 and 05 to carry the human steps stated in §6.2–6.4.	Gate G4 conformance	Open · procedural
D4	Define the Compass matrix criteria and weights, and version them.	Phase 04 reproducibility	Open · IP
D5	Complete sub-criteria to 4–6 per lens. Three lenses currently carry three.	Rubric v2	Partial
D6	Decide whether the Signal Brief and anomaly detection are in scope for Phase 05. Both are specified and priced but absent from SOP 06.	Aperture Atlas scope and pricing	Unclear
D7	Confirm whether the hosted Scoreboard is delivered or specified. Currently sold as live, delivered as a data feed.	Phase 05 guarantee in \$5.5	Open · material
D8	Add a Business X-Ray segment to the intake form. The wedge has no intake path.	The most common single purchase	Open
D9	Adopt or discard SOP 00 (Intake). If adopted, X-Ray Step 2 duplicates it.	Documentation set	Unclear
D10	Confirm SyncPoint AI terms — query limits, cost, redistribution rights.	Phase 03 external data	Open · commercial

How to close one

Make the decision, amend the affected artifacts in the same change, move the row to the amendment record at §11, and regenerate the Master Process Architecture. A decision closed in conversation but not written here is not closed.

A handwritten signature in black ink, reading "Fenwick How". The signature is written in a cursive, flowing style with a horizontal line extending from the end.

Fenwick How

Founder · Architect & Managing Principal

Adopted 10 August 2026