



THE APERTURE METHOD™ · MASTER PROCESS ARCHITECTURE

What we have **actually** built

Reconstructed from the SOPs, the Agent Build Spec, the Lumina engagement, the SOW, the contract and the live codebase. Nothing here is proposed — this is the system as it exists on 10 Aug 2026.

Bring your business into focus.

The Aperture Method moves an owner-run business from a first honest read to a living, visual intelligence platform. Each segment answers one executive question and produces one tangible deliverable the client keeps.

Who it is for: owner-run businesses, roughly \$1M–\$20M in revenue, place-based or data-rich. Senior-led throughout — the person who does the work is the person the client meets.

- Senior-led throughout.** No handoff to a junior team.
- Fixed-fee and phase-gated.** The client decides whether to continue at each gate — never an open-ended commitment.
- Plain language.** If we cannot explain why, we do not recommend on it.
- The client owns what we build.** Models, dashboards and playbooks built on their data are theirs to keep.
- A person reviews every output.** AI supports decisions; it does not make them.
- Measured on results.** We ask to be measured on whether the strategy shows up in the numbers.

The offer

The Business X-Ray is the fixed-fee front door and the sales wedge. It produces the same core deliverable as Insights; when a client engages the full Method it is delivered inside the Insights phase.

#	SEGMENT	VERB	EXECUTIVE QUESTION	DELIVERABLE	DURATION
—	Business X-Ray™	Diagnose	Where do we stand, and what do we fix first?	Business X-Ray + Aperture Score™	2–3 wk
01	Aperture Insights™	Understand	What is happening?	Business X-Ray + Aperture Score	2–3 wk
02	Aperture Analytics™	Quantify	Why is it happening?	Profit Map + scenario model	3–4 wk
03	Aperture Intelligence™	Reveal	What opportunities does it reveal?	Customer & Market Map + GIS package	2–3 wk
04	Aperture Compass™	Navigate	Where should we go next?	Opportunity Matrix + Focus Plan	1–2 wk
05	Aperture Live™	Perform	How do we sustain results?	Scoreboard + KPI system	Ongoing

How an engagement runs

1

Lead & fit call

Qualify the business and confirm the problem is one the Method fits.

2

Intake

Client completes the questionnaire. Save and resume; nothing is lost.

3

Proposal & contract

Scope a single segment, several, or the full arc. Fixed fee per phase.

4

Kickoff & data request

Issue the personalised data request and secure least-access to systems.

5

Deliver segment

Run the SOP. Present a readout. Secure the phase-gate decision before continuing.

6

Move to Live

Stand up the Scoreboard; shift to a light recurring cadence.

7

Steer & re-score

Monthly steering, quarterly review on live data, re-scored Aperture Score to prove progress.

What powers it

Human-led, AI-enabled. Five internal specialist agents compress the mechanical work — the reading, reconciling, modelling, mapping and monitoring — so judgment, context and accountability stay with a person the client can look in the eye. The agents produce findings, never conclusions. **Agent identities are internal; clients engage the Aperture modules, not the names.**


INTERNAL	CLIENT-FACING MODULE	OWNS	STANDARD OF EVIDENCE
Lena · Strategist	Aperture Insights™	Seven-lens diagnostic, Aperture Score, constraint candidates	Refuses to render an uncited figure
Marcus · Analyst	Aperture Analytics™	Financial model, Profit Map, scenarios	Never does arithmetic in prose; reconciles to the dollar
Natalie · Explorer	Aperture Intelligence™	Geocoding, segments, trade areas, site selection	Never sees a street address; PII drops at the door
Aria · Navigator	Aperture Compass™	Opportunity Matrix, sequencing	Every row cites its upstream finding
Thomas · Operator	Aperture Live™	KPIs, Scoreboard, re-score	Escalates; never acts

How to read the rest of this document. Level 1 is the whole Method on one screen. Level 2 opens each phase into inputs, methods, agents, outputs and deliverables. Level 3 shows how information actually moves between them. The Gap Register lists what is missing, conflicting or orphaned. Status marks throughout: ● exists · ◐ partial · ○ missing · ◆ unclear.

OBJECT TYPE	BUILD STATUS	READING IT
Input · client data Method · analysis AI agent · automation Output · internal Client deliverable Human gate · Fenwick	Exists — developed Partial — started Missing — required, absent Unclear — conflicting	Vertical spine = the client's journey. Left band = cross-cutting layers that serve every phase, not phases themselves. Dashed edge = loop or optional branch.


The whole Method on one screen

Five phases sit inside a much larger machine. The commercial wrapper, the intake, the ledger and Fenwick's gate layer are cross-functional — they touch every phase and are where most of the unbuilt work sits.



Cross-cutting layers

These serve every phase rather than sitting inside one. Four of the six exist only as specification — which is why the Method currently runs on Fenwick rather than on the machine.




Where the machine is thin

The client-facing surface is largely built. The layer underneath it — the thing that makes the Method repeatable rather than heroic — is mostly specification.

LAYER	STATUS	EVIDENCE
Client-facing deliverables	● Exists	8 real documents produced end-to-end for Lumina
SOPs 01–06	● Exists	Written, with inputs, steps, gates, handoffs
Intake + data request	● Exists	108 questions, save/resume, personalised data request, tracking views
Aperture Score engine	● Exists	Ratified v1.1.0, deterministic, version-stamped, 28 tests
Force5 agents	◀ Partial	Roles, tool allowlists and refusals written; tool bodies unimplemented
Phase tooling	○ Missing	~20 named tools; only <code>scoring.py</code> real
Ledger / provenance store	○ Missing	DuckDB store, token map, provenance envelope — spec only
Dispatcher + gate engine	○ Missing	2 of 7 gate checklists exist in YAML; 25 assertions unimplemented
Hosted Scoreboard	○ Missing	Delivered as CSV feed + HTML mock-up; nothing wired or hosted
Pricing / commercial terms	◆ Unclear	SOW and Contract carry two irreconcilable schedules

Component view — inputs → methods → agents → outputs → deliverables

One column per phase. Read left to right within each. The agent row is where the practice is most exposed: the roles are defined, the tools are named, and almost none of the tool bodies exist.




Human gates — where a person must act

Six steps across the six SOPs are marked as Fenwick's. Three phases contain none at all, while the Build Spec claims he delivers every readout. That contradiction is the single largest structural inconsistency in the system.

PHASE	FENWICK STEPS	CLAIMED VS WRITTEN
Business X-Ray	Kickoff · Interviews · Readout	Consistent
01 Insights	Validate constraint · Readout	Consistent
02 Analytics	None	Spec says he approves allocation basis + scenarios
03 Intelligence	None	Spec says he owns every site recommendation
04 Compass	None	Spec says he owns the sequence
05 Aperture Live	Set cadence	Spec says he reads every brief first

Dependency view — how information actually moves

Thick crimson edges are the load-bearing objects: things consumed by three or more downstream components. If one of these is wrong, everything after it is wrong and nothing will say so.



LOAD-BEARING OBJECT	CONSUMERS	STATUS
The binding constraint	Insights validation · Analytics tie-back · Compass scoring · Compass sequence · Live KPI selection	Named by Fenwick, no artifact records it
success_measures[]	Live KPIs · re-score baseline · Compass inputs · Focus Plan · Scoreboard	Defined in SOP 02, no tool produces it
Rubric + weights version	Every lens score · composite · every future re-score	Ratified, stamped, enforced in code
The ledger	All five agents read it directly	Store does not exist
Aperture Score baseline	X-Ray · Live re-score · client trend line	Engine live
Scenario model	Compass stress-test · Scoreboard forecast	Built by hand for Lumina

THE DIFFERENTIATOR

A marketing agency optimises ad clicks. A CPA reconciles last year.

We apply financial modelling, marketing science, management science, spatial econometrics and strategy — the analytical depth a \$1M–\$20M owner-run business can otherwise only buy from a firm that will not take its call.

What it costs the client, and what they get for it

Left column: the inputs a phase cannot run without. **Blocking** means work stops; **helpful** means the answer gets better. Right column: the methods applied to those inputs. ● applied in a delivered engagement · ◌ in the method catalogue, not yet evidenced in client work.

BUSINESS X-RAY™ + 01 INSIGHTS™				UNDERSTAND · 2–3 WK	
WHAT THE CLIENT PROVIDES	QUANTITATIVE ANALYSIS APPLIED				
	Financial		Customer		Management science
	• Ratio & DuPont decomposition		• Cohort & retention survival analysis		• Theory of Constraints
	• Unit economics — CAC, LTV, LTV:CAC, payback		• Customer lifetime value modelling		• Capacity & utilisation — Little's Law, queuing
	Marketing science		Strategy		Scoring
	• Marketing-mix modelling & attribution		• Porter's Five Forces & value chain		• Seven-lens rubric → weighted Aperture Score™
02 APERTURE ANALYTICS™				QUANTIFY · 3–4 WK	
WHAT THE CLIENT PROVIDES	QUANTITATIVE ANALYSIS APPLIED				
	Financial		Customer		Marketing science
	• Contribution-margin & CVP analysis		• RFM & cluster segmentation — k-means		• Price-elasticity analysis
	• Break-even analysis		• Market-basket / affinity analysis		
	• Scenario & sensitivity analysis				
	• Capital budgeting — NPV, IRR, payback				
	Statistics		Management science		
	• Regression — linear, multiple, logistic		• Optimisation / linear programming		
	• Time-series forecasting — smoothing, ARIMA, seasonal decomposition				

03 APERTURE INTELLIGENCE™				REVEAL · 2–3 WK	
WHAT THE CLIENT PROVIDES	QUANTITATIVE ANALYSIS APPLIED				
	Spatial • Site-selection suitability — weighted overlay / MCDA ▸ Trade-area and drive-time modelling ▸ Gravity / Huff trade-area models ▸ Hot-spot Getis-Ord Gi* & spatial autocorrelation Moran's I ▸ Geodemographic segmentation	Customer • Segmentation and CLV by geography	Statistics ▸ Churn / propensity modelling — logistic regression		
• Customer list with addresses or ZIP codes • Purchase history joinable by a shared customer ID • Every location — address, opened when, size ◦ Marketing spend by channel or area ◦ Competitors that actually cost you business					

04 APERTURE COMPASS™				NAVIGATE · 1–2 WK	
WHAT THE CLIENT PROVIDES	QUANTITATIVE ANALYSIS APPLIED				
	Decision science • Multi-criteria decision analysis — weighted matrix ▸ Decision trees & expected value • Scenario stress-testing against the Analytics model	Strategy ▸ Ansoff & growth-share matrices ▸ Real-options sequencing	Financial ▸ NPV / IRR on the shortlisted moves		
• Capital available, and how fast it releases • Whose time an initiative would actually take • What is genuinely off the table ◦ Appetite for a year-three payoff versus this year • Who must be persuaded, and who would own delivery					

05 APERTURE LIVE™				PERFORM · ONGOING	
WHAT THE CLIENT PROVIDES	QUANTITATIVE ANALYSIS APPLIED				
	Performance • Balanced Scorecard KPI architecture	Statistics ▸ Statistical process control ▸ Time-series forecasting & anomaly detection	Experimental ▸ A/B testing & experimental design		
• Live connections — accounting, POS, CRM, marketing • Who authorises system access • For each number: what means on track, what means call me • A named human owner per number ◦ The meeting this should feed, and when it runs ◦ Anything that cannot connect and needs manual upload	Marketing science ▸ Attribution & channel value	Scoring • Re-score on stored rubric version			

Honest read on the catalogue

Of roughly 33 methods named across the five phases, 11 are evidenced in the Lumina engagement. The remaining 22 are real capability claims that no delivered artifact yet demonstrates. That gap is not a problem with the positioning — the positioning is the right one — it is the build list that makes the positioning true. Each one becomes defensible the first time it appears in a client deliverable with its working shown.

Gap register

Grouped by kind. Ordered within each by how much damage it does if left alone.

Conflicts — two sources of truth

ITEM	CONFLICT
Pricing	SOW \$41,000 bundle / \$9,500 entry / \$2,950 mo — Contract \$25,000 / \$2,500 / \$1,000 mo
Atlas vs Live	Internal name appears in the signed contract, Exhibit A fee table, and the public web intake
Phase count	Contract lists 6 · SOW and deliverables show 5 · SOPs numbered 01–06 · deliverables numbered 1–5
Who runs the readout	Spec: Fenwick delivers every readout. SOPs: three phases mark readout as AGENT
Reconciliation tolerance	Intake accepts 0.5% · X-Ray 0.005 · Analytics 0.0. Marcus must tie to the dollar against a ledger accepted at half a percent

Resolved — closed since the first audit

ITEM	RESOLUTION
The seven lenses	Ratified. Finance 20 · Operations 18 · Customers & Retention 18 · Market & Competition 14 · Leadership & Team 12 · Processes 10 · Technology & Data 8. One source (<code>rubric.json v1.1.0</code>) with one site-side mirror (<code>src/lib/lenses.ts</code>). "Marketing & growth engine" is retired — marketing evidence scores inside Customers & Retention. The Lumina X-Ray, the kickoff pack, the intake form and the public pages were rebuilt to match.
Score weights	The ratified rubric governs; the Build Spec figures (Customers 20 / Ops 15) are superseded. <code>scripts/check-lenses.mjs</code> runs on <code>prebuild</code> and fails the deploy if the site's names, order or weights drift from the rubric, if the intake form's lens headings disagree, if a retired lens name reappears anywhere in <code>src</code> , or if the rubric is ever un-ratified.

Missing — logically required, absent

ITEM	CONSEQUENCE
~20 phase tools	Every number is produced by hand. The Method is not yet repeatable without Fenwick doing the work
Ledger / provenance store	Nothing enforces "every figure traces to source" outside the scoring engine
Hosted Scoreboard	Sold as live and wired; delivered as a monthly CSV drop and an HTML mock-up
Signal Brief · anomaly detection	Priced and described in the Build Spec; absent from SOP 06 entirely
Exhibits A and B	Blank in the signed contract template — no fee schedule, no record of system access granted
NDA	Referenced three times across intake and SOW; no artifact exists
Readout artifacts	Every SOP ends in a readout; no deck, agenda or notes template exists for any phase
Interview guide + record	SOP 01 requires 3–5 interviews; Lumina shows none
X-Ray in the web intake	The wedge — "the most common single purchase" — has no segment in the intake form
Gate assertions	25 of 31 SOP checklist items have no machine counterpart

Orphaned — defined, never used

ITEM	NOTE
~20 named frameworks	Porter, PESTLE, VRIO, BCG, MECE, SCQA, ICE/RICE, TAM/SAM/SOM, OKR, Balanced Scorecard — credentialed but outside every tool allowlist, so non-callable by rule
~45 advertised methods	Playbook advertises ARIMA, k-means RFM, churn models, Huff gravity, Moran's I, LP optimisation. Lumina evidences roughly seven
SyncPoint AI	Named in the contract and SOP 06; used nowhere in the delivered engagement
Seven intake questions	Feed no rubric criterion and no SOP step
Stale agent filenames	arthur · lin · marco · clara · darius — do not match the Force5 roster