

Method SOPs

Every segment: what it is, what it does, what it needs, and how we run it

- 01 **Aperture Snapshot™** · UNDERSTAND
 - 02 **Business X-Ray™** · UNDERSTAND
 - 03 **Aperture Analytics™** · QUANTIFY
 - 04 **Aperture Intelligence™** · REVEAL
 - 05 **Aperture Compass™** · NAVIGATE
 - 06 **Aperture Atlas™** · PERFORM
-

VERSION 2.0 **ISSUED** 25 September 2026 · 2:48 PM CT **SOURCE** Edition 01 of the Master Operating Manual

INTERNAL · Generated from the Master Operating Manual, which governs. Do not edit this book directly: change the Manual, then rebuild.

CONTENTS

What is in this book

No.	Segment	Phase	Page
00	House standards	Applies to every segment	3
·	How the segments fit together		13
01	Aperture Snapshot™	UNDERSTAND	14
02	Business X-Ray™	UNDERSTAND	21
03	Aperture Analytics™	QUANTIFY	45
04	Aperture Intelligence™	REVEAL	55
05	Aperture Compass™	NAVIGATE	63
06	Aperture Atlas™	PERFORM	70

Versions in this book

Book	Version 2.0, issued 25 September 2026 · 2:48 PM CT
Source	Master Operating Manual, Edition 01 (24 August 2026)
Manual Part 0	House standards · version 01 · In force
Manual Part 1	Aperture Insights™ · version 01 · In force · bench partially built
Manual Part 2	Aperture Analytics™ · version 01 · In force · bench built
Manual Part 3	Aperture Intelligence™ · version 01 · In force · bench to build
Manual Part 4	Aperture Compass™ · version 01 · In force · bench to build
Manual Part 5	Aperture Atlas™ · version 01 · In force · bench to build
Fees	Read from the website pricing file at build time
Questionnaire	Read from the intake field catalog used by the Manual

This book replaces the earlier "Method SOPs" binder (7 and 13 August 2026 PDFs). It is derived: change the Manual, the pricing file or the intake catalog, then rebuild with build_sop.py.

00

HOUSE STANDARDS · APPLIES TO EVERY SEGMENT

House standards

The rules that do not vary by phase

*The question it answers: **What is true of every piece of work we do?***

VERSION 2.0 ISSUED 25 September 2026 · 2:48 PM CT MANUAL PART 0 · version 01

0.1 · The operating principle

Human-led. AI-enabled.

The agents do not replace the consultant. They compress the mechanical work: the reading, the reconciling, the modelling, the mapping, the monitoring, so that judgment, context and accountability stay with a person the client can look in the eye.

Two consequences follow, and they are absolute:

- **The agents never close the loop themselves.** Every phase ends at a human gate.
- **The agent advises; Fenwick decides.** Agents surface findings, never conclusions. The conclusion is delivered by Fenwick How.

Force5 is internal

Force5 and the agent identities are internal to the practice. Clients engage the Aperture modules, not the names. This is a hard rule with practical consequences: never put an agent name or the term Force5 in a client-facing deliverable, email, readout, proposal or dashboard; never present an agent as a person to a client; no synthetic voices, no personas in meetings, no agent-authored correspondence sent under a human name. Internal working documents, this manual included, use the names freely. Before anything leaves the practice, check it against 0.9.

Internal	Client-facing
Lena · lena-strategist · SOP 01 and 02	Aperture Insights™
Marcus · marcus-analyst · SOP 03	Aperture Analytics™
Natalie · natalie-explorer · SOP 04	Aperture Intelligence™
Aria · aria-navigator · SOP 05	Aperture Compass™
Thomas · thomas-operator · SOP 06	Aperture Atlas™

0.2 · The evidence standard

Every figure carries an evidence reference. A number without a citation does not render, the render tools enforce it, and routing around them is not permitted.

Class	Meaning	What it must cite
Measured	Sourced from a client system or document.	The ledger ref, and the page, tab or column.
Derived	Computed by a named tool from measured inputs.	The tool and the refs it consumed.
Assumed	A judgment input.	Who set it and on what basis.
Claimed	Asserted in an interview or a questionnaire, not yet corroborated.	Labelled as a claim until the ledger supports it.

A forecast built on assumptions dressed as measurements is the fastest way to lose a client. Keep the classes visibly separate.

Three rules that follow from it

Never compute in prose. Any number that appears in an output comes from a tool Fenwick can re-run. A figure produced by reasoning rather than by a script is not evidence and it does not ship.

Insufficient evidence is a finding. Mark it and move on. Never estimate to fill a gap in a scored rubric. A short, fully-sourced report beats a long, partly-furnished one.

A claim is promoted to measured only when the ledger corroborates it, and the promotion is recorded rather than overwritten. Specific, confident and plausible are not corroboration.

0.3 · Human gates

Every step in every phase carries one of four markings. They are not advisory.

Marking	Who acts	What it means
TOOL	A script	Deterministic and re-runnable. The output is evidence because it can be reproduced.
AGENT	The phase agent	Judgment inside a defined boundary, producing findings with refs.
HYBRID	Agent proposes, Fenwick disposes	Produce the proposal, present the reasoning, then stop. Do not act on the likely answer.
FENWICK	Fenwick only	His step. When an agent reaches one, it stops and states what it needs.

What an agent must never do

Regardless of how obvious the answer looks or how much time it would save:

- Simulate an interview, or invent what a client said

- Draft a readout as though it had been delivered
- Name the binding constraint
- Apply a cost-allocation basis before approval
- Set or adjust rubric weights, scoring criteria or RAG thresholds
- Rank areas or sites before parameters are approved
- Decide the sequence that goes to a board
- Contact a client, or act on a detected signal

0.4 · Rubric and version discipline

The seven lenses, their sub-criteria, the Aperture Score weights and the decision-matrix criteria are Fenwick's intellectual property, not the model's. Agents read them from the rubric files and never invent, adjust or improve a criterion. **Every score stamps the weights and rubric version used.** Re-scoring in a later period uses the stored version, never the current one: comparability across periods is the entire value of a baseline. When a rubric changes, restate history explicitly and label the restatement.

Rubric in force: v1.3.0, ratified by Fenwick How on 10 August 2026. Weights, criteria and anchors are reproduced in full in Part 1.

0.5 · Confidentiality and data posture

Rule	What it means in practice
Least access	Pull the minimum fields required for the defined analysis. Never ingest a full customer table because it was easier than scoping the query.
PII drops at the door	Geocoding runs locally and street-level PII is dropped before analysis. Analysis works on aggregated geography, never individual records.
Suppress below threshold	Where an area contains too few customers to be non-identifying, suppress rather than publish, and show the cell as suppressed rather than as zero or absent.
Data stays local	Client data does not leave the local environment for third-party enrichment without explicit written approval.
Engagements do not mix	Evidence refs are scoped to their engagement. Benchmarking across clients uses only the anonymised aggregate set, never another client's ledger.
Folders are confidential	Every engagement folder is treated as confidential and not for redistribution.

0.6 · The engagement spine

The Method executes as a sequence, not a swarm. Each agent owns one phase, produces a defined output, and passes typed evidence to the next.

The engagement folder

Path	What lives there
engagement.yaml	Client, scope, dates, phase status, rubric version in force, confidentiality terms, and every gate decision with its date and reasoning.
ledger/sources/	Raw client documents, unmodified. Never edited. Working copies go in the phase folder.
ledger/ledger.yaml	Every evidence ref, its source, its date and vintage, and its class.
01-understand/ ... 05-perform/	Phase working folders.
deliverables/	Client-facing outputs only, and only after the pre-delivery check. Nothing is ever delivered from a phase working folder.

The evidence ledger

The ledger is the spine of the whole Method: every figure in every deliverable traces to an entry in it. Each entry carries a stable ref (LED-0142), the source document and location, the date and vintage, the class, and for derived entries the tool that produced it and the refs it consumed.

- **Refs are append-only.** Never renumber, never reuse, never delete. Supersede with a new ref and mark the old one superseded.
- **Promotions are recorded, not overwritten.** A claim that the ledger later corroborates gets a promotion entry.
- **Restatements are findings.** When a client restates a source document, the new version gets a new ref and both stay.

Gate packets

At each phase close, produce a gate packet for Fenwick containing four things, then stop:

#	The packet contains
1	What the phase found, with evidence refs.
2	What it could not establish, and why.
3	The decisions required from him, each stated as a specific question with the options and their consequences.
4	What the next phase will do with each possible answer.

Do not begin the next phase, do not draft its outputs, and do not assume the likely answer in order to keep working. Record the decision in engagement.yaml with the date and the reasoning. Six months later, on a retainer, the reasoning is what makes the Scoreboard interpretable.

Typed handoffs

A handoff states three things: the question, specific and answerable and scoped; the evidence, the refs supporting why the question matters; and the standard, what a good answer looks like and what would make it insufficient. Never hand over an unscoped "look into this". An untyped handoff is how a phase produces forty pages that answer nothing.

A Snapshot is not an engagement

None of the machinery above applies to an Aperture Snapshot: no engagement folder, no engagement.yaml, no evidence ledger, no gate packets, no handoffs. That is a cost decision, not a shortcut. At \$1,950 against a five-hour cap, standing up a folder and registering every source would consume the cap before any analysis began, and the client is not paying for a ledger they will never see.

What a Snapshot keeps: every figure still traces to a source, in the report's own "what we read" section; the rubric version is still stamped, so a later X-Ray or Atlas re-score stays comparable; and the single human gate stays, because nothing goes to a client unread.

If the client upgrades to a Business X-Ray, that opens an engagement, and the Snapshot's sources are registered into the new ledger at that point.

0.7 · Deliverable standards

Every deliverable carries the engagement and date; the module name in client-facing form; the rubric or criteria version where scored; an assumptions section where modelled; the confidence and vintage of external data where used; and the confidentiality footer.

Write for a hostile reader who could open the workbook and follow every number to its source. Build for that reader and the friendly ones are covered automatically.

Charts and maps

- **Every chart title states the finding, not the variable.** "Top 3 customers are 61% of revenue", never "Revenue by customer". A chart titled with its axes makes the reader do the analysis. If you cannot write that sentence, you do not yet know what the chart is for and it should not ship.
- One idea per chart. A second idea is a second chart.
- Make the comparison for the reader: sort by value, baseline against the prior period or the benchmark, highlight the series the finding is about, and annotate the moment that matters directly on the plot.
- Every map carries a legend, a scale bar, a north arrow, the data vintage, the projection and the source. Observed, modelled and estimated never share a symbology. Choropleths map rates, not counts.

Brand system

Black #000000, white #FFFFFF, and maroon #500000 as a sparing strategic accent, well under five percent of any page. Graphik Semibold for primary headlines and for "APERTURE", Graphik Regular or Light for "METHOD", supporting headings and body, with a legally embeddable web alternative where Graphik cannot be embedded. Never more than three type sizes on a page. Wide margins: whitespace is the practice's most consistent visual signal. Tabular figures in every financial table, never proportional.

Voice is accessible premium: confident, analytical, clear, practical, human, executive. Premium restrained visuals, plain and warm owner-friendly words. Avoid marketing-agency language, startup clichés, generic buzzwords, unsupported performance claims and overly technical language without business meaning.

Open discrepancy to settle

The brand reference inside the Force5 plugin (references/brand-system.md) describes a different system: crimson #C1444F on an ink #0B0709 ground, with typography marked TODO. It is labelled "inferred from the Force5 deck, Fenwick to ratify" and it has never been reconciled with the ratified black / white / maroon #500000 Graphik system above.

Until it is, an agent reading that file will produce off-brand output. The system in this section governs. Appendix F carries this as an open item.

Punctuation, and the AI tell

Em dashes read as AI-generated. This is a bias, not a ban: write the sentence without one first, and reach for one only when no other mark does the job. Where the dash was introducing a list,

definition or explanation, use a colon; where it set off an aside, use a comma pair or parentheses; where it joined two independent clauses, use a full stop or a semicolon; where it was doing nothing, delete it. En dashes in numeric ranges stay, as does the middot separator, which is the house separator for label · label · label.

0.8 · No borrowed frameworks

No SWOT grid and no five-forces diagram appears in any client artefact. Both produce unranked, unfalsifiable assertions, which is the opposite of what these instruments are for, and every competitor the client has ever hired has handed them one. **Where a force matters, it appears as a scored criterion with a figure and an evidence ref, or it does not appear.**

The classical question	Where it actually lives
Supplier power	ops-06, supplier and input dependency: the share of cost or volume through the largest supplier, and what switching would cost.
Buyer power	cus-03, customer value concentration. Do not score the same exposure twice.
Durability of advantage	mkt-05, advantage durability: what a determined competitor would find hardest to copy, and how long it would take.
Rivalry, substitution, entry	mkt-01 to mkt-03. Use the five-forces frame privately if it helps you cover the ground; never put it in a deliverable.

0.9 · The pre-delivery check

Run before any file leaves the practice. No exceptions for drafts: a draft sent to a client is a delivery. If any blocking item fails, the file does not go out. Agents never waive a blocking item, and never remove an internal name in order to pass this check without saying it was there.

1 · Internal identity scrub · BLOCKING

Search for every one of: Lena, Marcus, Natalie, Aria, Thomas, Force5, "the agent", "the model", the five agent slugs, and SOP 01 through SOP 06. Any hit blocks delivery. Check document metadata, file properties, sheet tab names, chart titles, image alt text, speaker notes, tracked changes, comments and the filename itself. Those are where internal names survive a body-text scrub.

2 · Evidence integrity · BLOCKING

☐ Check
☐ Every figure carries an evidence reference
☐ No number was computed in prose; each traces to a re-runnable tool
☐ Financial models reconcile to the P&L to the dollar
☐ recalc.py passed clean (Phase 02 outputs)
☐ feed.validate() passed (Phase 05 outputs)
☐ Every input labelled measured, derived, assumed or claimed
☐ Interview and questionnaire claims are labelled as claims, not stated as fact
☐ Insufficiently-evidenced items are marked, not estimated

3 · Gate compliance · BLOCKING

☐ Check
☐ Cost-allocation basis approved before figures were computed on it
☐ Trade-area parameters and scoring criteria approved before any ranking
☐ Rubric weights and matrix criteria agreed before scoring, not after
☐ The binding constraint was named by Fenwick, not by an agent
☐ The board sequence is Fenwick's, not the agent's proposal shipped as-is
☐ Nothing states a conclusion the client should reach
☐ Fenwick has read it, edited it and signed it

4 · Confidentiality, presentation and final

☐ Check
☐ No street-level PII anywhere, including appendices and source-data tabs
☐ Cells below the disclosure threshold are suppressed and shown as suppressed
☐ No other client's data, name or benchmark is identifiable

☐	Check
☐	Confidentiality footer present
☐	Hidden rows, columns, sheets and filtered-out data checked, and deleted rather than hidden
☐	House brand system followed; trademarks marked on first prominent use
☐	Rubric or criteria version stated where scored; assumptions visible where modelled
☐	External data carries source and vintage; maps carry legend, scale, north arrow, vintage, projection
☐	Capability language reads as a standard, never as a personal credential or a licensed opinion
☐	Spell-check, and the client's legal entity name checked
☐	Filename follows convention and contains no internal identifiers
☐	Correct version; superseded drafts removed; delivered from deliverables/

Where a licensed opinion is genuinely required, in audit, tax or law, say so and route it to a licensed professional. Agent capability descriptions state the standard the agent is built to. They are benchmarks, never credentials held by a person.

OVERVIEW

How the segments fit together

The Aperture Method moves through five phases, from understanding where a business stands to keeping its plan alive. There are two ways in, depending on size.

No.	Segment	The question	Investment
01	Aperture Snapshot™	What is happening in my business, and where should I look first?	\$1,950
02	Business X-Ray™	What is happening?	\$4,500
03	Aperture Analytics™	Why is it happening?	From \$8,500
04	Aperture Intelligence™	What opportunities does it reveal?	From \$8,500 · \$18,000
05	Aperture Compass™	Where should we go next?	From \$8,500
06	Aperture Atlas™	How do we sustain and improve results?	\$900/mo · \$1,500/mo · \$2,500/mo · Terms

Where to start

- **Under \$5M in revenue:** the Aperture Snapshot™ (\$1,950). It counts in full toward the Business X-Ray if you continue within 60 days.
- **\$5M to \$20M:** the Business X-Ray™ (\$4,500). It counts in full toward the Method if you continue within 60 days.
- **The Full Method** (X-Ray, Analytics, Intelligence and Compass) is \$25,000. Atlas follows as the ongoing monthly segment.

How we handle your data

- We ask for the least access that does the job: read-only, never edit rights.
- We do not need customer names, contact details or private notes. An ID and a ZIP code are enough.
- Every finding is tied to the evidence behind it, and where evidence is thin we say so rather than guess.

01

SEGMENT 01 OF 06 · UNDERSTAND

Aperture Snapshot™

A written first read, for businesses under \$5M

Question it answers	What is happening in my business, and where should I look first?
You receive	A written Snapshot report with a provisional Aperture Score™
Timing	Starts the day your last required item arrives. Written report only, no meetings.
Investment	\$1,950: Fixed fee. It counts in full toward the Business X-Ray if you continue within 60 days.
Your time	About twenty minutes for the questionnaire, plus the time to run the exports.
Runs it	lena-strategist, under SOP 01 and SOP 02
Bench	Aperture Insights Workbench v0.1
Manual source	Part 1 · version 01 · In force · bench partially built

VERSION 2.0 ISSUED 25 September 2026 · 2:48 PM CT SEGMENT 01 of 06

SEGMENT 01 · APERTURE SNAPSHOT™**What it is**

A written, document-based read of your business across the same seven areas we look at for every client: finance, operations, customers and retention, market and competition, leadership and team, processes, and technology and data.

It is built only from documents you already have and a short written questionnaire. There are no interviews and no meetings, which is what keeps it affordable.

It covers the same ground as a Business X-Ray. The difference is depth: in a Snapshot nothing you send is tested in conversation or checked against what the documents leave out. That is why its score is labeled provisional, and the report says so plainly.

What it does

- Scores all seven areas, and tells you how confident each score is
- Names the one area most likely holding profit back
- Gives you a provisional Aperture Score™ you can compare against later
- Lists what a Snapshot could not settle, and what it would take to settle it

What we need from you

Required. Without these the work cannot start.

Item	What it should contain	Why it matters
Profit & loss statements	Three full years plus year to date. Monthly if your system produces it. Cost of delivery shown separately from overhead where your accounts do that.	The starting point for every financial reading.
Balance sheets	The last four period ends, all at comparable dates.	Four dates show a pattern; one shows a single moment.
Transaction export	One row per sale for the last 24 months: date, customer ID, product or service, location, revenue, direct cost where you track it, and who delivered it. Exported straight from your system as CSV or Excel.	The single most valuable file you can send. It is where new findings usually come from.
Completed questionnaire	The sections listed below. Leave a question blank rather than guess; a blank tells us something too.	It answers what documents cannot: why customers leave, who you lose to, what breaks under pressure.

Strongly helpful. Each one sharpens a specific part of the reading.

Item	What it adds
Revenue by product or service, and by location	Shows which parts of the business carry the profit, not just which are busiest.
Customer list	Customer ID, ZIP code, how they found you, and first purchase date. Please do not send names or contact details.
Marketing spend by channel and month	Shows what a new customer costs you, and whether the channels you spend most on bring customers who stay.
Price list or service catalog	Current prices and, where you know it, the direct cost to deliver one unit.
Payroll or headcount by role, or an org chart	The shape of the team and where the pressure points are.
Staff skills or license register	Whether qualified people are what limits how much you can sell.
Debt and lease schedule	What the business is committed to, and how that affects cash.
Accountant's report and business tax returns	How much weight the figures can carry. Audited or internally prepared are both fine.
Cash flow statements	Whether profit is turning into cash.

Questionnaire sections to complete · 15 sections, 73 questions

Section	Questions
The business	14
The one question	3
Where you are headed	3
How decisions get made	4
Systems we would read from	6
What data exists	3
People we should talk to	2
Lens 1 · Finance	7
Lens 2 · Operations	6
Lens 3 · Customers & Retention	5
Lens 4 · Market & Competition	5
Lens 5 · Leadership & Team	6
Lens 6 · Processes	2
Lens 7 · Technology & Data	3
The constraint, and what "solved" means	4

How to send it

Send native exports	CSV or Excel straight from your system. A PDF of a spreadsheet has to be retyped; a PDF of a statement is fine.
One row per transaction	Monthly summaries hide the patterns the analysis is looking for.
Keep the same customer ID in every file	That is how sales and customers are linked. Please don't re-number them between files.
Leave out names and private notes	No customer names, contact details, or clinical or account notes. An ID and a ZIP code are all we need.
One delivery, one folder	Plain file names, all at once. Work starts when the last required item arrives.

If something is missing

- When your files arrive we check them before any analysis starts. If something important is missing or does not add up, we ask once, with a specific list and a date.
- If a gap remains, the report says so and tells you which part of the reading it affects.
- If there is not enough evidence for a reliable overall score, we withhold the single number rather than guess. You still receive every scored area and the named constraint.
- If it becomes clear the business needs more than a Snapshot can give, we will tell you and recommend a Business X-Ray rather than stretch the Snapshot.

OPERATING DETAIL · INTERNAL

How a Snapshot is run

Drawn from the Master Operating Manual, Edition 01, Part 1, version 01. The Manual governs; section numbers below are the Manual's own.

1.2 · Two depths

Aperture Insights is sold at two depths. Same seven lenses, same rubric, same weights, both times. What changes is how far the reading gets tested. Establish which one has been sold before doing anything, because delivering the wrong one is expensive in both directions.

	Aperture Snapshot™ · \$1,950	Business X-Ray™ · \$4,500
Sold to	Under \$5M in revenue, through partners and referral	The \$5M to \$20M core
Source	Submitted documents plus the written intake questionnaire	Documents and questionnaire, plus management interviews
All seven lenses scored	Yes	Yes
Management interviews	Skipped entirely	Required
Evidence ledger	Not produced	Required
Triangulation and disconfirming evidence	Not attempted	Required
Constraint output	The area most likely to be holding profit back, named as one area	Ranked candidates, each with what argues against it
Aperture Score	Provisional, stamped as scored from evidence as submitted	Baseline, stamped with weights and rubric version
Readout	Written only, no live readout	Live readout
Fenwick's time	5 hours, hard cap	~12 hours

The Snapshot is not fewer lenses. If it is ever cut to three or four, the score stops being comparable to a later X-Ray and to the Atlas re-score, and the whole upgrade argument collapses. Breadth is identical. Depth of testing is the entire difference, and that is what the client is told they are buying when they pay the \$2,550 to go further.

Snapshot mode, step by step: Step 1 becomes a document and questionnaire request with no kickoff call · Steps 2, 3 and 5 run as written · Step 4 is skipped · Step 6 names one area rather than

ranking candidates, and does not attempt disconfirming evidence · Step 7 produces a shorter report · Step 8 is skipped.

Two things every Snapshot deliverable must say, in the client's own language and not in a footnote

The score is provisional. It was scored from the evidence as submitted. Nothing in it has been tested against an interview, checked for what the documents leave out, or argued against.

What a Snapshot cannot settle, and what it would take to settle it. At this depth the "where to dig" register is doing double duty: it is both the honest limit of the work and the reason the X-Ray exists.

The failure mode to watch here is over-building, again, and it is worse at this depth. The five-hour cap is not a target, it is the condition on which the product exists: at twelve hours a \$1,950 Snapshot earns less per hour than almost anything else available. If a Snapshot is heading past the cap, the honest move is to say the business needs an X-Ray and stop, not to absorb the hours.

Where the rest lives

A Snapshot runs the Phase 01 process at its lighter depth. Operating standard, rubric, data required, the receipt and sufficiency check, deliverables and quality checks are in Segment 02, and they govern both depths.

02

SEGMENT 02 OF 06 · UNDERSTAND

Business X-Ray™

Aperture Insights™ · the full diagnostic, for businesses of \$5M to \$20M

Question it answers	What is happening?
You receive	The Business X-Ray™ report, a baseline Aperture Score™, and a live readout
Timing	2 to 3 weeks from the kickoff call
Investment	\$4,500: Fixed fee. It counts in full toward the Method if you continue within 60 days.
Your time	The questionnaire, a kickoff call, interviews with you and the people you name, and the readout.
Runs it	lena-strategist, under SOP 01 and SOP 02
Bench	Aperture Insights Workbench v0.1
Manual source	Part 1 · version 01 · In force · bench partially built

VERSION 2.0 ISSUED 25 September 2026 · 2:48 PM CT SEGMENT 02 of 06

SEGMENT 02 · BUSINESS X-RAY™**What it is**

The executive diagnostic that establishes where your business stands today. It reads the same seven areas as a Snapshot, at full depth.

Alongside your documents and questionnaire we interview you and your key people and read your core systems directly, with read-only access. What we are told is tested against what the evidence shows, and where the two disagree we say so.

It is deliberately broad rather than deep. It tells you where the real problem most likely sits and where deeper work would pay off. The deeper work itself is what the later segments are for.

What it does

- Scores all seven areas against evidence, with a confidence rating on each
- Ranks what is most likely holding profit back, with the evidence for and against each candidate
- Sets your baseline Aperture Score™, the benchmark later progress is measured against
- Shows where deeper analysis is warranted, and where it is not
- Ends with a live readout, not just a document

What we need from you

Required. Without these the work cannot start.

Item	What it should contain	Why it matters
Profit & loss statements	Three full years plus year to date. Monthly if your system produces it. Cost of delivery shown separately from overhead where your accounts do that.	The starting point for every financial reading.
Balance sheets	The last four period ends, all at comparable dates.	Four dates show a pattern; one shows a single moment.
Transaction export	One row per sale for the last 24 months: date, customer ID, product or service, location, revenue, direct cost where you track it, and who delivered it. Exported straight from your system as CSV or Excel.	The single most valuable file you can send. It is where new findings usually come from.
Completed questionnaire	The sections listed below. Leave a question blank rather than guess; a blank tells us something too.	It answers what documents cannot: why customers leave, who you lose to, what breaks under pressure.

Strongly helpful. Each one sharpens a specific part of the reading.

Item	What it adds
Revenue by product or service, and by location	Shows which parts of the business carry the profit, not just which are busiest.
Customer list	Customer ID, ZIP code, how they found you, and first purchase date. Please do not send names or contact details.
Marketing spend by channel and month	Shows what a new customer costs you, and whether the channels you spend most on bring customers who stay.
Price list or service catalog	Current prices and, where you know it, the direct cost to deliver one unit.
Payroll or headcount by role, or an org chart	The shape of the team and where the pressure points are.
Staff skills or license register	Whether qualified people are what limits how much you can sell.
Debt and lease schedule	What the business is committed to, and how that affects cash.
Accountant's report and business tax returns	How much weight the figures can carry. Audited or internally prepared are both fine.
Cash flow statements	Whether profit is turning into cash.

People and access

People to interview	You, plus the managers or advisors who see the business from a different angle. The questionnaire asks you to name them.
Read-only system access	Your accounting system and the one or two other systems that run the business, such as point of sale, scheduling or CRM. Read-only, never edit rights.

Questionnaire sections to complete · 15 sections, 73 questions

Section	Questions
The business	14
The one question	3
Where you are headed	3
How decisions get made	4
Systems we would read from	6
What data exists	3
People we should talk to	2
Lens 1 · Finance	7
Lens 2 · Operations	6
Lens 3 · Customers & Retention	5

Section	Questions
Lens 4 · Market & Competition	5
Lens 5 · Leadership & Team	6
Lens 6 · Processes	2
Lens 7 · Technology & Data	3
The constraint, and what “solved” means	4

How to send it

Send native exports	CSV or Excel straight from your system. A PDF of a spreadsheet has to be retyped; a PDF of a statement is fine.
One row per transaction	Monthly summaries hide the patterns the analysis is looking for.
Keep the same customer ID in every file	That is how sales and customers are linked. Please don't re-number them between files.
Leave out names and private notes	No customer names, contact details, or clinical or account notes. An ID and a ZIP code are all we need.
One delivery, one folder	Plain file names, all at once. Work starts when the last required item arrives.

If something is missing

- We check what arrives before analysis starts. Gaps are raised at the kickoff call, where most of them can be filled in conversation.
- Anything that still cannot be evidenced is written down as a finding, with the confidence it costs, rather than guessed at.

OPERATING DETAIL · INTERNAL

How a Business X-Ray is run

Drawn from the Master Operating Manual, Edition 01, Part 1, version 01. The Manual governs; section numbers below are the Manual's own.

1.1 · What it is, and what it is not

The purpose of this phase is to establish a defensible picture of where the business stands today. It is not to solve anything. Solving happens in Phase 04, and a Phase 01 that drifts into recommendations has skipped the work that makes the recommendations correct.

Insights is the client's way in and it is priced as one fixed, modest fee. It is deliberately **broad, not deep**: a seven-lens read of where the business stands, from documents the client already has plus a short interview. Meaningful and honestly scored, but not the full analysis. It does not do the heavy work. It finds where the heavy work is warranted, and names it.

Insights DOES	Insights does NOT · this is bought in Phases 02 to 04
Score all seven lenses from the evidence provided	Build financial models, cohorts, regressions or forecasts (Phase 02)
Name the likely binding constraint, with the evidence to hand	Run trade-area, drive-time, segmentation or site analysis (Phase 03)
Produce the Aperture Score as an honest baseline	Score, sequence or stress-test options (Phase 04)
Register strengths, risks and gaps	Exhaustively triangulate every management claim

The two failure modes, and they pull in opposite directions

Over-building. If this phase starts producing analysis that belongs to Phase 02, 03 or 04, stop. The client has not paid for it, the margin is gone, and the components that follow have nothing left to sell. When a lens looks interesting, write down why it warrants a deeper look and move on. That sentence is the product.

Under-calling. Breadth is not vagueness. Every lens still gets a score, a reason and a confidence, and the constraint candidates are still ranked. Shallow scope, full rigour of judgement.

1.3 · Operating standard

Rule	Why it exists
A score without an evidence ref is refused	Enforced by the scoring tool, not by discipline. A criterion with no ref does not count as scored.

Rule	Why it exists
Insufficient is excluded, never scored zero	Zero is a judgment that the business is failing at something. Absence of evidence is not that judgment.
Weight is never redistributed	A lens with no scored criteria drops out of the composite and its weight leaves the denominator. Redistributing it would hide the gap inside a score that looks complete.
Below the coverage floor the composite is withheld	Not caveated. A score built on a third of the evidence is not a cautious score, it is a fabricated one.
Findings, never conclusions	The named constraint is Fenwick's call. The agent proposes ranked candidates and stops.
A read, not an excavation	If a section is running long because the analysis is getting deep, it belongs to a later phase.

1.4 · Methodology

The seven lenses and their weights

Lens	Weight	Criteria	Why it carries that weight
Finance	20%	5	Highest weight. Cash and unit economics are what actually kill businesses in this segment, and it is the most reliably evidenced lens: the P&L exists whether or not anything else does.
Operations	18%	6	In place-based businesses the constraint most often binds physically: capacity, throughput, a bottleneck someone can point at.
Customers & Retention	18%	4	Retention is the highest-leverage and most commonly mismeasured area in owner-run businesses. Owners routinely believe their repeat rate is far higher than it is.
Market & Competition	14%	5	Sets the ceiling on everything else. A well-run business in a shrinking market has a strategy problem no operational fix reaches.
Leadership & Team	12%	5	Key-person risk is existential here and is the single most common reason an owner-run business cannot be sold. Weighted above Processes for that reason.
Processes	10%	3	Lower weight deliberately. Weak process is usually a symptom of Operations or Leadership rather than a root constraint.
Technology & Data	8%	3	Lowest weight, deliberately. Technology is almost never the binding constraint in this segment: it is an enabler. Weighting it higher produces 'buy software' recommendations, which is the failure mode of this entire industry.

Thirty-one sub-criteria in total. Sub-criteria are equally weighted within each lens unless stated otherwise.

The Aperture Score, and exactly how it is computed

The composite is deterministic and it stamps the weights and rubric version. Never hand-adjust it: if the result looks wrong, the input scores or the rubric are the place to look. Five rules, all enforced in both the browser scorer and the Python engine:

#	Rule
1	A score with no evidence ref is refused outright.
2	A criterion marked insufficient is excluded from the mean. It is never scored zero.
3	A lens score is the mean of its sufficient criteria.
4	The composite is the weight-normalised mean over the lenses that have a score. Weight is not redistributed to the lenses that happen to have evidence; the composite renormalises across the weight actually covered and the coverage figure is reported alongside, so the gap is visible rather than absorbed.
5	Below the coverage floor the composite is withheld, not published with a caveat.

Band	Meaning
0-20	Critical: actively threatens the business
21-40	Weak: a material drag; likely constraint territory
41-60	Adequate: functioning, not an advantage
61-80	Strong: a genuine asset
81-100	Distinctive: a durable competitive advantage

Band keys are integer-adjacent but a lens score is a mean to one decimal, so match on the lower bound rather than testing both ends. A score of 60.5 falls in no band if you test both.

The coverage floor, and the proposed corroboration floor

Test	Threshold	What happens below it
Coverage floor	22 of 31	70% of sub-criteria scored with evidence. Below it the composite is withheld. Ratified in rubric v1.3.0.
Corroboration floor	12 of 31	At least 12 criteria on measured or derived evidence, of which at least 6 in Finance and Customers & Retention. PROPOSED, not yet in the rubric. See 1.7 and Appendix F.
Lens survival	1 of n	A lens with no scored criteria drops out and its weight leaves the denominator.
Band ceiling	80	A criterion resting on claimed evidence alone cannot reach the Distinctive band. PROPOSED alongside the corroboration floor.

Worked reference: the Lumina FY2023 Snapshot scored six lenses over 86 points of weight, Market & Competition having dropped out. Its provisional score of 57 was $(71 \times 20 + 66 \times 18 + 47 \times 18 + 52 \times 12 + 44 \times 10 + 45 \times 8) \div 86$.

Structural pressure · internal checklist, never a client-facing artefact

Two sub-criteria carry the questions classical strategy analysis asks about a business's position, restated so they have to be evidenced rather than asserted. ops-06 is supplier power expressed as a number. mkt-05 is durability of advantage, with the honest follow-up: where the answer is "not much, and quickly", whether building one is a costed plan with an owner and a date. Buyer power is already carried by cus-03, so do not score the same exposure twice. Rivalry, substitution and threat of entry inform mkt-01 to mkt-03. Use the frame privately if it helps you cover the ground; see 0.8 for why it never appears in a deliverable.

1.5 · The rubric in full

Rubric v1.3.0, ratified by Fenwick How on 10 August 2026. Reproduced here in full so the instrument and the manual cannot drift. Each criterion carries its definition, the evidence classes it requires, and the intake fields the rubric resolves against. **Intake refs are stable field names** from the intake system, never positional question numbers: positional refs silently rot the moment a question is inserted, and `validate_rubric.py` fails the build if a name stops resolving.

The anchors describe what a 20, a 50 and an 80 look like. They are the scoring instrument, not illustrations: score against the anchor, then write the reason.

Finance · 20% · 5 sub-criteria

ID	Criterion	What it assesses	Evidence	Intake refs
01	Margin visibility	Whether the business knows where it makes and loses money.	Measured, Derived	ins_fin1, ins_fin6, an_num2, an_num3
02	Cash position and cycle	Whether cash is predictable and adequate through the year, and whether the operating cycle that drives it is measured in days.	Measured	ins_fin2, ins_fin4, ins_fin7
03	Unit economics	Whether the economics of one unit, job or customer are understood and positive.	Measured, Derived	ins_fin5, an_price1, an_price3, an_price4, an_dec3
04	Cost trend control	Whether cost increases are understood and acted on.	Measured, Derived	ins_fin3
05	Books reliability	Whether the accounts can be relied on for decisions.	Measured, Claimed	an_num4, data_missing, data_basis

ID	Anchor 20	Anchor 50	Anchor 80
01	Profit is known only at whole-business level. The owner is guessing about product or location profitability.	Gross margin known by major line; overhead unallocated. The owner has a view but low confidence in it.	Contribution and fully-absorbed profit known by product/service and location, on a documented allocation basis.
02	Cash is unpredictably tight. The owner discovers shortfalls rather than forecasting them. The cash conversion cycle has never been calculated and the inputs to calculate it are not to hand.	Seasonal tight points are known from experience. No rolling forecast exists, and the cash cycle is held as an impression, “customers pay slow”, rather than as a number of days.	A rolling cash forecast exists. The cash conversion cycle is calculated from the accounts (DSO + DIO – DPO), tracked over time, and its three components are managed separately.
03	Unit economics unknown. Pricing set by feel or by reference to a competitor.	Rough unit margin known for the main lines, untested against fully loaded cost.	Unit economics known and tested, including cost to serve and acquisition cost.

ID	Anchor 20	Anchor 50	Anchor 80
04	Costs have grown materially and the owner cannot explain why.	The main areas of cost growth are identified; no systematic review.	Cost trends reviewed on a cadence with named owners; drivers understood and challenged.
05	Books closed late or not at all. The owner does not trust them, and is right not to.	Closed monthly with some lag. Broadly trusted, with known issues that get worked around.	Closed promptly, reconciled, and actually used by the owner to make decisions.

Operations · 18% · 6 sub-criteria

ID	Criterion	What it assesses	Evidence	Intake refs
01	Capacity headroom	Whether the business can absorb demand growth without a step change in cost.	Measured, Derived	ins_ops2, ins_ops3, ins_ops4
02	Bottleneck identification	Whether the constraining step is known, measured and actively managed.	Measured, Claimed	ins_ops1, ins_ops2
03	Throughput reliability	Whether work completes predictably, on time and to standard.	Measured, Claimed	ins_ops1, ins_ops4, an_vol1
04	Demand and capacity matching	Whether capacity is planned against demand, including seasonality.	Measured, Derived	ins_ops3, ins_ops4, an_vol3
05	Qualified capacity	Whether the supply of licensed, certified or specifically-trained people is the binding constraint on billable volume, and how long that constraint takes to relieve.	Claimed, Measured	ins_ops5, ins_ops4
06	Supplier and input dependency	Whether the business depends on a supplier, subcontractor or platform it could not replace quickly, and whether that exposure is sized rather than assumed.	Claimed, Measured	ins_ops6

ID	Anchor 20	Anchor 50	Anchor 80
01	At or over capacity now. Demand is being turned away or service quality is degrading. No headroom in the binding resource.	Roughly 10-25% headroom in the binding resource. Peaks already strain it.	40%+ headroom, and the next step-change in capacity is identified and costed.
02	Nobody can name where work backs up, or different people name different places. Treat disagreement here as a finding in its own right.	The bottleneck is named consistently but not measured. Management is reactive.	Named, measured, and explicitly protected and scheduled around.

ID	Anchor 20	Anchor 50	Anchor 80
03	Frequent rework, missed dates or expediting. On-time completion is not measured at all.	Generally reliable, but exceptions are absorbed by individual effort rather than by system. Measured informally.	On-time and quality both measured, variance understood, exceptions rare and root-caused.
04	Capacity is fixed and demand swings are absorbed by overtime or lost sales.	Seasonality is known from experience and informally planned for.	Capacity is flexed against a demand forecast; the plan is written and reviewed on a cadence.
05	Billable volume is capped by qualified headcount, the cap is not measured, and the lead time to qualify another person is long or unknown.	The constraint is understood and roughly measured, but hiring or qualification is reactive: it starts when someone leaves or demand already exceeds supply.	Qualified capacity is measured against forecast demand, the qualification pipeline is deliberate, and lead time to add a qualified person is known and planned for.
06	A single supplier, subcontractor or platform carries a large share of cost or volume, the share is not measured, and no alternative has been identified, priced or tested.	The main dependencies are known and roughly sized. Alternatives exist in principle but have not been priced or trialled, and terms are accepted as given.	Concentration of spend and volume is measured by supplier, switching cost and lead time are known, and either a second source is live or the exposure is a deliberate, priced decision.

Customers & Retention · 18% · 4 sub-criteria

ID	Criterion	What it assesses	Evidence	Intake refs
01	Acquisition understanding	Whether it is known where customers come from and what each channel costs.	Measured, Derived	ins_cust1, ins_cust5, int_grow2, int_grow3
02	Retention measurement	Whether repeat behaviour is measured rather than assumed.	Measured	ins_cust2
03	Customer value concentration	Whether value concentration is known and deliberately managed.	Measured, Derived	ins_cust3, b_buyers, data_have
04	Churn causes	Whether the reasons customers leave are known and addressed.	Measured, Claimed	ins_cust4

ID	Anchor 20	Anchor 50	Anchor 80
01	Sources unknown or purely anecdotal. No spend attribution.	Sources ranked by impression. Marketing spend known in total but not by channel.	Sources measured, spend attributed by channel, acquisition cost known per channel.
02	Repeat rate unknown. The owner believes customers return but cannot show it. Expect the belief to be optimistic: test it early.	Repeat rate estimable from the system but not tracked over time or by cohort.	Cohort retention tracked over time; changes are noticed and investigated.

ID	Anchor 20	Anchor 50	Anchor 80
03	No view of which customers carry the revenue. Concentration risk unassessed.	Top customers known by name. No systematic segmentation or lifetime value.	Segmented with LTV. Concentration quantified and managed as a deliberate risk position.
04	Departures go unnoticed. No exit information of any kind.	Reasons known anecdotally. Nothing systematic.	Churn reasons captured systematically and fed back into the offer or operations.

Market & Competition · 14% · 5 sub-criteria

ID	Criterion	What it assesses	Evidence	Intake refs
01	Competitive position	Whether there is a clear, defensible reason customers choose this business.	Claimed, Measured	ins_mkt1, int_comp3
02	Competitive awareness	Whether the business knows who it loses to and why.	Measured, Claimed	ins_mkt1, ins_mkt4, int_comp1, int_comp2
03	Demand outlook	Whether the direction of the market is known and reflected in the plan.	Measured, Derived	ins_mkt2
04	Geographic and segment position	Whether the served market, its penetration and its whitespace are understood.	Measured, Derived	ins_mkt3, int_geo2, int_geo3, int_grow1
05	Advantage durability	Whether the reason customers choose this business would survive a determined competitor setting out to copy it, and, where it would not, whether building one is a costed plan with an owner rather than an aspiration.	Claimed, Measured	ins_mkt5, int_comp3

ID	Anchor 20	Anchor 50	Anchor 80
01	No articulable difference. Competing on price by default.	A real difference exists but is not consistently communicated or defended.	A clear, evidenced advantage that customers can name and competitors struggle to copy.
02	Competitors unknown or unexamined. Losses unexplained.	Competitors named. Loss reasons anecdotal.	Losses tracked with reasons. Competitor moves monitored and responded to deliberately.
03	No view of market direction. Continuity assumed.	Directional view held from experience, not evidenced externally.	Market trend evidenced from external data, and the plan visibly reflects it.
04	No view of catchment, penetration or whitespace.	Catchment known by impression. Strong and weak areas sensed but never measured.	Catchment and penetration measured. Whitespace quantified and sized.

ID	Anchor 20	Anchor 50	Anchor 80
05	Nothing here that a competent competitor could not reproduce within a season, and no plan to build anything that would hold. Competing on price and effort.	A real advantage exists (location, relationships, a licence, accumulated data, switching cost) but it is undefended, inconsistently communicated, and how long it would take to copy has never been considered.	The advantage is named, its replication cost and lead time are estimated, and it is deliberately reinforced. Where a gap remains, a costed plan with an owner and a date exists.

Leadership & Team · 12% · 5 sub-criteria

ID	Criterion	What it assesses	Evidence	Intake refs
01	Key-person risk	What happens when the owner or a key person is absent.	Claimed, Measured	ins_lead1, ins_lead4
02	Role clarity and capability	Whether the right roles are defined and filled with capable people.	Claimed, Measured	ins_lead2
03	Decision-making and execution	Whether decisions get made and executed, or stall.	Claimed, Measured	d_approve, ins_lead3, comp_done1, comp_done2, comp_done3
04	Succession and capability transfer	Whether the ability to do critical work exists in more than one person, and whether it is written down well enough to transfer. Assesses the system that holds capability, never an individual's competence.	Claimed, Measured	ins_lead5, ins_proc1, ins_proc2
05	Structure and span of control	Whether the shape of the organisation fits its size: how much reports directly to the owner, and whether anyone else carries real management load.	Claimed, Measured	ins_lead6, b_employees

ID	Anchor 20	Anchor 50	Anchor 80
01	The business stops or degrades badly within days. The owner cannot take a month off. This is the most common reason an owner-run business cannot be sold.	Runs for weeks, but decisions queue and some functions have no backup. Single-person dependencies are known but not quantified.	Runs a month without the owner. Decision rights are delegated and documented, and the revenue exposed to any single-person dependency is quantified and being deliberately reduced.
02	Roles undefined or badly mismatched to the people in them. Critical gaps unfilled and unacknowledged.	Roles broadly defined. One or two critical gaps acknowledged but not being actively closed.	Roles defined against the plan, filled with capable people, gaps identified and actively recruited.
03	Initiatives routinely stall. Nobody can say why, or who decides.	Decisions get made but slowly. Execution depends on the owner personally pushing.	Clear decision rights, named delivery owners, and evidence of initiatives actually completed.

ID	Anchor 20	Anchor 50	Anchor 80
04	Critical work lives in one head. Nothing is documented, nobody is cross-trained, and a resignation would remove capability the business cannot rebuild quickly.	Some cross-cover exists, unevenly. Documentation is partial or out of date, so transfer depends on the departing person's goodwill and time.	Every critical role has a named cover and current documentation good enough to work from. Capability transfer has actually been tested, not merely assumed.
05	Everyone reports to the owner. There is no management layer, so the owner is the constraint on every decision and the business cannot grow past their personal bandwidth.	A partial layer exists, but spans are uneven and the owner still carries direct reports who should sit under someone else.	Spans are deliberate and appropriate to size. Management load is genuinely distributed, and the owner's direct reports are the ones that should be.

Processes · 10% · 3 sub-criteria

ID	Criterion	What it assesses	Evidence	Intake refs
01	Process documentation	Whether the work survives the person currently doing it.	Measured, Claimed	ins_proc2, ins_lead1
02	Manual and fragile steps	How much of the work depends on memory, a spreadsheet or paper.	Measured, Claimed	ins_proc2
03	Improvement capability	Whether broken processes actually get fixed.	Measured, Claimed	ins_proc1, ins_lead3, comp_done1

ID	Anchor 20	Anchor 50	Anchor 80
01	Nothing documented. Process lives entirely in individuals' heads.	Key processes documented but out of date, or documented and not followed.	Documented, current, and actually followed. New staff can be trained from them.
02	Core commercial processes run on memory, a spreadsheet or paper. Single points of failure in the critical path.	Manual steps remain in non-critical paths. Known and tolerated.	Critical paths systematised. Remaining manual steps are deliberate choices and documented.
03	Known problems have persisted for years. No mechanism exists to fix them.	Problems get fixed when they become acute, by whoever has time.	A named owner, a cadence, and evidence of processes measurably improved.

Technology & Data · 8% · 3 sub-criteria

ID	Criterion	What it assesses	Evidence	Intake refs
01	System coverage	Whether the systems capture what the business needs to run.	Measured	s_accounting, s_pos, s_crm, s_marketing, s_other, atl_wire2

ID	Criterion	What it assesses	Evidence	Intake refs
02	Data trustworthiness	Whether the numbers the business relies on are trusted, and rightly so.	Measured, Claimed	ins_tech2, an_num4, data_missing, data_basis
03	Decision visibility	Whether decision-makers can see what they need, when they need it.	Measured, Claimed	ins_tech1, ins_tech3, atl_watch1, atl_watch3

ID	Anchor 20	Anchor 50	Anchor 80
01	Core activity is not captured in any system. Reporting is rebuilt by hand each time.	Systems in place per function but not integrated. Manual joins required to answer basic questions.	Systems cover the core and data flows between them without manual re-entry.
02	The owner does not trust the numbers, and is right not to.	Numbers broadly trusted, with known problem areas that get worked around rather than fixed.	Data quality actively managed. Reconciliations routine, discrepancies rare and traced to source.
03	Key questions cannot be answered without a manual exercise taking days.	Standard reports exist. Anything ad-hoc requires the owner or an analyst.	Decision-makers self-serve the numbers that matter, on a cadence that fits the decision.

1.6 · Process

Eight steps, then the closing steps that open the engagement. Snapshot mode alters four of them; see 1.2.

Step	Marking	What happens	Output
1 · Kickoff & data request	FENWICK	A 45 to 60 minute kickoff to align on goals and the one question the owner most wants answered. Confirm scope, timeline and the single success question. Issue the data request and the questionnaire, set a return date, set up the secure workspace and least-access credentials. Stop here if it has not happened; do not draft an agenda as though the meeting occurred.	Signed-off scope; data request issued
2 · Assemble & sanity-check	TOOL	Ingest sources into ledger/sources/ unmodified and register each with a ref, date and class. Sanity checks that must pass before proceeding: statements internally consistent, period coverage complete, revenue and headcount plausible against each other, obvious restatements flagged. A failed check is reported to Fenwick, never silently corrected.	Registered ledger; checks passed or reported
3 · Seven-lens diagnostic	AGENT	Score each criterion against the rubric with <code>scoring.score_lens()</code> . Every score requires evidence refs; the tool refuses an unsupported score. For each lens produce the score, the refs, the strengths, the risks, the gaps and an explicit confidence. Where evidence is insufficient, mark the lens insufficient rather than estimating. Run <code>benchmarks.lookup()</code> where comparables exist, and say so where they do not rather than reaching for an unrelated sector.	Seven scored lenses with evidence and confidence
4 · Management interviews	FENWICK	He runs every interview personally. Once transcripts or notes exist, run <code>interviews.triangulate()</code> to test claims against the ledger, producing three buckets: corroborated, contradicted, and untestable. Never smooth a contradiction, and never resolve one by deciding which party is more credible. Surface both and let Fenwick judge. SKIPPED ENTIRELY IN A SNAPSHOT.	Corroborated / contradicted / untestable

Step	Marking	What happens	Output
5 · Score & weight	TOOL	Run <code>scoring.compute_aperture_score()</code> . Deterministic, stamps weights and rubric version. Never hand-adjust a composite. Run <code>measures.baseline()</code> to pull ledger-sourced baselines and target ranges for each success measure.	Aperture Score, overall and by lens
6 · Name the binding constraint	HYBRID	Run <code>constraints.propose()</code> . It returns ranked candidates, each with the evidence for it, what argues against it, what would have to be true for it to be the binding constraint, and the test that would settle it with the component that runs that test. Do not run the test here; scoping it is the deliverable. You propose; Fenwick names.	Ranked candidates with the case against each
7 · Synthesis & report	AGENT	Run <code>render.insights_doc()</code> , which refuses to render an uncited figure. Contents are listed in 1.9. Findings, never conclusions.	The Business X-Ray report
8 · Readout	FENWICK	He edits, signs and delivers, in a 60-minute session. Recommend the next step and confirm the phase-gate decision. SKIPPED IN A SNAPSHOT, which is written only.	Client decision on next components

Closing steps, which open the engagement

Step	Marking	What happens
Validate the constraint with the owner	FENWICK	Pressure-test the named constraint against the owner's lived experience and any new data. Agree it in writing, and define what "solved" looks like in numbers.
Define success measures	HYBRID	Three to five measurable outcomes tied to the constraint, each with a ledger-sourced baseline and a target range. These become the Scoreboard KPIs in Phase 05, so vagueness here is expensive later.
Sequence the components	HYBRID	Propose which of Phases 02 to 05 the engagement needs and in what order, with the data prerequisites for each and the phase-gate dates. Fenwick scopes and prices.

If a Business X-Ray already exists when a full engagement opens, do not redo it. Extend it. That is the whole of SOP 02: the X-Ray is the baseline, and Insights frames the engagement around it.

1.7 · Data required

Two admissible sources: the business documents submitted, and the completed intake questionnaire. For an X-Ray, add the interviews and read-only system access. Nothing else is admissible, and a gap that remains after both sources are read is a finding written down as one.

A change carried in at Edition 01: the questionnaire counts as evidence in a Snapshot

The Snapshot was originally specified as "submitted documents only", and the Lumina sample was built that way, which is why its Market & Competition lens is unscored. Documents cannot answer why customers leave, what breaks if demand doubles, who the business loses to, or what would fall over if the owner took a month off. Four lenses were being penalised for asking questions no P&L answers.

Questionnaire answers enter as claimed evidence: Low confidence, capped below the Distinctive band, and never the sole support for the area named as the likely constraint. A questionnaire answer is promoted to measured when a document says the same thing, and then the document is the citation.

The consequence, and the reason the corroboration floor is proposed: once the questionnaire is admissible, all 31 criteria become reachable and the coverage floor stops discriminating.

Tier 1 · required. Without all four there is no Snapshot, and no X-Ray worth the name.

Source	What it must contain	Period and granularity	Criteria it feeds
Profit & loss	Revenue and cost of delivery separated from overhead, with whatever segmentation the accounts already carry.	Three full years plus year to date. Monthly where the system produces it.	fin-01, fin-04, ops-04
Balance sheets	Receivables, payables, inventory, debt. Comparable period ends, not a mix of dates.	Last four period ends. Four is what makes a cycle rather than a photograph.	fin-02
Transaction-level export	One row per transaction: date, customer ID, product or service, location, revenue, direct cost where held, and the person or resource that delivered it.	Twenty-four months minimum. Native export, never re-keyed or pre-aggregated.	cus-01 to cus-03, ops-01, ops-05, ldr-01, mkt-04, tec-02
Completed questionnaire	The shared sections plus the whole Aperture Insights section. Blanks are acceptable and informative; a wholesale blank lens is not.	Returned with the documents, not after them.	All seven lenses; sole source for 13 of the 31 criteria

The transaction export is the highest-value single item in the entire request: it feeds eight criteria across five lenses and it is the only document from which a genuinely new finding tends to emerge.

It is what produced the Lumina finding, 56% of patients with no second visit inside 90 days, from documents alone.

Tier 2 · strongly wanted. Each one lights a specific criterion.

Source	What it closes	Criteria
Revenue by product or service, and by location	Which parts of the business carry the profit rather than which are busiest.	fin-01, ops-04
Customer list	Customer ID, ZIP or postcode, acquisition channel, first transaction date. Names are not required and should not be sent.	cus-01, cus-03, mkt-04
Marketing spend by channel and month	Cost per acquired customer, and whether the heaviest channels bring the customers who stay.	cus-01
Product or service catalogue	Current prices and, where held, the direct cost to deliver one unit.	fin-03
Payroll or headcount by role, or an org chart	The shape of the organisation, spans of control, and where the pressure points are.	ldr-01, ldr-05
Staff list showing who can do what, plus any licence register	Whether qualified supply is the cap on billable volume, and how long it takes to relieve.	ops-05, ldr-04
Debt and lease schedule	What the business is committed to, and how that interacts with the cash cycle.	fin-02
Accountant's report and business tax returns	How much weight the figures can carry. Audited and internally prepared are both fine; they mean different things.	fin-05
Cash flow statements	Whether profit is turning into cash. Where none exist, the balance sheets carry fin-02 alone.	fin-02

Tier 3 · opportunistic. Ask, do not chase.

Source	Why it is worth a line
Top-vendor spend or AP summary	The only document that turns ops-06, supplier dependency, from a claim into a number. PROPOSED addition to the intake checklist.
Operating documentation, if any exists	The only document behind prc-01 and part of ldr-04. Its absence is as informative as its content. PROPOSED addition.
Chart of accounts	The account structure behind the P&L. Feeds fin-01 and fin-05.
Fixed asset or depreciation schedule	Feeds fin-02.
Prior reports, plans or valuations	Context. Not scored directly, but it tells you what the owner has already been told.

Format rules, and why each exists

Rule	Why
Native exports, CSV or XLSX	A PDF of a 12,000-row export has to be re-keyed, and re-keying inside a five-hour cap is where the cap goes. A PDF of a statement is fine.
One row per transaction	Aggregation destroys the analysis the export exists for. A monthly summary cannot produce a repeat rate, a concentration figure or a provider dependency.
The same customer ID across every file	The join between the transaction export and the customer list is what makes cus-03 and mkt-04 derivable. A re-anonymised ID that breaks the join costs two criteria.
No names, contacts, or clinical or account notes	Least access. An ID and a postcode do everything the analysis needs.
One delivery, one folder, plain filenames	The clock starts when the last Tier 1 item lands. Partial deliveries make that date ambiguous and the cap unmanageable.

Receipt and sufficiency check

Thirty minutes, run the moment the data lands and before a single lens is read. Its purpose is to score the intake rather than the business. Hard gates for a Snapshot; advisory for an X-Ray, where the kickoff call makes a gap recoverable.

#	Check	Passes when	If it fails
1	Manifest	Every file listed with its name, period and what it is. This list becomes the "what we read" section.	Nothing fails here. The list is the deliverable.
2	Reconciliation	Transaction export revenue agrees with P&L revenue for a matched period within 2%.	2 to 10%: proceed and state the variance. Over 10%: stop and query. The result is also the evidence for tec-02.
3	Period coverage	Financials cover three years; the export covers at least twenty-four months.	Under three years, fin-04 is unscorable. Under twenty-four months, say plainly that a true new-versus-returning split is not available.
4	Join	Twenty sampled customer references resolve in the customer list.	A broken join costs cus-03 and mkt-04. One re-export request is worth making.
5	Granularity	One row per transaction; direct cost present or absent and known which.	A pre-aggregated export cannot carry cus-02, cus-03 or ops-05.
6	Plausibility	Revenue against headcount and location count is credible for the sector.	Implausible pairs mean a scope problem: several entities, or books that are not the whole business.

#	Check	Passes when	If it fails
7	Questionnaire	Shared plus all seven lens sections returned, four fifths of fields answered.	One chase, with a date. A blank questionnaire removes thirteen criteria and takes Processes out entirely.
8	Projected coverage	Mark each of the 31 criteria scoreable or not. Coverage at or above 22; corroboration at or above 12.	Below either floor the composite is withheld. Take the decision now, not at hour four.

When the intake is short

Situation	Decision	What is said
Both floors clear, Tier 2 gaps	Proceed	The normal case. Gaps are named in "what we read and what was missing", each tied to the confidence it costs a named lens.
Coverage clears, corroboration does not	Return once	One request naming the two or three documents that would move it, with a date. If they do not arrive, the composite is withheld and the report ships as a written read with no Aperture Score.
Questionnaire largely blank	Return once	Twenty minutes of the owner's time carrying thirteen criteria. Worth one chase and only one.
No transaction export exists in any form	Convert or decline	The business is being asked to buy a diagnostic its own systems cannot support. An X-Ray, where the interview substitutes, or nothing.
A Snapshot heading past five hours	Convert	Say that the business needs an X-Ray, and stop. Do not absorb the hours.

Insufficient evidence is a finding, not a failure

A withheld composite is a legitimate outcome and it is not a refund event. The client still receives seven scored lenses with confidence ratings, the area most likely to be holding profit back, and the register of what would settle the rest.

What they do not receive is a single number built on a third of the evidence, and the reason they do not is the same reason the number is worth having when it does ship.

1.8 · How it is developed

The Aperture Insights Workbench, v0.1

The Phase 01 bench scores an X-Ray against the rubric and produces the Aperture Score. It lives in aperture-operations/Operational Process/01 Insights (X-Ray and Score)/.

Pipeline	src/template.src.html + rubric.json → python3 build.py → two editions.
Editions	Aperture Insights - STANDALONE Workbench.html (offline, sign-in gate), and site-dropin/private/method-lab/insights-workbench-tool.html (no gate, because the Method Lab passphrase already guards it).
The instrument is inlined	build.py reads rubric.json and injects it between the APERTURE:RUBRIC markers, so the tool and the Python engine read the same instrument. The build refuses to run if the lens weights do not total exactly 1.0, and warns loudly if the rubric is not ratified.
Never edit a built file	Change src/template.src.html or rubric.json and rebuild both. That rule is what kept the Analytics source 360 lines behind its own output for a month before anyone noticed.
The cross-check	python3 crosscheck.py. The Aperture Score is computed twice, in browser JavaScript and in scoring.py, which is the arrangement that drifts. The cross-check builds a deliberately awkward fixture across all 31 sub-criteria, drives the JS through headless Chromium, computes the same result in Python, and fails on any divergence in a lens score, the composite, the coverage or the withholding decision. Run it in the cloud container; it needs Playwright and Chromium.
Console	build_master.py assembles every bench into the Operations Console. Each bench is embedded exactly as its own build.py produced it, as an iframe with a blob URL. Rebuild a bench, re-run the master.

What v0.1 does, and what it does not

Built	Not yet built
All 31 sub-criteria with definitions, evidence classes and anchors	Intake import from the website questionnaire
An evidence ledger; a score without a reference is refused	The data-request tracker
Insufficient excluded from the mean; lens rollup; renormalised composite	Constraint candidates
Composite withheld below the coverage floor	The X-Ray report writer
Depth check: Snapshot below \$5M, score marked provisional	The corroboration floor, pending ratification

Built	Not yet built
Session save and load, with a rubric-version warning	

1.9 · Deliverables

The Business X-Ray

Section	What it contains
Executive summary	Leading with the answer to the owner's one question.
The seven-lens diagnostic	A score, a reason and an explicit confidence for each lens, in plain language, jargon removed.
Strengths, risks and gaps register	What is working, what to watch, and what could not be established.
Ranked constraint candidates	Each with the evidence to hand and what argues against it. A candidate with no disconfirming evidence has not been tested.
The baseline Aperture Score	Overall and by lens, with the weights and rubric version stamped.
Identified data gaps	What was missing, and what it cost.
The "where to dig" register	For each lens that warrants depth: the question, which component answers it, and what a good answer would be worth. This is the piece that earns the next engagement.
The prescription	Which components will move the needle, in what order, and the expected payoff.

The Aperture Snapshot

Seven pages, shorter and blunter, in this order: what a Snapshot is and is not, first page and in the client's language · what we read and what was missing · the seven-lens table with a confidence column · the finding · three other things the evidence shows · what a Snapshot cannot settle, four open questions with why they cannot be answered at this depth · the "where to dig" register, ranked · what it costs to take further. The parts that must not soften are the first, the sixth and the seventh.

1.10 · Quality checklist and failure modes

☐ Before readout
☐ Every headline number reconciles to the client's financials
☐ Each of the seven lenses has evidence, not just an opinion
☐ The Aperture Score is reproducible from the recorded inputs and weights
☐ The #1 constraint is singular, specific and defensible, and was named by Fenwick
☐ The report is readable by a non-analyst; jargon removed

☐ Before readout
☐ The prescription names concrete next components and why
☐ No borrowed framework appears: no SWOT grid, no five-forces diagram. Supplier power is scored at ops-06, buyer power at cus-03, durability at mkt-05, each with a figure and a ref
☐ All PII handled per the data-security standard

How this phase goes wrong

- Accepting the owner's diagnosis as the starting frame instead of as a hypothesis
- Scoring a lens on impression when the evidence is thin. Mark it insufficient instead
- Producing a constraint candidate with no disconfirming evidence, which means it was not tested
- Letting the report drift into recommendations, which pre-empts Phase 04 and usually anchors it wrongly
- Importing a strategy framework instead of scoring the criterion that already carries the same question
- Scoring mkt-05 on the strength of an advantage while ignoring its durability. "Customers love us" is not an answer to how long it would take to copy
- Mistaking the loudest problem for the binding one. Check whether each candidate is a constraint or a symptom of one upstream

1.11 · Handoff

Typed questions only. Quantitative to Marcus in Phase 02. Customer and geographic to Natalie in Phase 03. Strategic to Aria in Phase 04. Each states the question, the evidence refs behind it, and what a good answer looks like. Phases 02 and 03 may then run concurrently.

03

SEGMENT 03 OF 06 · QUANTIFY

Aperture Analytics™

Where the money is actually made and lost

Question it answers	Why is it happening?
You receive	Profit Map™ report and the .xlsx scenario model
Timing	3 to 4 weeks
Investment	From \$8,500: As a standalone segment. Included in the Full Method at \$25,000.
Your time	The questionnaire section below, answers to follow-up questions on the numbers, and the readout.
Runs it	marcus-analyst, under SOP 03
Bench	Aperture Analytics Financial Analysis Workbench v2.7
Manual source	Part 2 · version 01 · In force · bench built

VERSION 2.0 ISSUED 25 September 2026 · 2:48 PM CT SEGMENT 03 of 06

SEGMENT 03 · APERTURE ANALYTICS™

What it is

The quantitative segment. It shows exactly where your business makes and loses money: by product or service, by location and by customer.

It turns your financials into the handful of profit drivers you can actually move, and builds a model so decisions can be tested before they are made rather than defended after.

Every figure ties back to your own profit & loss statement, to the dollar.

What it does

- Maps profit by product or service, location and customer group
- Separates the few drivers carrying the profit from the costs quietly eroding it
- Measures how much cash is tied up in receivables and inventory, in dollars, not just days
- Finds your break-even, and the one variable the business is most sensitive to
- Models the specific decisions you are weighing, with every assumption visible and adjustable

What we need from you**What to provide**

Item	What it should contain	Why it matters
Everything from the diagnostic	The documents already provided for your Snapshot or X-Ray.	This segment builds on them rather than starting again.
Detailed revenue and cost	At the finest level your systems hold, by transaction or by job: date, line, location, customer, revenue, and cost where recorded.	Profit can only be cut as finely as the data allows.
Chart of accounts and current cost allocations	Your account structure, and how shared costs are currently split, if they are.	How shared costs are assigned decides which lines look profitable.
Volumes	Units, jobs, visits or orders in a typical month, by product or service and by location, with any seasonality.	Volume is what turns revenue into unit economics.
Pricing	Current prices, discounts, packages, memberships and negotiated rates, and when prices last changed.	Discounts are where margin most often leaks unseen.
Owner pay and unusual items	Owner compensation, one-off costs, and anything personal or related-party running through the books.	So the numbers reflect how the business really performs.
The decisions on the table	What you are weighing, and what you would need to see to be convinced either way.	The model is built to answer your question, not a generic one.

Questionnaire sections to complete · 4 sections, 15 questions

Section	Questions
How your numbers are kept	5
Price and mix	4
Volume	3
The decisions you want tested	3

OPERATING DETAIL · INTERNAL

How Aperture Analytics™ is run

Drawn from the Master Operating Manual, Edition 01, Part 2, version 01. The Manual governs; section numbers below are the Manual's own.

2.1 · What it is, and what it is not

The purpose of this phase is to separate what management believes from what the evidence demonstrates. Everything below serves that.

It shows exactly where the business makes and loses money, by product or service, by location and by customer, and turns the financials into the handful of profit drivers the owner can actually move. It then builds a scenario model so decisions can be tested before they are made rather than defended after.

The defining rule

Never do arithmetic in prose. Every number in every output comes from a tool Fenwick can re-run. A figure produced by reasoning is not evidence and it does not ship.

The second rule follows from it: the model reconciles to the P&L to the dollar. There is no such thing as close enough.

2.2 · Operating standard

Rule	What it means in practice
Reconcile to the dollar	On a reconciliation break: stop, report the break with its magnitude and where it sits, and wait. Never plug, never round away, never describe an unreconciled model as "broadly consistent". An unreconciled model that ships is how a practice loses a client permanently.
Normalise before analysing	Strip one-offs, owner compensation adjustments, related-party transactions and accounting-policy changes, and document every normalisation on the assumptions sheet. A normalisation that is not visible is a manipulation.
The allocation basis is a gate	Allocation choices decide who looks profitable: which line, which location, which customer. That is a judgment with commercial consequences, not a technical detail. Fenwick approves the basis before any figure is computed on it.
Contribution and absorption are different questions	Report contribution margin and fully-absorbed profit separately. Conflating them produces bad decisions: a line that is unprofitable fully-absorbed but contribution-positive should usually not be cut.

Rule	What it means in practice
Verify before rendering	recalc.py independently re-computes the workbook and verifies every formula. If recalc disagrees with the workbook, the workbook is wrong until proven otherwise. Nothing renders until recalc passes clean.
Stop before the recommendation	You draft the synthesis. Fenwick presents the readout and makes the recommendation.

2.3 · Methodology

Profitability and the Profit Map

Cut profit by product or service, by location and by customer segment, at the finest level the data actually supports and no finer. Identify loss-making and margin-diluting segments, and separate the engine, what carries the business, from the drag. Where a result is sensitive to the allocation choice, model both bases and show the difference: if the owner's favourite line is profitable under one basis and not the other, that is the finding.

80/20 and cost structure

Identify the few drivers carrying the profit and the costs quietly eroding it. Report both. The cost side is the one clients have usually never seen. Separate fixed, variable and step-fixed costs explicitly, because step-fixed costs are where capacity decisions hide.

Working capital and the cash conversion cycle

The formula	$CCC = DSO + DIO - DPO$, where $DSO = \text{receivables} \div \text{revenue} \times 365$, $DIO = \text{inventory} \div \text{COGS} \times 365$, and $DPO = \text{payables} \div \text{COGS} \times 365$.
Use period averages	Never a single closing balance. One balance-sheet date in a seasonal business is an artefact of the date, not a measure of the business.
Size it in money	$\text{Cash released} = \Delta \text{days} \times (\text{annual revenue} \div 365)$. Days do not move an owner. At \$5M revenue a fifteen-day improvement releases roughly \$205,000, routinely more than a year's net profit, and it arrives without debt, dilution or a new customer.
Report it against what they know	Put that figure beside their net profit and their headroom on the facility. The comparison is the finding. The ratio on its own is trivia.
Trend the components	DSO, DIO and DPO separately, over at least eight quarters. A CCC only means something against its own history and a sector comparable: 171 days is alarming for a distributor and unremarkable for capital equipment. Where no honest comparable exists, say so.

The DPO caution, and this is the one to get right

Do not recommend stretching payables as a default lever. The famous negative cycles belong to businesses with real supplier power, which no client in this segment has. A \$4M contractor who goes

to 90 days loses the merchant who fronts materials on a Saturday morning, and that relationship is worth more than the float.

For this segment the levers are DSO and DIO; treat DPO as close to fixed. The legitimate DPO finding is terms being left on the table: the supplier offers 45 days, the client pays in 10, and gets nothing for it. Where early payment buys a settlement discount, price the discount before recommending anything.

Sector	How much the cycle is worth leading with
Distribution and wholesale	Highest. The cycle alone can carry an engagement.
Home and trade services	High. Retainage, progress billing, van stock.
Professional and B2B services	Real but receivables-only.
Healthcare practices	Genuine where billing is through insurance; negligible where cash-pay.
Restaurants and cash-at-till retail	Usually negative already. Say so plainly and move on rather than manufacturing a finding to fill the section.

Break-even, sensitivity and scenarios

Compute break-even and sensitivity to price, volume, cost and mix, and report which variable the business is most sensitive to. That single fact reframes more owner conversations than any other output of this phase.

Model only the scenarios Fenwick has specified. Do not invent scenarios and do not model twelve when he asked for three. **Every scenario is driver-based:** no percentage growth applied to a total. Drivers are things the business can actually pull, such as price, volume, mix, capacity, conversion and retention. For each scenario state the drivers and their values, the class of each input, what would have to be true for it to occur, and the leading indicator that would show it happening.

2.4 · Process

Step	Marking	What happens
1 · Build the financial model	TOOL	build_financials.py assembles revenue and costs into a model that reconciles to the P&L to the dollar. Normalise, and document every normalisation.
2 · Approve the allocation basis	FENWICK	allocations.propose() returns a proposed shared-cost basis with the reasoning for each choice. Stop. Fenwick approves before any figure is computed on it.
3 · Map profitability	TOOL	build_profitmap.py cuts profit by product or service, by location and by customer segment. Contribution and fully-absorbed reported separately.
4 · 80/20 and cost structure	TOOL	pareto.rank_drivers() identifies the drivers carrying the profit and the costs eroding it. Fixed, variable and step-fixed separated explicitly.
4b · Working capital and CCC	TOOL	The cycle computed on period averages, converted to money, trended by component, and read against the sector guidance in 2.3.
5 · Break-even and sensitivity	TOOL	breakeven.sensitivity(). Report the variable the business is most sensitive to.
6 · Scenario and forecast model	AGENT	scenario.run() for the specified scenarios only, driver-based, each input labelled by class.
7 · Verify	TOOL	recalc.py, mandatory. It independently re-computes the workbook and verifies every formula. Nothing renders until it passes clean.
8 · Synthesis and readout	HYBRID	render.analytics_doc() builds the Profit Map report and the .xlsx model. You draft the synthesis; Fenwick presents and recommends. Stop before the recommendation.

2.5 · Data required

Phase 02 consumes everything Phase 01 collected, plus the detail that a diagnostic did not need. The intake questionnaire's Analytics section covers the shape of the accounts; the documents below carry the substance.

Source	What it must contain	Why
Transaction- or job-level revenue and cost	The finest level the systems hold. Date, line, location, customer, revenue, and cost where recorded.	Contribution margin cannot be computed above the level the data supports, and false precision is worse than a coarser honest cut.
Chart of accounts and existing cost allocations	The account structure, and any allocation already in force with its basis.	Feeds the allocation proposal at the Fenwick gate.
Volumes	Units, jobs, visits or orders in a typical month, by product or service, split by location, with seasonality.	an_vol1 to an_vol3. Volume is what turns revenue into unit economics.

Source	What it must contain	Why
Pricing structures	Current prices, discounts, packages, memberships and negotiated rates, and when prices last changed and what happened.	an_price1 to an_price4. Discount structures are where margin leaks invisibly.
Owner compensation and unusual items	Owner comp, one-off or unusual costs, and anything personal or related-party running through the business.	an_num5. This is the normalisation input, and it is a delicate question asked plainly.
Books context	Export granularity, whether delivery cost is separated from overhead, who closes the books and their confidence in them.	an_num1 to an_num4. Feeds fin-05 in Phase 01 and the assumptions sheet here.
The decisions on the table	What is actually being weighed, and what the owner would need to see to be convinced either way. Break-even if they believe they know it.	an_dec1 to an_dec3. This is what the scenario model is for; without it the model answers nobody's question.

The Private Company Intake Template (.xlsx), with its worked example, lives beside the bench in Operational Process/02 Analytics/ and is the structured way to collect the above. A copy is published for download on the website.

2.6 · How it is developed

Bench	Aperture Analytics Financial Analysis Workbench, v2.7, in Operational Process/02 Analytics/.
Pipeline	src/template.src.html (~812 KB) → python3 build.py → two editions, both around 3.1 MB because SheetJS and pdf.js are inlined. Fully offline.
Four copies exist and build.py writes two	After every build, copy the standalone to StandAlone/ and the site drop-in to aperture-website/private/method-lab/. Verify all four by checksum. The website copy is tracked in git and needs committing and pushing. Then re-run build_master.py so the Operations Console carries the new build.
Version stamp	build.py holds VERSION and a UTC timestamp with the time, and both appear in five places: the file comment, the sign-in screen, the masthead, the report cover footer, and the "which build produced this" block. The filename never changes between builds, so nothing else distinguishes the copy in a folder from the one last published. Bump it on every meaningful change.
The trap	The built HTML was once hand-edited, leaving the source 360 lines behind it. Building from that source would have silently deleted the EDGAR import, the multi-filing merge, the Import-tab reorganisation and the route cards. Never edit a built file.

Library recovery	build.py needs /tmp/libs/ and /tmp does not survive a restart. The CDN is blocked, so recover the libraries out of an existing built file, reading between each inlined marker and the next closing script tag.
Verification standard	Load both editions headless in the cloud container, run both worked examples and the Ashford session, and check: zero console errors, no page taller than the Letter box, printed sheet count equal to on-screen page count, contents entries matching their pages, the gate locked on the standalone and absent on the site copy, and the build stamp in all five places.

Three rules learned the hard way

- **Explanations are read from where they already live, never rewritten.** Help panels read the Method appendix and the ratios table, the contents page is read off the pages actually built, and both editions build from one template. Anything maintained twice disagrees with itself by the third revision.
- **Pagination measures with `getBoundingClientRect`, not `offsetHeight`.** `offsetHeight` rounds, and a page 1056.55px tall reports 1057 and slips a tolerance test. A page at min-height is insensitive to content until it overflows, so always fit-test the candidate including the piece being added. A block taller than a page must be splittable.
- **Never repeat a data URI per page.** The running-header mark was an image with a 19,800-character data URI, so a 57-page report carried 1.13 MB of identical base64 and the tab died on Print. Declare it once as a background image.

2.7 · Deliverables

The Profit Map report and the .xlsx scenario model. The workbook ships with a visible assumptions sheet, a reconciliation sheet showing the tie to the P&L, the allocation basis in plain language, and every input labelled by class.

The bench report runs in this order, and the order is deliberate: **cover · contents · executive summary · entity profile · valuation · statements · vertical · horizontal · DuPont · ratios · charts · findings · analyst commentary · synthesis · levers · sensitivity · Now/Next/Later · investment thesis · verification and sources · evidence and assumptions · method and formulas, last.** Method is reference, not argument, so it sits behind the ledger.

The investment thesis scores six weighted components out of 100: profitability 20, return on capital 20, quality of earnings 20, resilience 20, growth 10, evidence confidence 10. Going-concern doubt, or evidence under 35, forces "insufficient evidence". Framing switches to acquire for a private entity. Growth reads the compound rate at two thirds and the latest year's revenue direction at one third, after a fix in August 2026: reading CAGR alone once scored a business 100 on growth in a year revenue fell 5.9%.

2.8 · Quality checklist and failure modes

- ☐ **Before readout**
- ☐ The model reconciles to the client's financials to the dollar

☐ Before readout
☐ Cost allocations are documented, defensible, and were approved before use
☐ Profit cuts are shown at the level the data actually supports, with no false precision
☐ The scenario model's assumptions are explicit and adjustable
☐ recalc.py passed clean
☐ Recommendations tie back to the #1 constraint

How this phase goes wrong

- Reconciling "close enough". There is no such thing
- Applying an allocation basis before it is approved because it seemed obvious
- Presenting fully-absorbed profit as the profit, prompting a client to cut a contribution-positive line
- Percentage-growth forecasts dressed as driver-based models
- Skipping recalc.py because the model obviously works
- Softening the uncomfortable cut. If the owner's favourite line loses money that is the finding, and the allocation basis is shown so it can be challenged honestly
- Reporting a cash conversion cycle in days without converting it to cash, leaving the owner with a ratio instead of a decision
- Recommending a payables stretch because a case study did it, or computing the cycle off closing balances in a seasonal business

2.9 · Handoff

Pass Aria each opportunity with its size, its confidence, the assumptions it depends on, and what would have to be true for it to fail. Flag geographic or segment-shaped findings for Natalie rather than guessing at the spatial dimension.

04

SEGMENT 04 OF 06 · REVEAL

Aperture Intelligence™

Customers, market and geography

Question it answers	What opportunities does it reveal?
You receive	Customer & Market Map™, the layer guide, and a load-and-go ArcGIS package the client keeps
Timing	2 to 3 weeks
Investment	From \$8,500: As a segment of the Method. Included in the Full Method at \$25,000. \$18,000: Site Selection Study, a standalone fixed-fee engagement.
Your time	The questionnaire section below, a conversation about your locations and competitors, and the readout.
Runs it	natalie-explorer, under SOP 04
Bench	Not yet built
Manual source	Part 3 · version 01 · In force · bench to build

VERSION 2.0 ISSUED 25 September 2026 · 2:48 PM CT SEGMENT 04 of 06

SEGMENT 04 · APERTURE INTELLIGENCE™**What it is**

The segment that reveals where opportunity lives: who your best customers are, where they are, and where the next ones will come from, laid over real geography.

It answers the questions a spreadsheet cannot, such as how far customers really travel, where you are under-represented, and which competitors actually cost you business.

What it does

- Identifies your most valuable customers and where they come from
- Maps trade areas and drive times around each location
- Measures market penetration and competitive pressure by area
- Finds white space: places with demand you are not yet reaching
- Scores candidate sites when a new location is on the table
- Hands you the map package to keep, not just a report

What we need from you

Demographic, household and business-listing data is sourced by us, and its date is always shown.

What to provide

Item	What it should contain	Why it matters
Customer records with ZIP codes	Exported from your system and linked to purchase history by the same customer ID.	Without the link we can map where customers are but not what they are worth.
Location list	Every site: address, when it opened, size, how it performs in your view, and which one you would copy.	Your best site is the model new sites are measured against.
Competitors	The ones that actually cost you business, where they are, and anyone who has opened, closed or expanded nearby.	We can find competitor locations; only you know which ones matter.
Marketing spend by area or channel	Where spend goes geographically, and where referrals come from.	Turns a map into a spending decision.
Areas under consideration	Any towns, areas or sites you have been thinking about, however tentative.	Each one is scored against demand rather than dismissed.
How far customers travel	Roughly how far a typical customer comes, and the furthest you would call normal.	A reality check before drive-time bands are set.

Questionnaire sections to complete · 4 sections, 12 questions

Section	Questions
Your customers, geographically	4

Section	Questions
Your footprint	2
The competition	3
Where you might grow	3

OPERATING DETAIL · INTERNAL

How Aperture Intelligence™ is run

Drawn from the Master Operating Manual, Edition 01, Part 3, version 01. The Manual governs; section numbers below are the Manual's own.

3.1 · What it is, and what it is not

The purpose of this phase is to reveal where opportunity lives: the patterns that are invisible in rows and columns. Who the best customers are, where they are, and where the next ones will come from, laid over real geography.

It is the one capability no CPA and no marketing agency in the client's orbit can offer at all, which is why the Site Selection Study is priced on the value of the decision rather than on effort. An owner signing a ten-year lease is making a six-figure commitment.

The defining rule

You never see a street address. Geocoding runs locally, PII drops before analysis, and all work is on aggregated geography. Client address data does not leave the local environment for a third-party geocoding service. That is a confidentiality commitment to the client, not a preference.

3.2 · Operating standard

Rule	Why
Report the geocode match rate	A rate below roughly 95% biases every downstream map. Unmatched records are rarely random: they cluster in newer developments, rural routes and commercial addresses. State the rate and where the misses concentrate.
Suppress below the disclosure threshold	Any output cell containing fewer customers than the threshold is suppressed, and shown as suppressed rather than as zero or absent.
Drive time, not radius	Wherever a road network exists. Radius rings are a cartographic convenience that misrepresent real catchments across water, motorways and terrain.
Penetration over addressable households	Never over population, which flatters every result.
Cite vintage and source for every external dataset	An undated market figure is not evidence. A five-year-old estimate in a growth corridor is worse than no estimate.
Observed, modelled and estimated stay separate	Actual customer locations, estimated demand and competitor draw never share a symbology without a legend that separates them.

Rule	Why
Parameters are gated	Fenwick approves the trade-area parameters before any analysis runs, and the scoring criteria before any area is ranked. Drive-time bands and catchment definitions decide the answer; never choose a band because it produces a tidier map.

3.3 · Methodology

Segmentation before mapping

Segment on behaviour before demography. What customers do predicts value; who they are explains it. A segmentation that only splits by age and income is a demographic exercise, not a commercial one. Report revenue concentration explicitly: the share of revenue in the top decile of customers is often the single most consequential number in the phase.

Trade areas, penetration and competition

Build primary and secondary trade areas with drive-time rings, measure where demand concentrates, overlay demographics (age, sex, income, households, spend), then compute penetration by area and map competitors to find contested and open demand.

Site selection, and why it carries the heaviest caution

A site recommendation can be a seven-figure decision. Therefore three things are mandatory on the face of the output:

- **The sensitivity of the ranking to the parameters.** If candidate A beats candidate B only under one weighting, that appears on the output, not in a footnote.
- **What would disqualify each top candidate.** Lease cost, labour availability, zoning, cannibalisation of an existing location.
- **Cannibalisation modelled explicitly** wherever an existing location sits inside the trade area. A new site that moves revenue rather than adding it is a cost, not a growth move.

3.4 · Process

Step	Marking	What happens	Output
1 · Assemble and geocode	TOOL · LOCAL	geocode.py standardises and validates addresses, geocodes them, and drops street-level PII before anything else runs. Report the match rate.	Geocoded, de-identified customer layer
2 · Customer analytics	TOOL	segments.clv() segments the base and computes lifetime value by segment. Behaviour first, demography second.	Segments and CLV by segment
3 · Trade area and drive time	TOOL · LOCAL	Parameters approved by Fenwick first. gis_analysis.py for trade areas, drive-time rings and demand concentration; gis_layers.py for Voronoi tessellation, buffers, choropleths and H3 aggregation to GeoJSON.	Trade areas and demand concentration

Step	Marking	What happens	Output
4 · Penetration and competition	TOOL	demographics.overlay() for age, sex, income, households and spend by area; syncpoint.query() for competitor locations and business or household data.	Penetration and competitive map
5 · Site-selection modelling	AGENT	Scoring criteria approved by Fenwick first. site_model.score() ranks candidate ZIPs and areas by opportunity score, with sensitivity, disqualifiers and cannibalisation.	Ranked site-selection model
6 · Package and readout	HYBRID	render.intelligence_doc() builds the Market Map report, layer guide and ArcGIS package. Fenwick reviews, owns and presents every site recommendation. Stop before delivery.	GIS package delivered; gate decision

3.5 · Data required

Source	What it must contain	Why
Customer records with addresses or postcodes	Exportable from a named system, joinable to purchase history by a shared customer ID, one row per transaction where history is included.	int_geo1 and int_geo1b. Without the join, this phase can map where customers are but not what they are worth.
Location list	Every site: address, opened when, size, and how it is performing in the owner's view. Which one is the model they would repeat, and why.	int_fp1, int_fp2. The reference site is what the site model is calibrated against.
Competitor locations	The competitors that actually cost the client business and where they are; anyone who has opened, closed or expanded nearby; what they do better.	int_comp1 to int_comp3. We can source locations, but not which ones actually cost them business.
Marketing spend by area or channel	Spend broken down geographically where it exists, and where referrals come from.	int_grow2, int_grow3. Spend by area is what turns a penetration map into a decision.
Candidate areas	Any towns, areas or sites already under consideration, however half-formed.	int_grow1. Listed even as hunches; they are scored against demand rather than dismissed.
Travel behaviour	Roughly how far customers travel, and the furthest the owner would call normal.	int_geo2, int_geo3. The reality check on the drive-time bands before they are set.

External data (demographics, household counts, business listings) is sourced by the practice and always carries its vintage. The disclosure threshold and the local-geocoding rule apply to everything in this table.

3.6 · How it is developed

No bench exists yet. Operational Process/03 Intelligence is an empty folder, and the tooling named in 3.4 is specified rather than built. Until a bench exists, this phase runs on the analyst's GIS environment directly, which means the standards in 3.2 are enforced by discipline rather than by code. That is the weakest link in the phase and it is the reason to build it.

When it is built it follows the pattern the other two benches set: **src/template.src.html plus a parameter file → build.py → a standalone edition and a Method Lab drop-in, then build_master.py to fold it into the Operations Console**. The instrument gets inlined rather than copied, built files are never edited, and anything computed twice gets a cross-check.

3.7 · Deliverables

Artefact	What it contains
Customer & Market Map™	Customers mapped, trade areas, demand clusters, segments and lifetime value, penetration and competitive benchmarks, and the where-to-grow recommendation.
The layer guide	What each layer is, how it was built, its vintage and its symbology. The document that makes the package usable a year later.
The ArcGIS package	Load-and-go, with recommended symbology. The client keeps it. It must open cleanly with the layer guide beside it.
Site-selection model	Where a site question was in scope: ranked areas, the sensitivity of the ranking, the disqualifiers, and the cannibalisation estimate.

Cartography follows 0.7. Every map carries a legend, a scale bar, a north arrow, the data vintage, the projection and the source. A client should be able to read the map without you standing next to it.

3.8 · Quality checklist and failure modes

☐ Before readout
☐ Geocoding accuracy checked; unmatched records reported, not hidden
☐ PII removed from the analysis layer after geocoding
☐ Trade areas and drive times use documented, approved parameters
☐ The growth recommendation is backed by the client's own demand data
☐ Cannibalisation modelled wherever an existing location sits in the trade area
☐ The ArcGIS package opens cleanly with the layer guide

How this phase goes wrong

- Radius rings where a road network exists
- Penetration computed over population instead of addressable households

- Undated demographic data presented as current
- A choropleth of raw counts rather than rates, which just maps population
- Ranking sites before the criteria are approved
- Ignoring cannibalisation and presenting moved revenue as new revenue

3.9 · Handoff

Combine external market intelligence with Marcus's quantitative findings and Lena's diagnostic before passing evidence to Aria. Every opportunity carries its geography, its size, its data vintage and its confidence.

05

SEGMENT 05 OF 06 · NAVIGATE

Aperture Compass™

The few moves that matter, in order

Question it answers	Where should we go next?
You receive	Opportunity Matrix™ and the one-page Focus Plan™
Timing	1 to 2 weeks
Investment	From \$8,500: As a segment of the Method. Included in the Full Method at \$25,000.
Your time	A candid working session on your limits and priorities, and the readout.
Runs it	aria-navigator, under SOP 05
Bench	Not yet built
Manual source	Part 4 · version 01 · In force · bench to build

VERSION 2.0 ISSUED 25 September 2026 · 2:48 PM CT SEGMENT 05 of 06

SEGMENT 05 · APERTURE COMPASS™**What it is**

By this point we understand the business, the numbers, the customers and the market. This segment answers the harder question: what should you actually do about it?

The aim is not a long list of recommendations. It is the few moves that matter, sequenced so your team can execute them.

What it does

- Brings every finding from earlier segments into one set of options
- Scores each option on value, effort, risk and its effect on the main constraint
- Tests the leading options against a downside scenario
- Sequences the winners into Now, Next and Later, each with an owner

What we need from you

Most of what this segment works from is our own earlier work, so there is little new to send. It starts once the main constraint has been agreed with you in writing.

What to provide

Item	What it should contain	Why it matters
What is already on the table	Every move you are considering, however half-formed, and which one your instinct favors.	So nothing you are already weighing is missed, and your instinct can be checked honestly.
Your real limits	Capital available and how fast it can be released; whose time each move would take; what is off the table; and how long you can wait for a payoff.	A plan that ignores these fails at execution.
How things get done here	An initiative that stalled and what stopped it; who needs persuading and what persuades them; who would own delivery day to day.	The difference between a roadmap and a plan.

Questionnaire sections to complete · 3 sections, 9 questions

Section	Questions
What's already on the table	2
Your real limits	4
How things get done here	3

OPERATING DETAIL · INTERNAL

How Aperture Compass™ is run

Drawn from the Master Operating Manual, Edition 01, Part 4, version 01. The Manual governs; section numbers below are the Manual's own.

4.1 · What it is, and what it is not

By the time work reaches this phase the practice understands the company, the numbers, the customers and the market. This phase answers the harder question: what should we do about it?

The mission is not a hundred recommendations. It is **the few moves that matter**, sequenced so they can actually be executed. The job of this phase is subtraction, and it is the only phase whose quality is visible in what it leaves out.

The prerequisite is absolute

Without a constraint named by Fenwick there is nothing to sequence against and this phase cannot run. Do not proceed on a "likely" constraint, and do not name one in order to get started.

4.2 · Operating standard

Rule	Why
Every candidate cites its upstream finding	A move with no traceable origin does not enter the list. It goes to Fenwick as a question about what is missing.
Weights are agreed with the owner before scoring, never after	Setting weights after seeing the options is how a predetermined answer gets dressed as analysis. Fenwick sets the criteria and weights and agrees them with the owner.
The matrix stamps its criteria version and weights	Same reason the rubric does. A ranking that cannot be reproduced is an opinion with a table around it.
Constraint first, expansion last	Expanding through a constraint multiplies the constraint. A growth initiative sequenced ahead of the fix that enables it converts a good idea into an expensive failure. This is the single most valuable judgment the phase makes.
The Focus Plan is one page	That is the discipline, not a formatting preference. If the matrix has thirty rows and the plan has fifteen items, this phase has failed at the only job that distinguishes it.
Propose, do not decide	You propose the order. Fenwick owns the order that goes to a board.

4.3 · Methodology

Consolidation, and why the same move arrives three times

The same opportunity often arrives wearing different clothes: as a margin gap from Phase 02, an underserved trade area from Phase 03, and a process weakness from Phase 01. Merge them and keep all three evidence trails. A move supported from three directions is not three moves, and it is stronger than any of them alone.

The weighted matrix

Criterion	What it asks
Value or return	What it is worth, sized from the Phase 02 model rather than estimated here.
Complexity or effort	What it actually takes, in whose time.
Risk or confidence	How sure we are, and what the downside looks like.
Impact on the binding constraint	The criterion that distinguishes this Method from a generic prioritisation exercise. Weight it accordingly.

Stress-testing

For each top move, **state the plausible future in which it is the wrong call**. If you cannot construct one, it has not been stress-tested; go back. Flag moves that are only attractive in the base case. A move that fails in the downside is not necessarily wrong, but the owner must choose it knowingly.

Sequencing, and the three things a roadmap must surface

- **Dependencies.** What must finish before what.
- **Resource conflicts.** Two initiatives needing the same person in the same quarter are one initiative and a delay.
- **Displacement.** For every "Now" item, what it pushes out. A roadmap assuming infinite capacity is a wish list.

4.4 · Process

Step	Marking	What happens	Output
1 · Consolidate	AGENT	opportunities consolidate() de-duplicates candidate moves from every prior phase into one defined list, each citing its upstream finding.	Master opportunity list
2 · Score	TOOL	matrix.score() against value, complexity, risk and impact on the constraint, on weights agreed beforehand. The matrix stamps the criteria version.	Scored, weighted matrix
3 · Stress-test	TOOL	scenario.stress_test() against the Phase 02 scenario model and plausible futures, naming the future in which each top move is wrong.	Scenario-tested ranking

Step	Marking	What happens	Output
4 · Sequence	HYBRID	sequencer.order() by dependency and gate, constraint first and expansion last, surfacing dependencies, resource conflicts and displacement. You propose; Fenwick owns the order.	Sequenced roadmap
5 · Focus Plan and readout	HYBRID	render.compass_doc() builds the Opportunity Matrix and the one-page Focus Plan. Fenwick presents. Stop.	Approved Focus Plan

4.5 · Data required

This phase consumes outputs rather than client documents, with one exception: the owner's real limits, which nobody else has asked for and which decide whether the plan is executable.

Input	What it must contain	Where it comes from
The named constraint	Singular, specific, defensible, agreed with the owner in writing.	Phase 01, named by Fenwick
Success measures	Three to five, each with a ledger-sourced baseline and a target range.	Phase 01 closing steps
Shortlisted moves	Each with size, confidence, the assumptions it depends on, and what would have to be true for it to fail.	Phase 02 handoff
Growth options	Each with its geography, size, data vintage and confidence.	Phase 03 handoff
What is already on the table	Every move under consideration, however half-formed, and which one the owner's instinct favours. Recorded so it can be checked later, not to flatter it.	comp_tab1, comp_tab2
The owner's real limits	Capital available and how fast it releases; whose time this would take and whether they have it; what is genuinely off the table; appetite for a payoff in year three rather than this year.	comp_lim1 to comp_lim4
How things get done here	An initiative that stalled and what actually stopped it; who needs persuading and what persuades them; who would own delivery day to day.	comp_done1 to comp_done3

The last two rows are the difference between a roadmap and a plan. An initiative with no owner, or one that needs the person who is already the bottleneck, fails at execution regardless of how it scored.

4.6 · How it is developed

No bench exists yet. Operational Process/04 Compass is an empty folder. The matrix and the sequencer are specified in 4.4 and are currently run by hand, which means the weights-before-options rule and the criteria version stamp depend on discipline rather than on a tool that refuses. Of the three benches still to build, this is the cheapest: a weighted matrix with a stamped criteria version, a displacement column and a one-page renderer is a small piece of software with a disproportionate effect on defensibility.

4.7 · Deliverables

Artefact	What it contains
Opportunity Matrix™	Every option scored on the same criteria and weights, with the criteria version stamped, and a defensible ranking.
Focus Plan™	One page. Now / Next / Later. Each entry carries the move, why now, the owner, the dependency, what it displaces, the success measure, and the expected effect on the constraint.

4.8 · Quality checklist and failure modes

☐ Before readout
☐ Every opportunity scored on the same criteria and weights
☐ The ranking is reproducible and the weights are recorded
☐ The sequence puts the constraint first; expansion is gated on evidence
☐ Each item has an owner, milestones, dependencies and a displacement
☐ Every top move has a stated future in which it is the wrong call
☐ The plan is genuinely board-ready: one page tells the story

How this phase goes wrong

- Recommendations with no traceable upstream evidence
- Weights adjusted after the options were visible
- Expansion sequenced ahead of the constraint fix
- A Focus Plan that is a list of everything, ranked, which is not a plan
- No displacement stated, so every item is "Now" and nothing is
- Success measures too vague to instrument in Phase 05

4.9 · Handoff

Hand Thomas the approved initiatives with their intended outcomes and the measures that would show them working. Vague success criteria here become an unusable dashboard three months later: **"improve customer experience" is not a measure. "First-response time under 4 hours on 90% of tickets" is.**

06

SEGMENT 06 OF 06 · PERFORM

Aperture Atlas™

Keeping the plan alive, month after month

Question it answers	How do we sustain and improve results?
You receive	The live Scoreboard™ and KPI system, and the monthly one-page Signal Brief
Timing	Stand-up of 2 to 3 weeks, then recurring
Investment	\$900/mo: Single location \$1,500/mo: 2–4 locations \$2,500/mo: 5+ locations Terms: 6-month minimum, 30 days' notice thereafter. Or pay annually and get 12 months for the price of 10.
Your time	A setup session, authorizing read-only connections, then a few minutes a month with the Signal Brief.
Runs it	thomas-operator, under SOP 06
Bench	Not yet built
Manual source	Part 5 · version 01 · In force · bench to build

VERSION 2.0 ISSUED 25 September 2026 · 2:48 PM CT SEGMENT 06 of 06

SEGMENT 06 · APERTURE ATLAS™**What it is**

The ongoing segment. It turns the plan into a living management system: a live Scoreboard™ of the few numbers that matter, connected to your own systems.

Each month you receive a one-page Signal Brief on what changed and what needs attention, and your Aperture Score™ is re-scored against the original baseline so progress is measured, not assumed.

What it does

- Defines the few KPIs that matter, each with a target, an owner and an alert threshold
- Connects them to your systems with read-only access
- Flags variances and anomalies against your baseline and plan
- Delivers a one-page Signal Brief every month
- Re-scores the business so you can see progress against where you started

What we need from you**What to provide**

Item	What it should contain	Why it matters
What you watch today	Which numbers you check weekly and monthly, and where you look them up. The report you receive and do not read, and why.	Tells us what the Scoreboard must be, and what it must not be.
What good looks like	For each number that matters, the value that means on track and the value that means call me. How quickly you need to know.	Sets the targets and alert thresholds.
Rhythm and ownership	Who owns each number, by name. Which existing meeting this should feed, and who else should see it.	A number without an owner does not get acted on.
System access	Which systems can be connected, and who authorizes it. Any that cannot connect and would need a manual upload.	Read-only, least access, always.

Questionnaire sections to complete · 4 sections, 10 questions

Section	Questions
What you watch today	3
What good looks like	2
Rhythm and ownership	3
Wiring it up	2

OPERATING DETAIL · INTERNAL

How Aperture Atlas™ is run

Drawn from the Master Operating Manual, Edition 01, Part 5, version 01. The Manual governs; section numbers below are the Manual's own.

5.1 · What it is, and what it is not

This phase turns a consulting engagement into a living management system. It is also the phase that converts a project into a retainer, so treat the Scoreboard as a product, not as a leftover artefact.

It keeps the strategy live: an always-on executive view of the few KPIs that matter, wired to the client's systems, with the cadence that keeps the plan honest quarter after quarter, and a re-scored Aperture Score that proves progress against the baseline set in Phase 01.

The defining rule

Escalate, never act. Surface the signal and route it. Acting is a human decision every time, and contacting a client is never the agent's to make.

5.2 · Operating standard

Rule	Why
Every KPI traces to an initiative in the approved Focus Plan	A metric that measures nothing anyone decided to do is noise on a screen.
Leading indicators, not only lagging	A dashboard of lagging measures reports history. Leading indicators let someone intervene while it still matters.
Every KPI has a named human owner	A person, not a department. An unowned metric is never acted on.
Keep the set small	Fifteen KPIs an operator reads beats sixty they ignore.
Thresholds are a judgment, not a formula	What counts as red is proposed with reasoning and approved by Fenwick. Never silently adjusted because a client was uncomfortable with the colour.
Re-score on the stored rubric version	Never on current or improved weights. Comparability across periods is the entire value of a baseline. If the rubric has changed, restate history explicitly and label the restatement.
Reconcile on every refresh	Every dashboard value ties to its source system on build and on every refresh. A dashboard that quietly drifts from the accounting system is worse than no dashboard, because it is trusted.
Build it to be handed over	The client owns the Scoreboard and can operate it themselves. A client who cannot operate it will stop valuing it, and the retainer dies quietly.

5.3 · Methodology

Designing for an operator with four minutes and no analyst

- Lead with what needs a decision, not with a wall of green tiles.
- Every number shows its trend and its comparison basis. A value with no context is trivia.
- State the refresh time and the source on the face of the dashboard.
- Charts follow the house system: the title states the finding, one idea per chart, no chart junk.

Signal over noise

Before raising a deviation, ask whether a competent operator would change a decision because of it. If not, log it, do not raise it. **A brief that cries wolf monthly gets skimmed by month four, and then the one that mattered gets skimmed too.**

For each signal state: what moved, by how much, against what expectation, what changed, and the question it raises. Where the cause is unknown, say so plainly. That is often the most useful line in the brief.

Escalation routing

Signal type	Routes to	Which means
Something structural is wrong in the business	Lena	Phase 01 diagnosis, or a re-score
The numbers do not behave as modelled	Marcus	Phase 02 analysis
A market, customer or geographic shift	Natalie	Phase 03 intelligence
The plan itself needs adjusting	Aria	Phase 04 strategic adjustment

Fenwick reads every brief before the client does, and he decides when a signal becomes a conversation. Fresh data re-opening the next constraint is how the loop restarts and how a new component engagement begins.

5.4 · Process

Step	Marking	What happens	Output
1 · Define the KPIs	HYBRID	kpi.define(). Each KPI carries a definition, source system, target, cadence, owner and RAG thresholds. Propose thresholds with reasoning; Fenwick approves.	Approved KPI definitions
2 · Wire the data	TOOL	pipelines.wire() for least-access pipelines from accounting, POS, CRM and marketing, pulling the minimum fields the defined KPIs require. feed.validate() reconciles every value to its source.	Live, validated feed

Step	Marking	What happens	Output
3 · Build the Scoreboard	TOOL	build_dashboard.py assembles KPIs against targets, trends, and the Market Map and forecast where relevant.	Published Scoreboard
4 · Set the cadence	FENWICK	His step. Reporting rhythm and steering-meeting design are his call: monthly steering, quarterly review, KPI owners assigned.	Operating cadence live
5 · Re-score and steer	AGENT	anomaly.detect() for variance against baseline and plan; scoring.compute_aperture_score() on the stored weights and rubric version; render.signal_brief() for the one-page monthly brief.	Proven progress; next constraint identified

5.5 · Data required

Input	What it must contain	Where it comes from
The Focus Plan and its success measures	Initiatives, intended outcomes, and measures specific enough to instrument.	Phase 04 handoff
The Aperture Score baseline	The overall and per-lens scores, with the weights and rubric version stamped.	Phase 01
What they watch today	Which numbers are checked weekly and monthly, and where each one is looked up. Which report they receive and do not read, and why. What they would check every morning if it were one click away.	atl_watch1 to atl_watch3
What good looks like	For each number that matters, the value that means on track and the value that means call me. How quickly they need to know when something goes wrong.	atl_good1, atl_good2
Rhythm and ownership	Who owns each number, as a person. Which existing meeting this should feed and when it runs. Who else should see the Scoreboard.	atl_rhythm1 to atl_rhythm3
System access	Which systems can be connected directly and who authorises it. Any system that cannot be connected and would need a manual upload.	atl_wire1, atl_wire2. Also feeds tec-01 in Phase 01.

The report they receive and never read is the most useful question in this set. It tells you what a dashboard must not be.

5.6 · How it is developed

No bench exists yet. Operational Process/05 Atlas is an empty folder. The Scoreboard is intended to run on the Aperture platform (SyncPoint AI), and until that exists the phase cannot be delivered at the price it is sold at, because a hand-built dashboard per client does not carry a \$900 a month tier.

This is the highest-value bench still to build, and not only for delivery. Atlas is the recurring revenue line and the reason the practice is worth more than the sum of its projects: **a project ends, a Scoreboard renews**. Every month it runs is also a month of fresh data that re-opens the next constraint and originates the next component engagement.

5.7 · Deliverables

Artefact	What it contains	Cadence
The Scoreboard™	A live executive KPI dashboard the client keeps and can operate. KPIs against targets, trends, RAG, and the Market Map or forecast where relevant.	Live, continuous
Automated reporting	The agreed report, to the agreed people, on the agreed rhythm.	Per the cadence
The Signal Brief	One page. What moved, by how much, against what expectation, what changed, and the question it raises.	Monthly
The re-scored Aperture Score	Computed on the stored weights and rubric version, as before-and-after proof.	Each period

5.8 · Quality checklist and failure modes

☐ Before it goes live, and on every refresh
☐ Every KPI has a definition, source, target, owner and cadence
☐ Dashboard numbers reconcile to the source systems
☐ Access is least-privilege and documented
☐ The re-scored Aperture Score uses the identical stored baseline method
☐ The client can operate the Scoreboard themselves
☐ Fenwick has read the brief before the client does

How this phase goes wrong

- KPIs that measure what is easy to pull rather than what was decided
- Re-scoring on updated weights, destroying period comparability
- A dashboard of lagging indicators only
- Alert fatigue from raising every variance
- Silently adjusting a threshold because a client was uncomfortable with the colour
- Building a Scoreboard the client cannot operate without the practice

5.9 · Handoff, and the loop

Atlas is the durable end state, and it is also the front door to the next engagement. Fresh data re-opens the next constraint and can trigger a new component. Confirm the recurring agreement is in

place, and when a signal warrants a phase rather than a conversation, route it per 5.3 and open the work properly rather than absorbing it into the retainer.